



Blink Memory Reduction (@BlinkOn5)

{haraken, bashi, tasak, hajimehoshi}@chromium.org


2015 Nov



Memory team projects




- Oilpan
 - haraken@, keishi@, peria@, yutak@
- **Memory reduction** (\Leftarrow This talk is about this)
 - haraken@, bashi@, hajimehoshi@, tasak@
 - Working with primiano@, ssid@, ruuda@
- Tab resource reclaiming
 - kouhei@, tzik@



History




- As of 2015 June, memory reduction was raised as one of the priorities of the web-platform team
 - At that point, no one knew even what percentage of renderer's memory is consumed by Blink
 - Since then, we've been working on **understanding Blink's memory** and **identifying sweet spots** to reduce it
- 

Strategy




- Three steps for memory reduction:
 1. Tooling
 2. Understanding
 3. Reducing



Strategy



- Step 1: Tooling
 - Build up [memory-infra](#) and provide **consistent data** to break down the renderer's memory into pieces
 - Provide **reproducible telemetry benchmarks and metrics** with which developers can keep track of their reduction efforts



Strategy




- Step 2: Understanding
 - Using the tools, identify sweet spots that need reduction



Strategy



- Step 3: Reducing
 - Brainstorm ideas with teams who own the code area
 - Enable the teams to get involved in the reduction
 - Getting more teams involved is key to have a bigger impact



Agenda of this talk



1. Tooling


- Provide consistent data
- Provide reproducible benchmarks and metrics

2. Understanding

- Identify sweet spots

3. Reducing

- Brainstorm ideas
- Enable more teams to get involved




Tooling



Allocator unification

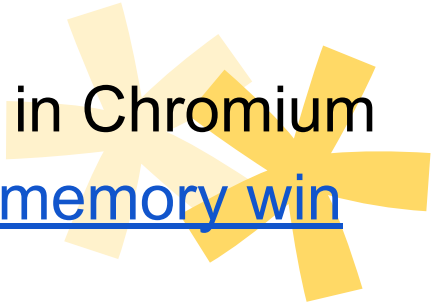


- As of 2015 June, Blink mixed four memory allocators
 - PartitionAlloc
 - Oilpan
 - tcmalloc
 - system allocator
- It was crazy...



Allocator unification



- We removed tcmalloc and system allocator from Blink
 - PartitionAlloc
 - Oilpan
 - ~~—tcmalloc~~
 - ~~—system allocator~~
 - Note: Why only from Blink?
 - We experimented with unifying all allocators in Chromium but couldn't observe any clear performance/memory win
- 

memory-infra



- We supported PartitionAlloc and Oilpan in [memory-infra](#)
 - We're adding more detailed profiling information to **get more understanding**

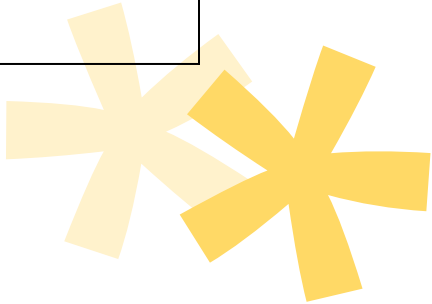
Process ▾	Total resident ▾	Peak total resident ▾	PSS ▾	Private dirty ▾	Swapped ▾	blink_gc ▾	cc ▾	discardable ▾	gpu ▾	gpumemorybuffer ▾	malloc ▾	partition_alloc ▾	skia ▾	v8 ▾
■ Browser (pid 16542)	99.8 MiB	105.5 MiB ↗	68.2 MiB	40.9 MiB	0.0 B		36.4 MiB	0.0 B	10.0 MiB	0.0 B	23.7 MiB		71.7 KiB	2.1 MiB
■ Renderer (pid 16594)						4.6 MiB		2.6 MiB			4.2 MiB	3.3 MiB	5.5 MiB	13.5 MiB
■ GPU Process (pid 16605)									6.8 MiB		53.5 MiB			
■ Renderer (pid 16707): chrome://tracing						3.4 MiB	15.1 MiB	4.0 MiB	3.1 MiB		3.6 MiB	4.7 MiB	132.6 KiB	16.2 MiB
■ Renderer (pid 16734): Kentaro Hara (@xharaken) Twitter						9.5 MiB	8.3 MiB	4.7 MiB	3.1 MiB		3.4 MiB	14.3 MiB	18.2 MiB	27.9 MiB
Total	99.8 MiB	105.5 MiB	68.2 MiB	40.9 MiB	0.0 B	17.5 MiB	59.9 MiB	11.3 MiB	22.9 MiB	0.0 B	88.4 MiB	22.3 MiB	23.9 MiB	59.8 MiB
Allocator details														
Allocator ▾	size ▾	effective_size ▾	allocated_objects_size ▾	decommittable_size ▾	discardable_size ▾	slot_size ▾	total_pages_size ▾	virtual_committed_size ▾	virtual_size ▾	active_pages ▾	decommitted_pages ▾	empty_pages ▾		
*partition_alloc	14.3 MiB	14.3 MiB	12.4 MiB	660.0 KiB	252.0 KiB	8.6 MiB	14.4 MiB	17.9 MiB	30.6 MiB	174	114	10		
allocated_objects	12.4 MiB	12.4 MiB →												
*partitions	14.3 MiB ①	2.0 MiB ←	12.4 MiB	660.0 KiB	252.0 KiB	8.6 MiB	14.4 MiB	17.9 MiB	30.6 MiB	174	114	10		
▶ buffer	7.5 MiB	1.0 MiB	6.0 MiB	660.0 KiB	240.0 KiB	8.5 MiB	11.5 MiB	9.3 MiB	18.6 MiB	87	74	10		
▶ fast_malloc	4.9 MiB	678.3 KiB	4.5 MiB	0.0 B	12.0 KiB	87.6 KiB	2.1 MiB	6.3 MiB	8.0 MiB	71	38	0		
▶ layout	2.0 MiB	280.7 KiB	1.9 MiB	0.0 B	0.0 B	3.0 KiB	828.0 KiB	2.3 MiB	4.0 MiB	16	2	0		
node	0.0 B	0.0 B	0.0 B	0.0 B	0.0 B			0.0 B	0.0 B					

More detailed profiling



- What types of objects and how many objects are allocated in Blink ([per-object-type profiler](#))

Object type	Size	Count
Vector<RuleData>	12256 bytes	3
Vector<RuleFeature>	800 bytes	6
ImmutableStylePropertySet	26 bytes	2112
...



More detailed profiling




- Where large objects are allocated ([allocation-site profiler](#))

```
StringImpl 1020 KB
```

```
<--- AtomicString::AtomicString()
```

```
<--- ScriptResource::script()
```




More detailed profiling



- Retaining relationships that cross allocator boundaries
 - Blink retains a lot of memory in other allocators in the renderer process

[HTMLLinkElement::m_resource](#) => locked discardable memory (620 KB)

[FontCustomPlatformData::m_typeface](#) => Skia (700 KB)



Q4/Q1 Roadmap



- Add **more profiling information** to memory-infra
 - Per-object-type profiler
 - Allocation-site profiler
 - Cross-allocator retaining relationships
- Add **reproducible telemetry benchmarks and metrics** with which developers can keep track of their reduction efforts




Understanding



Scope




- Scope:
 - Understand the breakdown of renderer's memory
- The following slides are focusing on memory only inside Blink, but our scope is to understand and reduce renderer's memory (especially renderer's memory retained by Blink)



Scope




- Short-lived apps or long-lived apps?
 - For short-lived apps: memory after 10 seconds after a page load matters
 - For long-lived apps: memory after 1 hour after a page load & a lot of user interactions matters
 - We're now focusing on short-lived apps because short-lived apps would be more important in mobile devices, but long-lived apps are on our radar too
- 

Break down the memory into pieces




- Break down renderer's memory
 - **Renderer = V8 + PartitionAlloc + Oilpan + Skia + ...**
- Break down the PartitionAlloc's memory
 - **PartitionAlloc = Buffer + FastMalloc + Node + Layout**
- Break down each partition into per-object memory
 - **Buffer = StringImpl + Vector + HashTable + ...**
 - **FastMalloc = SharedBuffer + ...**



Break down of renderer's memory




Page	Total (*1)	PartitionAlloc	Oilpan	V8	Skia	Discardable	cc	malloc
pinterest.com	78.06 MB	34.35%	1.76%	13.32%	4.32%	20.50%	17.93%	7.82%
facebook.com	70.13 MB	33.02%	1.25%	8.63%	16.64%	17.11%	14.97%	8.38%
worldjournal.com	56.71 MB	41.57%	1.76%	11.24%	6.19%	7.05%	18.74%	13.45%
wikipedia.org	44.32 MB	34.82%	2.26%	11.92%	2.28%	9.03%	23.69%	16.01%
reddit.com	72.10 MB	21.15%	1.56%	9.72%	14.08%	27.74%	14.63%	11.12%
wordpress.com	77.31 MB	26.83%	1.78%	31.67%	5.10%	5.17%	14.55%	14.89%
plus.google.com	41.84 MB	27.11%	2.99%	13.98%	3.66%	19.12%	14.11%	19.03%
blogspot.com	57.57 MB	25.23%	1.74%	32.35%	3.43%	6.95%	18.24%	12.08%
mail.google.com	50.73 MB	17.91%	1.72%	7.10%	16.39%	23.65%	21.07%	12.15%



Break down of renderer's memory

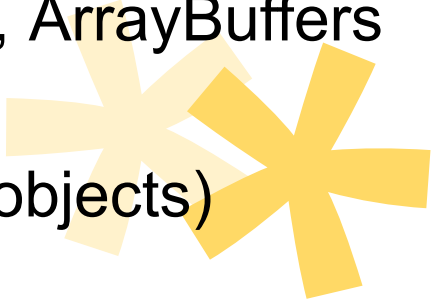


- In common cases V8 is the largest consumer, but in these pages **Blink is the largest consumer (20 - 40%)**
 - because the page set is chosen that way (i.e., they are key 10 pages where Blink's reduction is a key)
- Some (a lot?) of the memory in Discardable, CC and Skia is **retained by Blink**



Break down of Blink's memory




- Since we've not yet fully shipped Oilpan, most Blink objects are allocated in PartitionAlloc
 - PartitionAlloc has four partitions:
 - Node (for Nodes)
 - Layout (for LayoutObjects)
 - Buffer (for Vectors, HashTables, StringImpls, ArrayBuffers etc)
 - FastMalloc (for SharedBuffers and all other objects)
- 

Break down of PartitionAlloc's memory




Page	Total of PartitionAlloc	Buffer	FastMalloc	Node	Layout
pinterest.com	26.76 MB	75.74%	21.93%	1.17%	1.17%
facebook.com	23.30 MB	35.90%	58.46%	1.44%	4.19%
worldjournal.com	23.16 MB	34.08%	40.92%	9.11%	15.89%
wikipedia.org	15.56 MB	37.25%	37.15%	9.74%	15.86%
reddit.com	15.26 MB	47.88%	39.89%	3.89%	8.35%
wordpress.com	20.92 MB	52.88%	38.61%	4.33%	4.18%
plus.google.com	11.41 MB	66.19%	27.10%	3.08%	3.63%
blogspot.com	13.86 MB	51.69%	36.87%	3.55%	7.89%
mail.google.com	9.14 MB	58.46%	31.71%	3.33%	6.50%



Break down of PartitionAlloc's memory




- The largest consumer is the Buffer partition
 - StringImpls, Vectors, HashTables, ArrayBuffers etc
 - The second largest consumer is the FastMalloc partition
 - SharedBuffers, all DOM objects except for Nodes and LayoutObjects
 - sizeof(Node) and sizeof(LayoutObject) don't really matter
- 

Break down of the Buffer partition




Page	Total of the Buffer partition(*1)	StringImpls	Vectors	HashTables	Others
pinterest.com	9.51 MB	90.19%	2.64%	6.98%	0.19%
facebook.com	2.91 MB	41.23%	40.50%	17.00%	1.26%
worldjournal.com	3.22 MB	45.61%	39.00%	14.68%	0.71%
wikipedia.org	2.00 MB	65.18%	7.30%	26.63%	0.89%
reddit.com	2.89 MB	41.21%	42.39%	16.21%	0.20%
wordpress.com	4.94 MB	55.32%	29.88%	14.50%	0.29%
plus.google.com	2.62 MB	81.56%	7.23%	10.77%	0.44%
blogspot.com	2.89 MB	66.53%	14.02%	17.88%	1.57%
mail.google.com	2.11 MB	30.77%	54.24%	13.99%	1.00%



Break down of the Buffer partition




- The largest consumer is StringImpls
 - We confirmed that more than 90% of the large StringImpls are JavaScript source code
- The second largest consumer is Vectors and HashTables
 - We confirmed that 25% of the Vectors are consumed by the unused region
 - We confirmed that 70% of the HashTables are consumed by the unused region of small HashTables



Break down of the FastMalloc partition




Page	Total of the FastMalloc partition	SharedBuffers	Other objects
pinterest.com	2.38 MB	15.93%	84.07%
facebook.com	6.45 MB	74.69%	25.31%
worldjournal.com	2.90 MB	53.12%	46.88%
wikipedia.org	2.19 MB	6.25%	93.75%
reddit.com	2.32 MB	13.47%	86.53%
wordpress.com	3.58 MB	27.36%	72.64%
plus.google.com	0.92 MB	13.22%	86.78%
blogspot.com	2.32 MB	22.88%	77.12%
mail.google.com	1.01 MB	0.00%	100.00%



Break down of the FastMalloc partition




- The largest consumer is SharedBuffers
 - They are used as backing storage of Resources
- Other objects have a very long tail
 - The 2nd - 5th largest consumers depend on webpages
 - Reducing sizeof(DOM object) won't contribute to reducing Blink's memory



Summary




- The largest consumers inside Blink are:
 1. StringImpls
 2. SharedBuffers
 3. Vectors + HashTables
- BTW, how much memory is consumed by Vectors and HashTables as **absolute values**?
 - Only < 1 MB... (even in the Blink-heavy pages)
 - It's not worth reducing...



What does it mean?




- Inside Blink, the only objects we should really reduce would be **StringImpls and SharedBuffers**
- Reducing other objects won't contribute to reducing renderer's memory



What does it mean?




- After reducing StringImpls and SharedBuffers, **we should look at outside Blink**, not inside Blink
 - “Outside Blink” consumes **60 - 80%** (even in the Blink-heavy pages)
 - “Outside Blink” doesn’t mean that Blink is not related
 - Blink retains a bunch of memory in Discardable, Skia, CC etc
- Remember that our goal is to reduce renderer’s memory



What we have learned




- Memory reduction based on a “guess” is NOT likely to contribute to a meaningful reduction
 - Reducing sizeof(DOM object) won't have an impact
 - We experimented with discarding various Blink caches each TL listed in [this spreadsheet](#), but most of them didn't have an impact



What we have learned




- It is important to drive the reduction efforts **more tactically**
 1. Understand the breakdown
 2. **Identify key problem areas (sweet spots)**
 3. Reduce the memory at a pinpoint



What we have learned




- It is really hard to collect **consistent and reproducible** results
 - It is a substantial amount of work to confirm the reasonableness of various memory metrics
 - Example:
 - Render's private memory >> sum(memory usage of each allocator)
 - Where does the dark matter come from...?
 - Various "size" notions
 - resident, effective, active, committed...



Q4/Q1 roadmap



- Get more understanding on **outside Blink**
- **Understand the real world**
 - UMAs
 - Deep reports
 - OOM
- **Identify X projects** that would have the biggest impact
- Profile long-running apps



Reducing



Strategy




- Our strategy is:
 - Provide good tooling and understanding
 - Brainstorm ideas for reduction
 - And **enable more teams to get involved** in the reduction (instead of only the memory team working on the reduction)
 - That way we can move faster



On-going reduction projects



- [MemoryPurgeController](#) (\Rightarrow Memory team)
- [CompressableString](#) (\Rightarrow Memory team)
- More aggressive Resource pruning (\Rightarrow Loading team)
- [Making V8 \$\Leftrightarrow\$ Blink cycles collectable](#) (\Rightarrow V8 team)




MemoryPurgeController



- When Blink should save memory, MemoryPurgeController dispatches purgeMemory()


```
class MemoryPurgeController {  
    void purgeMemory(...) { // Dispatched when Blink should save memory  
        for (auto& client : m_clients)  
            client->purgeMemory(...);  
    }  
    HashSet<MemoryPurgeClient*> m_clients;  
};  
  
class MemoryCache : public MemoryPurgeClient {  
    void purgeMemory(...) override { pruneAll(); } // Example  
};
```



MemoryPurgeController




- When does MemoryPurgeController dispatch the purgeMemory()?
 - After 10 seconds after a tab goes inactive
 - When Blink receives MemoryPressureListener's notifications



MemoryPurgeController




- What should the `purgeMemory()` do?
 - Discard [discardable caches in Blink](#)
- Even though the `purgeMemory()` is dispatched against only inactive tabs & memory-pressure situations, we must be careful about its performance impact



CompressableString



- The largest consumer of PartitionAlloc is large Strings
- More than 90% of the large Strings are JavaScript source code
 - They are not discardable (used by V8)
 - The only option to reduce the memory is to compress the Strings (\Rightarrow Introduce **CompressableString**)



CompressableString



- We confirmed that if we gzip large Strings, we can reduce PartitionAlloc's memory **by 10 - 60% (17% in average)**


Page	Reduction rate
pinterest.com	63%
facebook.com	15%
theverge.com	11%
worldjournal.com	3%
wikipedia.org	19%
reddit.com	10%
wordpress.com	14%
plus.google.com	5%
blogspot.com	9%
mail.google.com	21%



More aggressive Resource pruning




- The second largest consumer of PartitionAlloc is SharedBuffers
- SharedBuffers are used as backing storage of Resources
 - Memory reduction of Resources is important



More aggressive Resource pruning



- Actually, memory reduction of Resources is more important than the data implies, because **Resources also retain memory in Discardable and Skia**
 - Some Resources use a locked memory in Discardable
 - Image/FontResources can hold caches in Skia



More aggressive Resource pruning




- Low-hanging fruits:
 - [HTMLLinkElement should clear CSSStyleSheetResource](#) when it finishes parsing the stylesheet
 - This reduces **1.2 MB from pinterest.com**, **800 KB from theverge.com**
 - [GlyphPage should discard GlyphPageTree](#)
 - This reduces **200 KB from theverge.com**
 - [FontCustomPlatformData should discard cached font-faces](#) in Skia



More aggressive Resource pruning




- Goals:
 - Don't keep alive ResourcePtrs longer than needed
 - Discard various caches retained by Resources (unless it regresses performance)
 - Visualize the retaining relationships in memory-infra



Making V8 ⇔ Blink cycles collectable



- In common cases, V8 is the largest consumer of renderer's memory
- At the very least, **Blink should not be the culprit of V8 memory leaks**

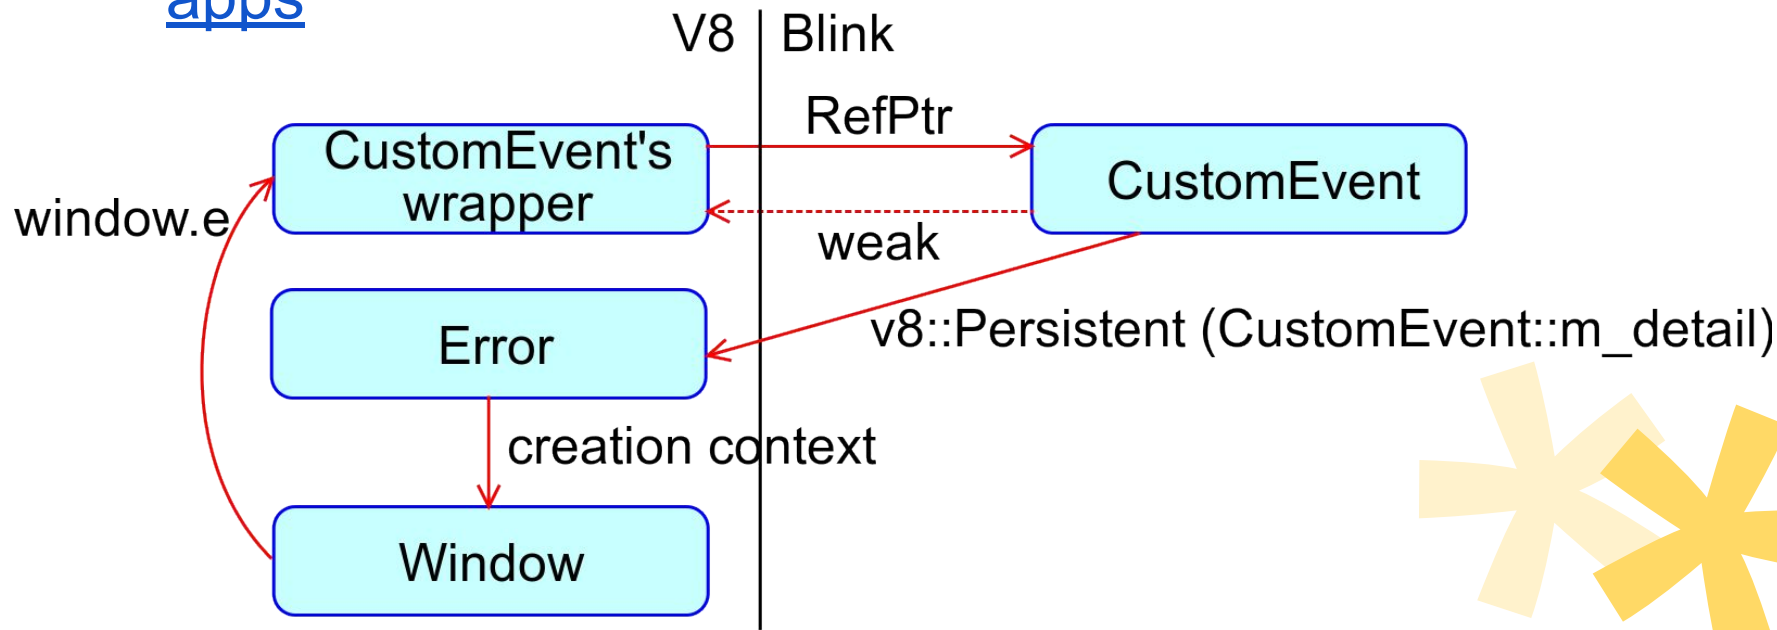


Making V8 \leftrightarrow Blink cycles collectable



- Example: `e = new CustomEvent("foo", {detail: new Error()});`
 - This was a cause of [a lot of window leaks in Polymer apps](#)

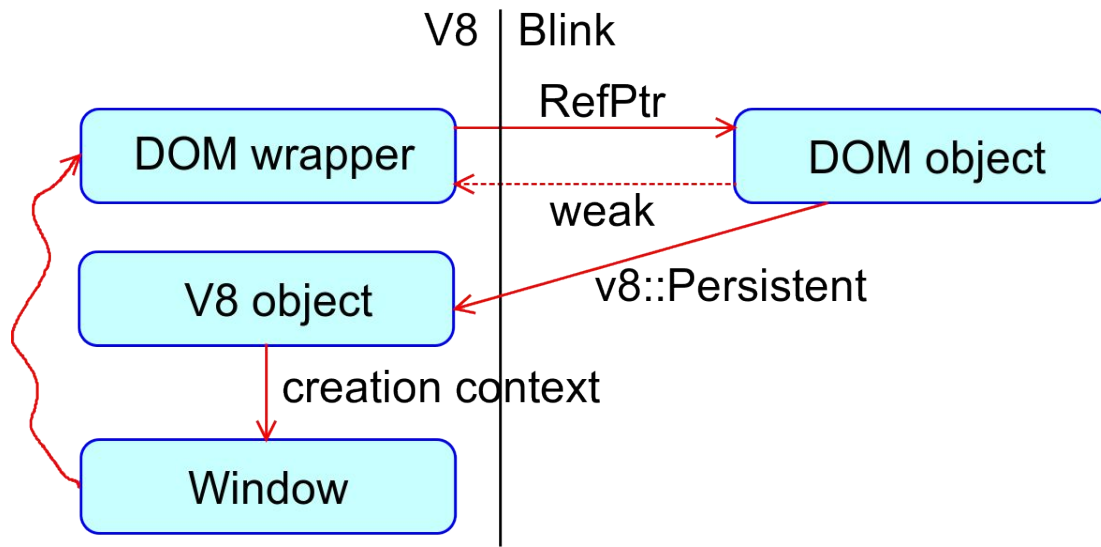
[apps](#)



Making V8 \leftrightarrow Blink cycles collectable



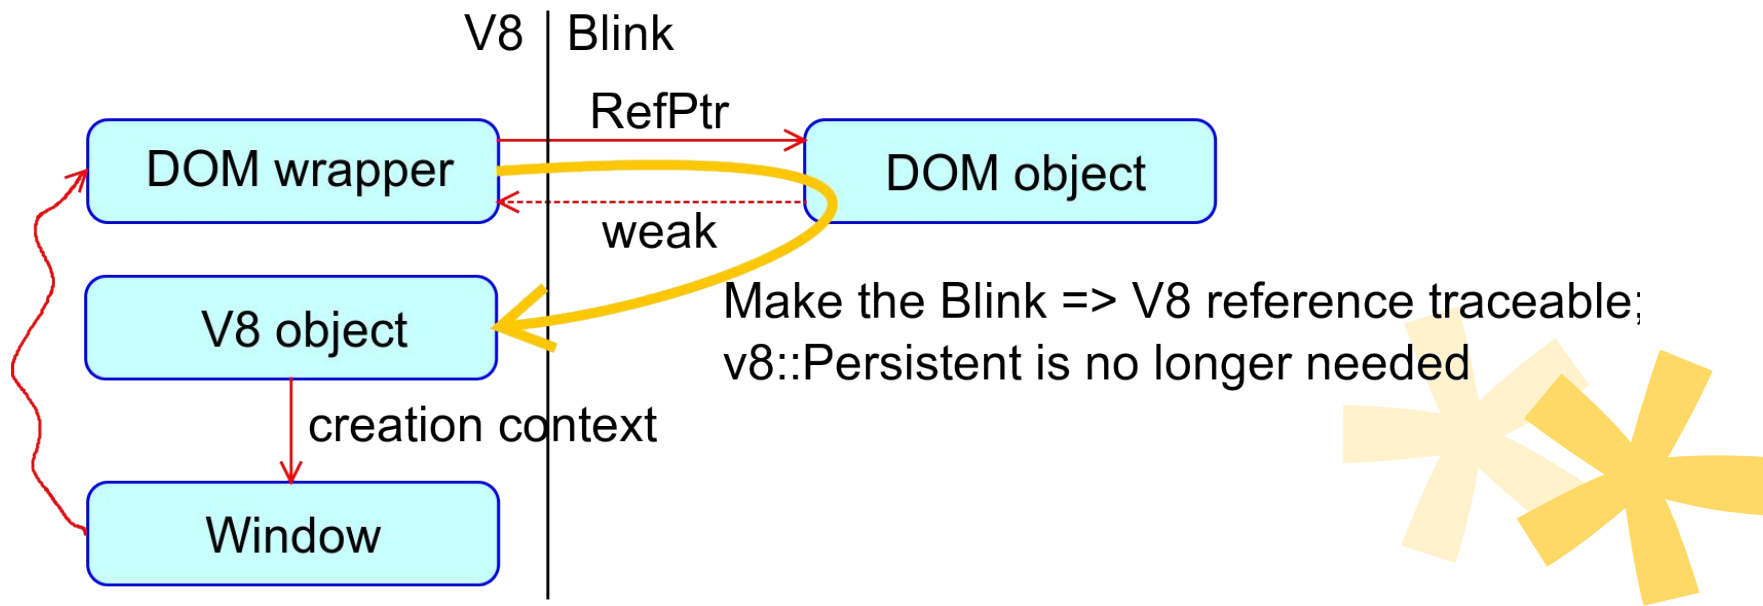
- In general, a V8 object has a strong reference to the window object that created the V8 object
- Thus **if a Blink object retains a `v8::Persistent` handle to the V8 object, it retains the entire window object...**



Making V8 \Leftrightarrow Blink cycles collectable



- Idea: Remove all v8::Persistent handles by **making V8 \Leftrightarrow Blink cycles collectable**
 - All V8 leaks caused by Blink will be gone



Summary




- Various reduction projects are going on:
 - MemoryPurgeController
 - CompressableString (for StringImpls)
 - More aggressive Resource pruning (for SharedBuffers)
 - Making V8 \Leftrightarrow Blink cycles collectable
- As we identify more key problem areas, we'll work with teams who own the area and brainstorm ideas for reduction



Q4/Q1 roadmap



- Establish the following development flow:
 - Provide good tooling
 - Identify key problem areas
 - Work with teams who own the area and brainstorm reduction ideas
 - Enable the teams to set memory-reduction OKRs and work on the reduction




Conclusions



Scope

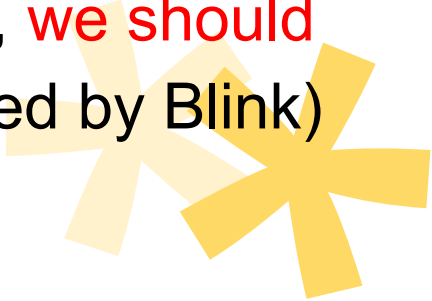


- Understand and reduce renderer's memory



Current understanding



- In Blink-heavy pages, Blink can be the largest consumer (20 - 40%)
 1. StringImpls
 2. SharedBuffers
 3. Vectors and HashTables
 - After reducing StringImpls and SharedBuffers, **we should look at outside Blink** (some of them are retained by Blink)
- 

Strategy



1. Tooling


- Provide consistent data in memory-infra
- Provide reproducible telemetry benchmarks and metrics

2. Understanding

- Identify sweet spots

3. Reducing

- Brainstorm ideas
- Enable more teams to get involved



Strategy



- Our goal is to establish a development flow that **enables each team to set memory-reduction OKRs and get involved in the reduction**
 - Instead of only the memory team working on the reduction :)
- **Empowering TRIM** for that direction!

