

# The Blink Scheduler

*London is minding your tasks*

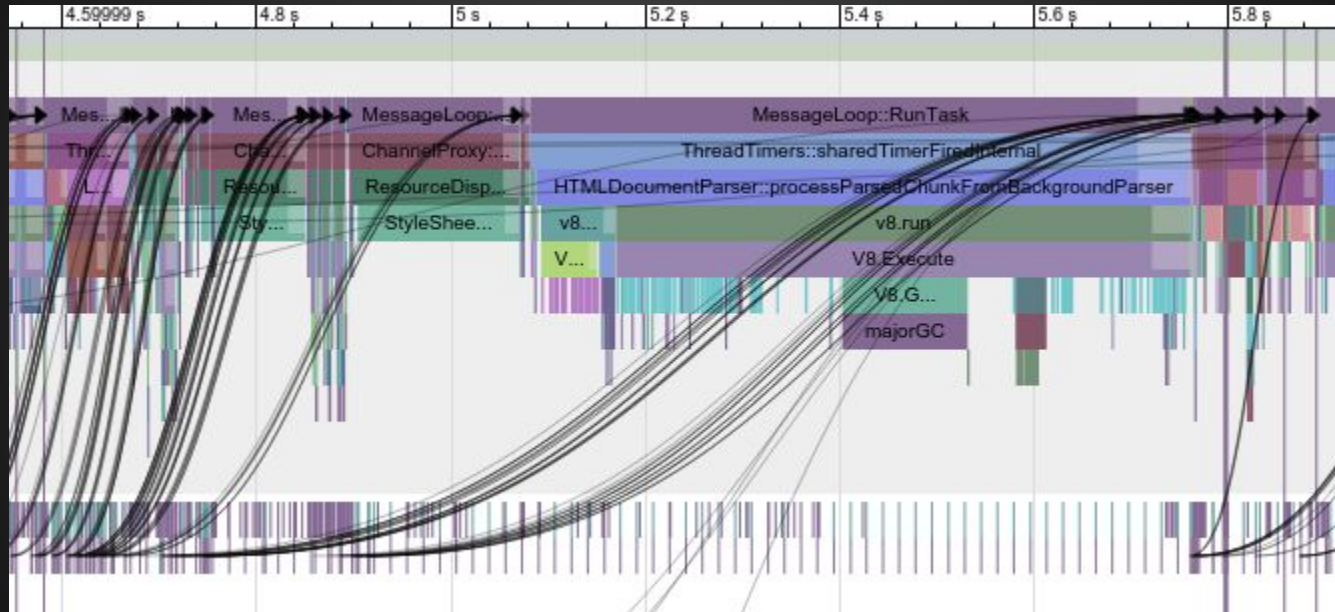
BlinkOn 3, Nov 5th 2014

{alexclarke, petrermak, picksi, rmcilroy, skyostil}@chromium.org

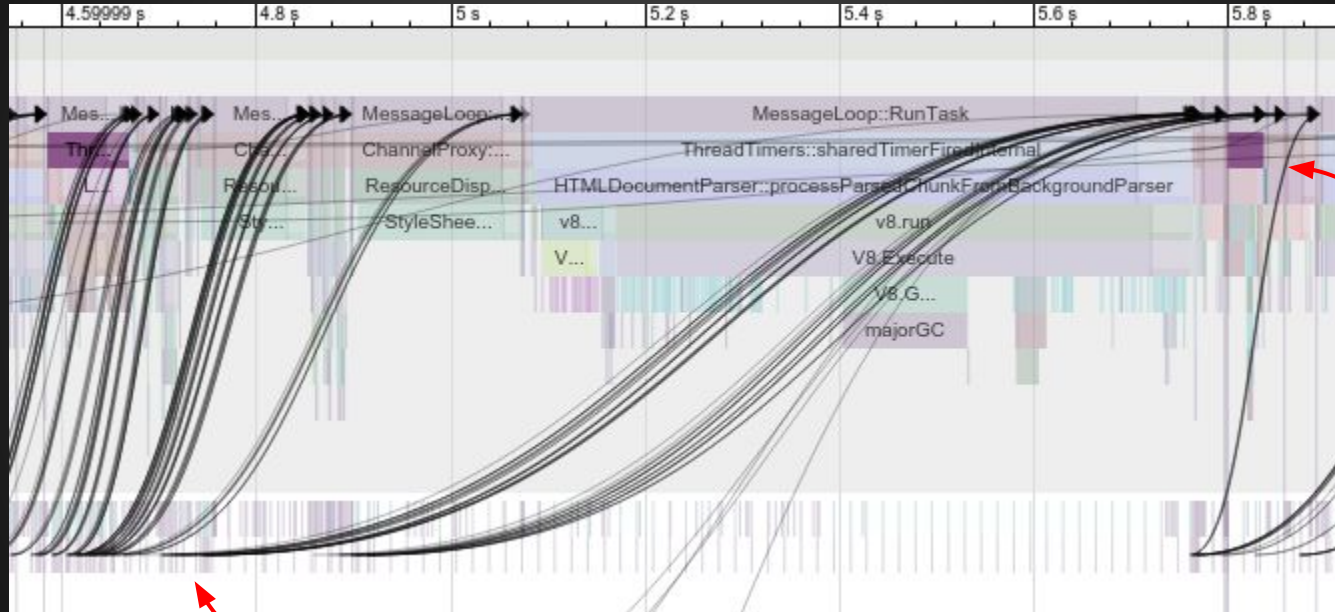
# The main problem: Traffic jams

- High priority tasks (e.g. user input) can get stuck behind slow tasks.
- Delays of up to 1 second when scrolling

# The main problem: Traffic jams



# The main problem: Traffic jams

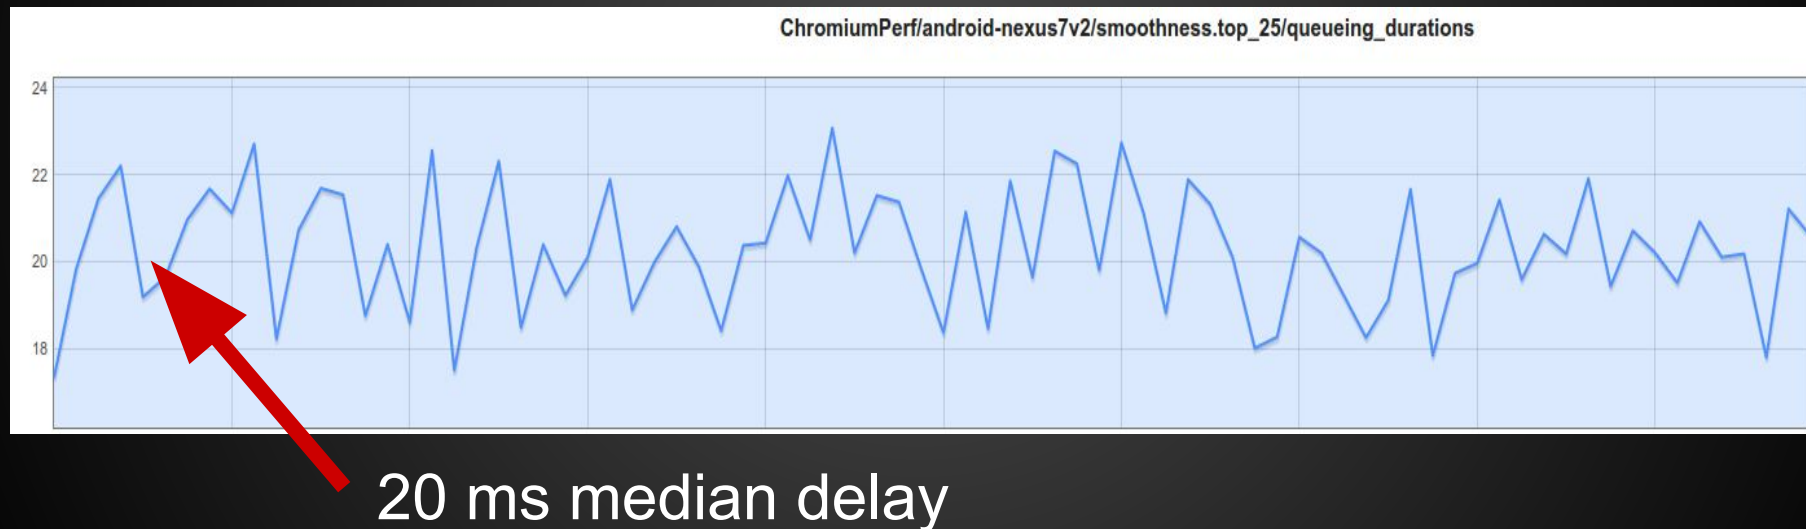


Composer task  
posted here

...and  
executed  
here, after 1  
second

# Issue: How do we measure traffic jams?

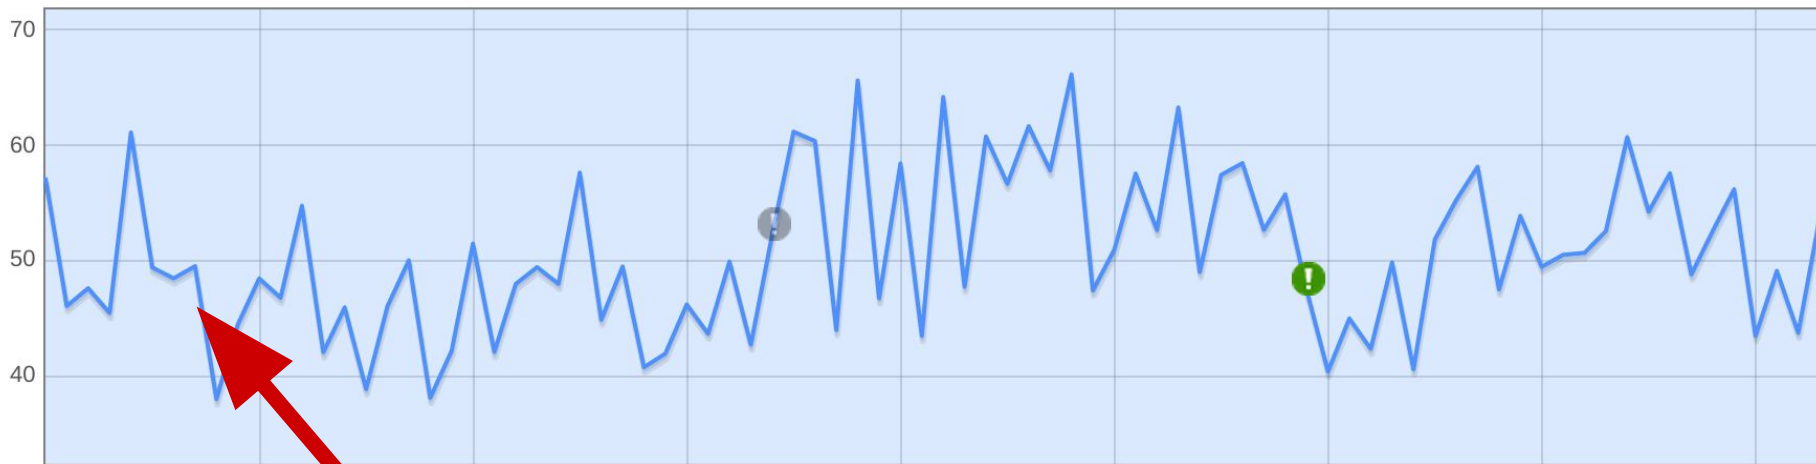
`queueing_durations` track how long compositor tasks are stuck in the message loop queue.



# Issue: How do we measure traffic jams?

`mean_input_event_latency` tracks how long user input takes to hit the GPU.

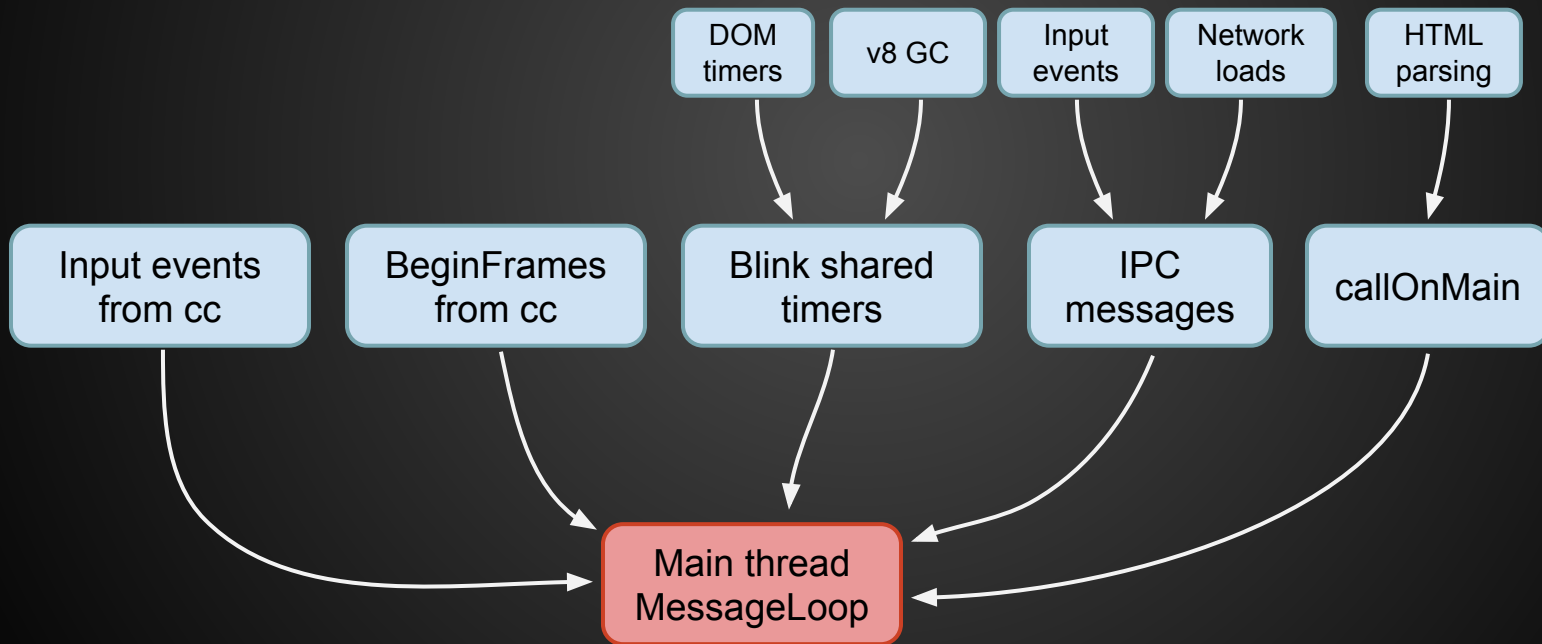
ChromiumPerf/android-nexus4/smoothness.sync\_scroll.key\_mobile\_sites/mean\_input\_event\_latency



50 ms median delay!

# Task sources

Renderer main thread tasks come from many sources:



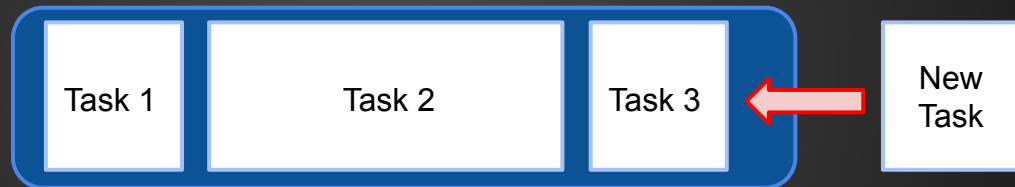
# Goal: Perfect Scheduling

- *Perfect scheduling* makes sure the right things get done at the right time.
- We're doing the same stuff, just in a sensible order.
- Enable well-written content to take advantage of good scheduling.

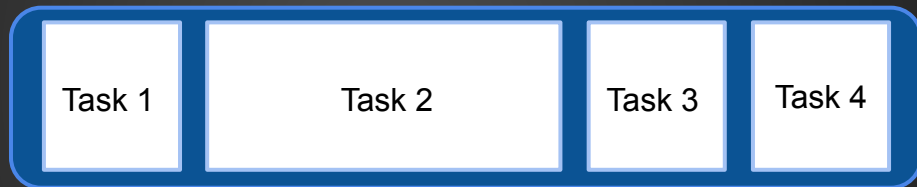


**All about scheduling**

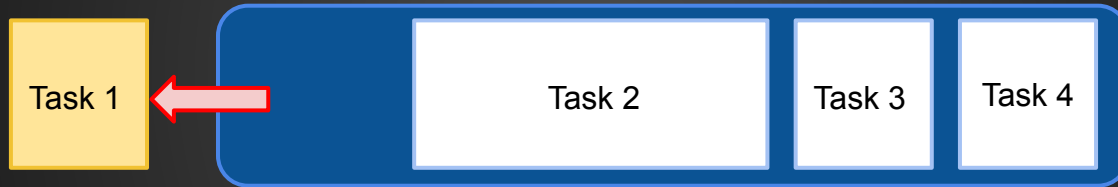
# Pre-Scheduler world



# Pre-Scheduler world

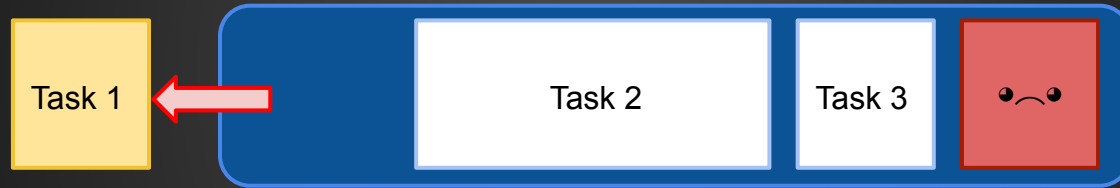


# Pre-Scheduler world



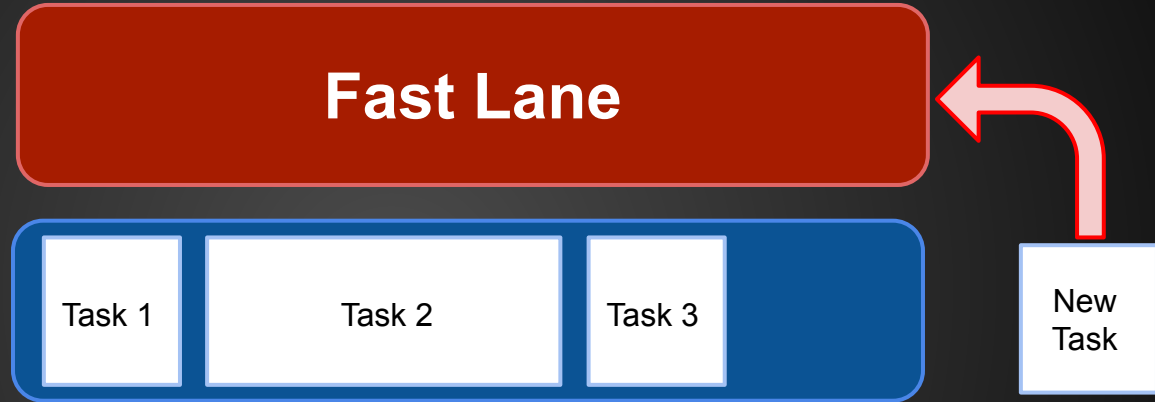
Task 1 taken from the front of the queue and executed.

# Pre-Scheduler world



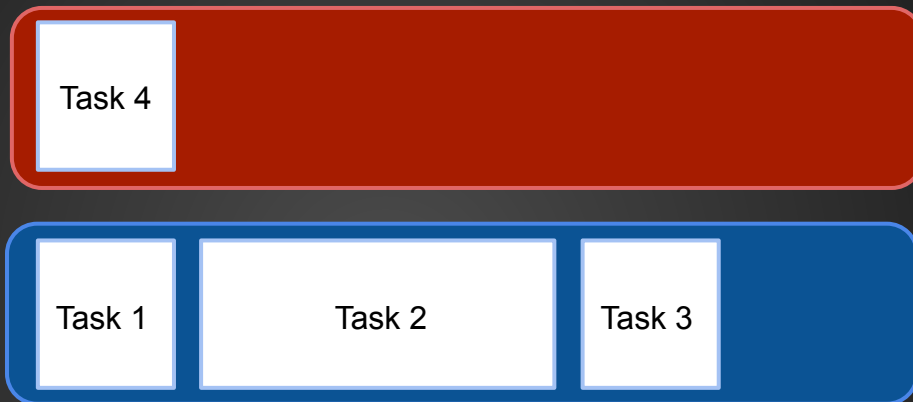
A new task will need to wait until all the preceding tasks have been executed.

# Scheduler world

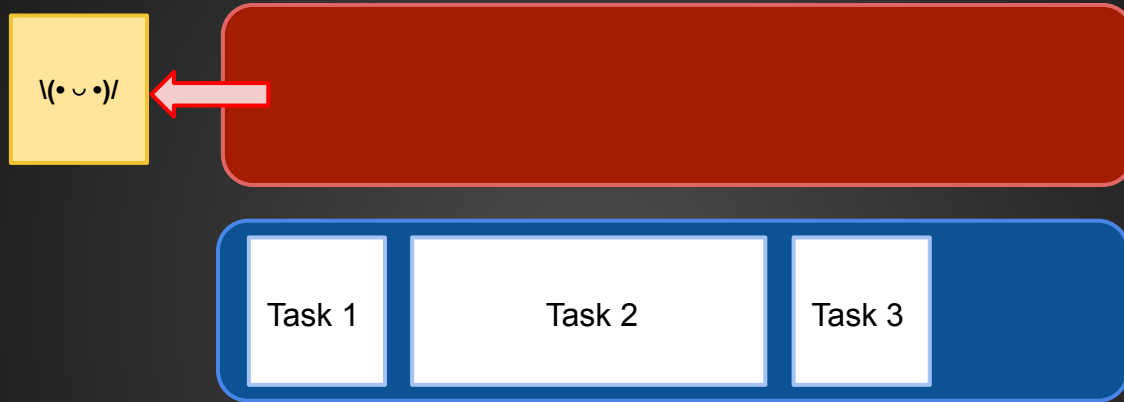


Ideally we'd like to be able to fast-track high priority tasks.

# Scheduler world

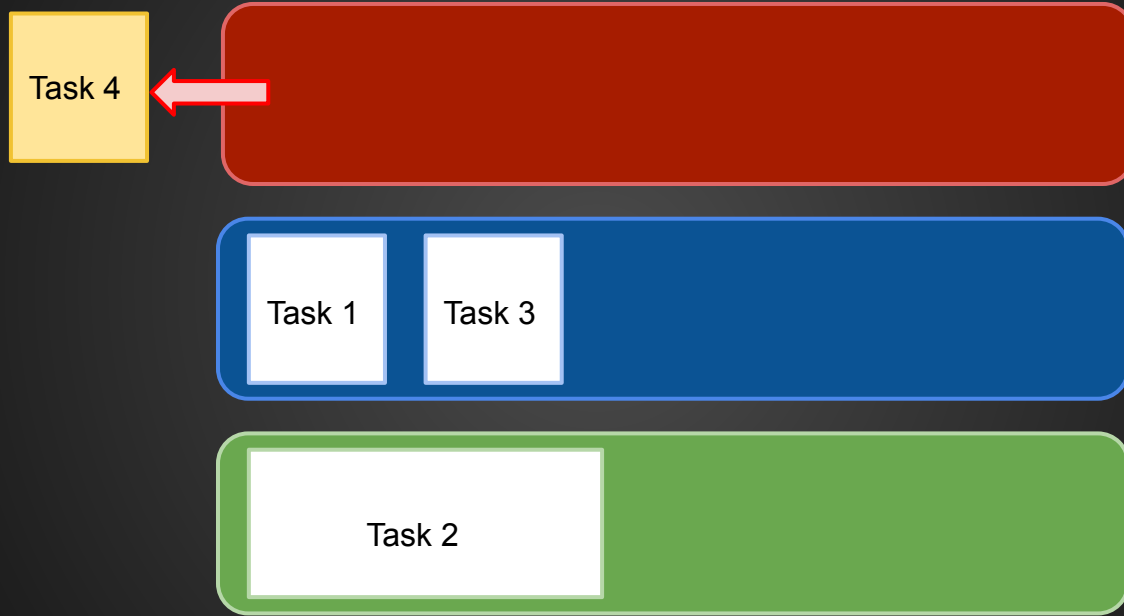


# Scheduler world





# Multiple queues give more flexibility



# Issue: Maintaining execution order

Most tasks need to have their relative execution order maintained.

- Can't composite the results of input processing before the input has been processed
- Ordering assumptions baked into existing code

Only reorder tasks that have opted into prioritization.

# Solution: Tasks are grouped by type

- Maintains relative execution order
- Provides a clear API for posting tasks

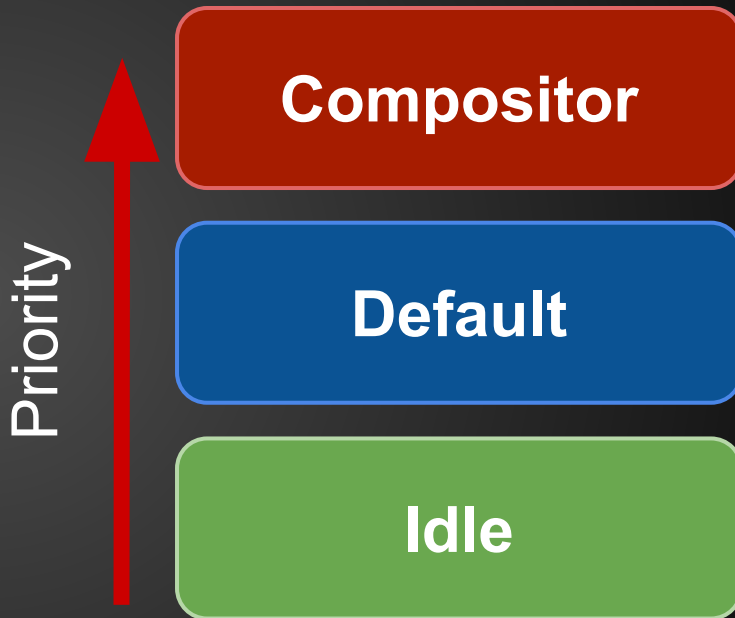
**Composer**

**Default**

**Idle**

# Solution: Tasks are grouped by type

Always prioritise  
composer?



# Issue: Static prioritisation doesn't work

Always prioritising  
compositor tasks  
regressed page loading  
time by 14%.

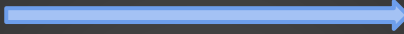
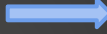
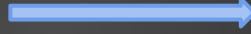
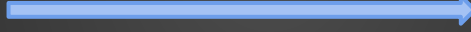
# Solution: Dynamic prioritisation

Use contextual awareness to determine prioritisation *policy*.

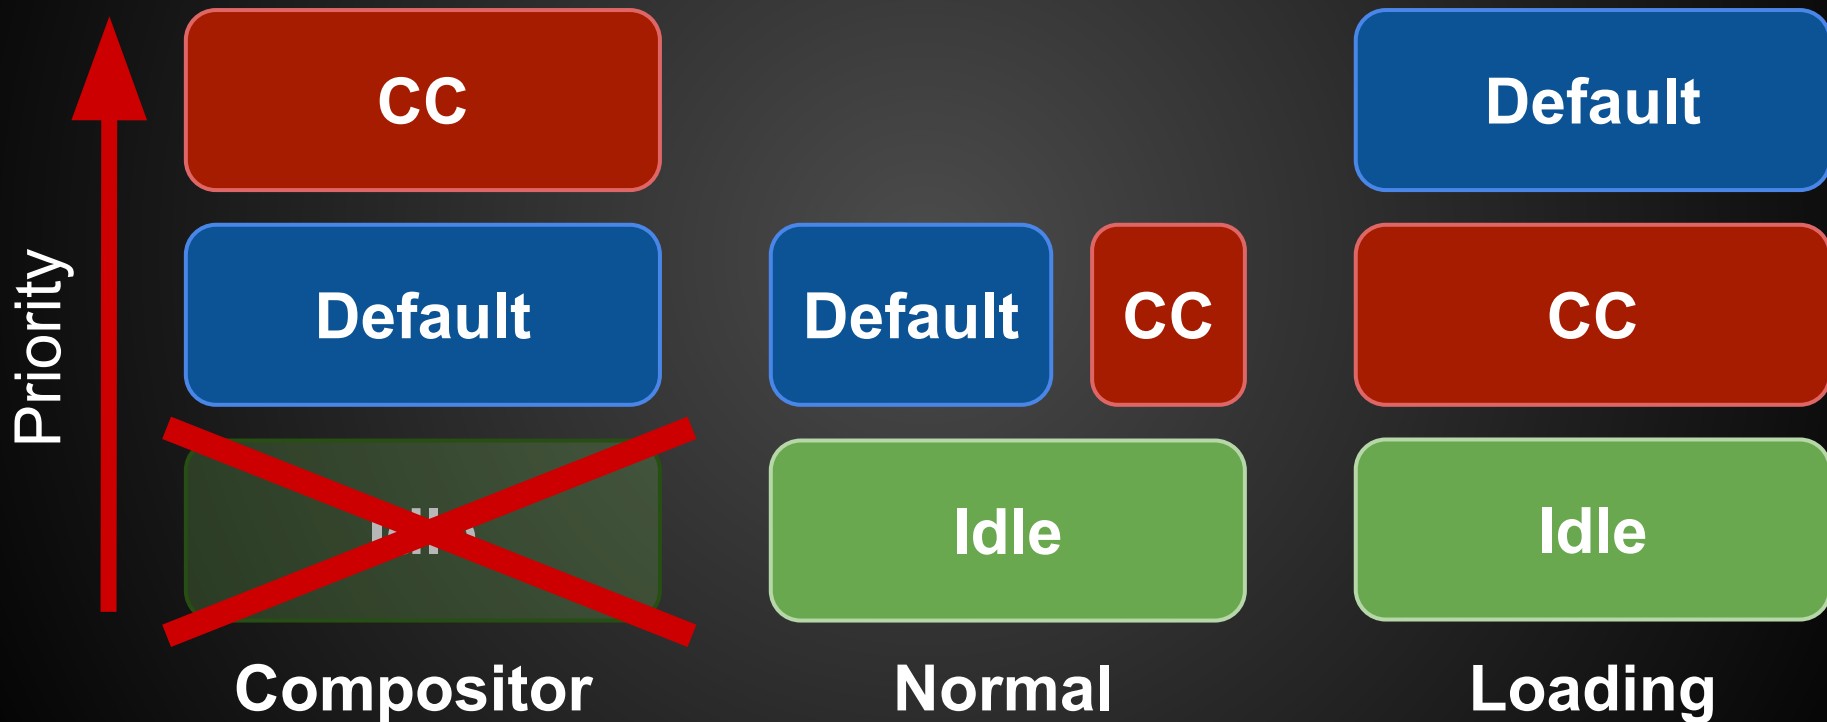
- User interaction makes:
  - **Compositor tasks** *more* important
  - **Resource loading** *less* important
- When finished drawing a frame, allow background tasks:
  - pro-active **GC**
  - background **HTML parsing**

# Policy selection

Events will change the active policy:

- User input events  Prioritize Compositor
- Compositor committed frame  Idle
- Compositor begin frame  Normal
- Page navigation  Page Load

# Queue priorities set by policy





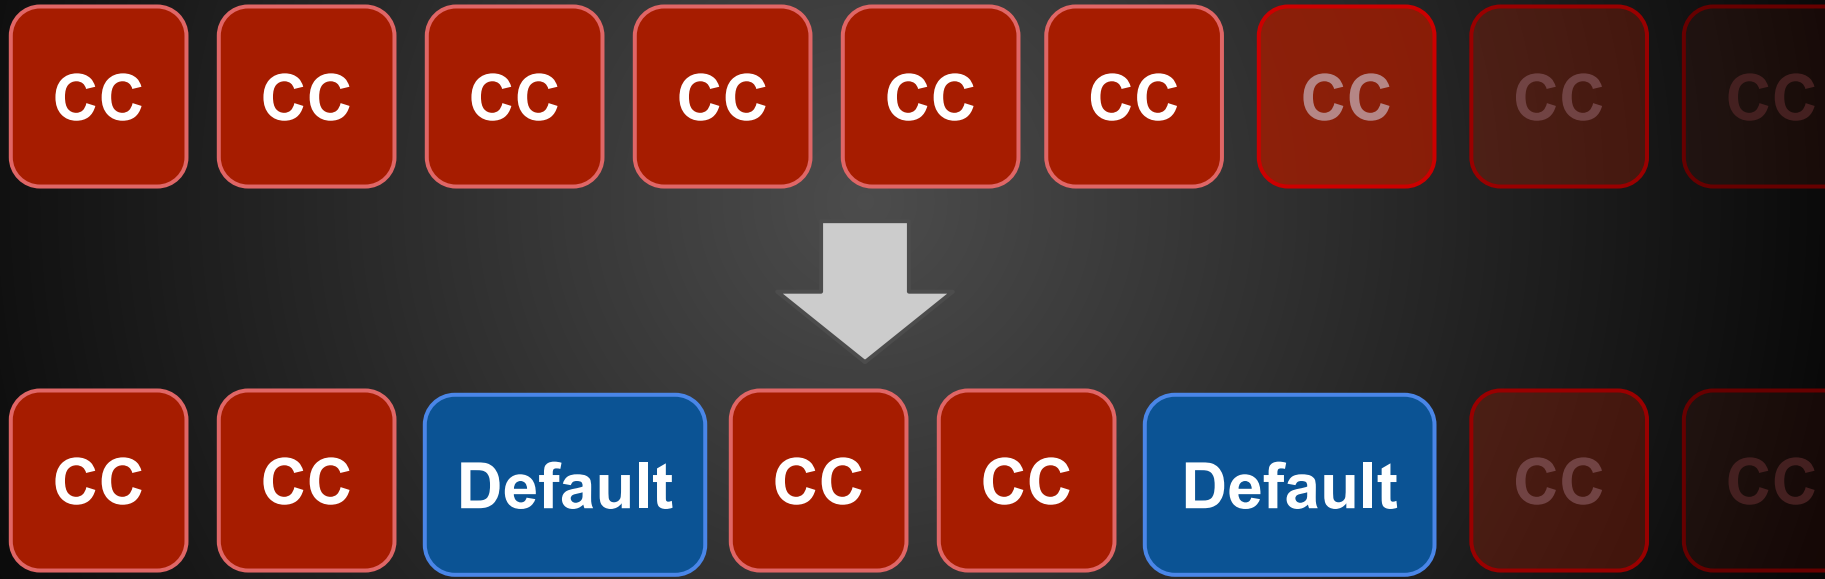
# Issue: Task starvation

We need a scheduler which cannot be 'gamed' by sociopathic tasks to completely starve out other tasks.

# Issue: Task starvation



# Solution: Weighted round robin to avoid starvation



# Architecture

# Finding a home for the Scheduler

**Version 1** implementation was in **Blink**

- Re-implemented much of the message loop
- Posting tasks required a lot of Blink-Chromium plumbing
- Task execution required Russian doll closures

# Finding a home for the Scheduler

**Version 2** implementation was in `base/message_loop`

- Required low level changes in `base`
- Changing mission critical but cruffy code
- Changes would be inherited by all message loops

# Finding a home for the Scheduler

**Version 3** (current) lives in `content/renderer/scheduler/`

- Sits on top of `base::MessageLoop`
- Lets us use more powerful primitives from base
- Avoids crossing layers unnecessarily

# Scheduler architecture overview

- Scheduler provides N task queues
- Policies determine which task queues are prioritised
- System behaviour determines policies

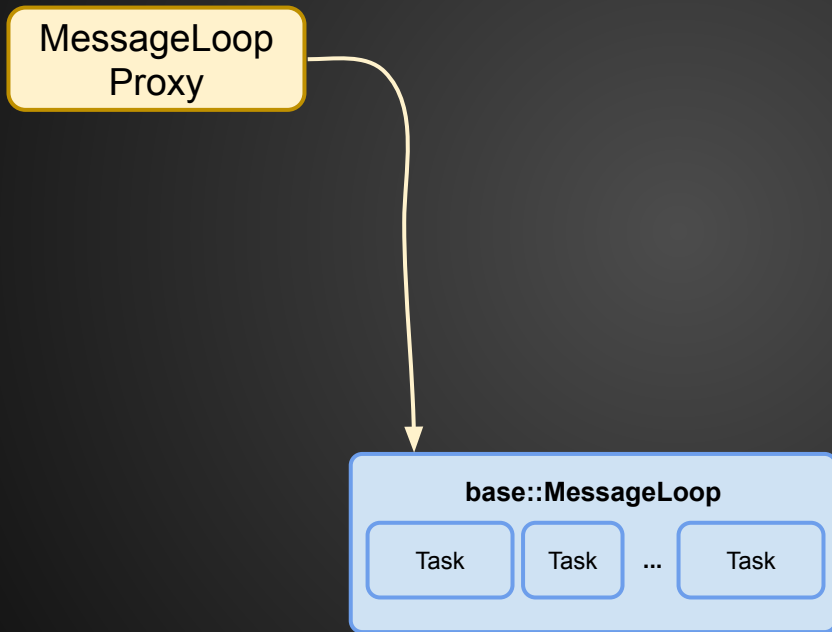


# Scheduler architecture overview

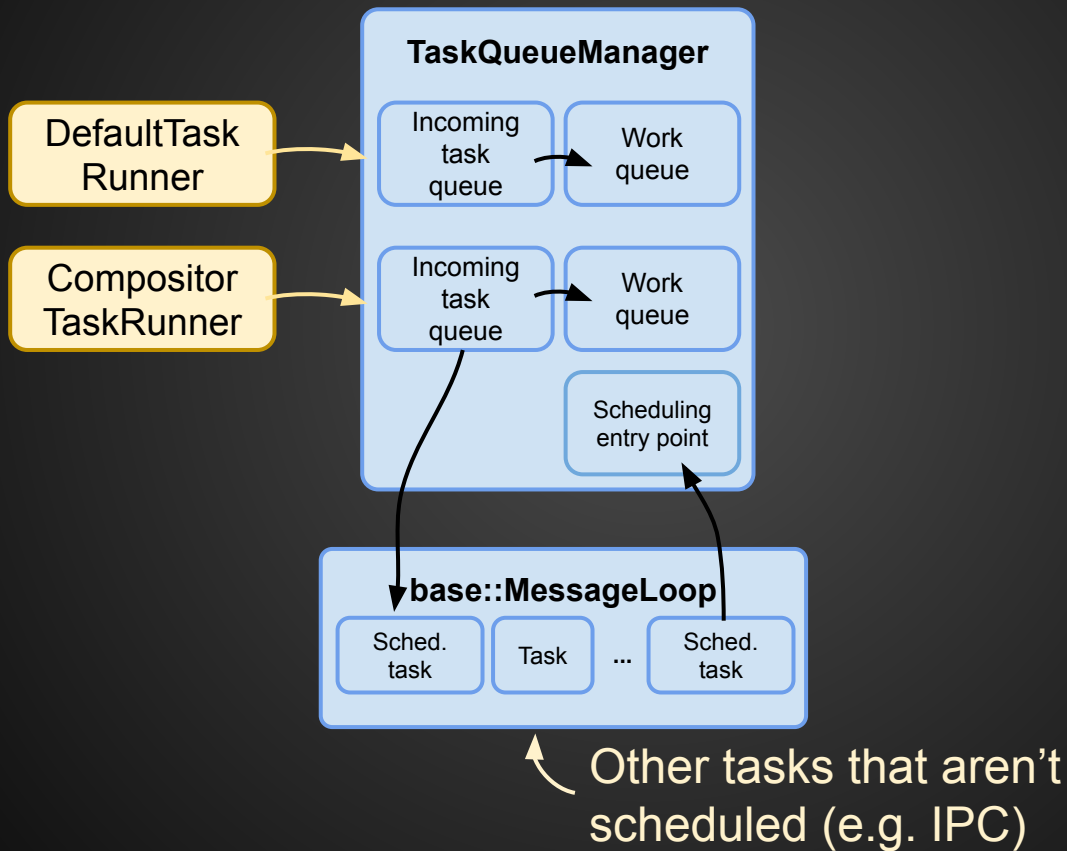
## Three key elements

1. API for posting to different queues
2. Interface that selects the next queue to process
3. Mechanism for changing policies

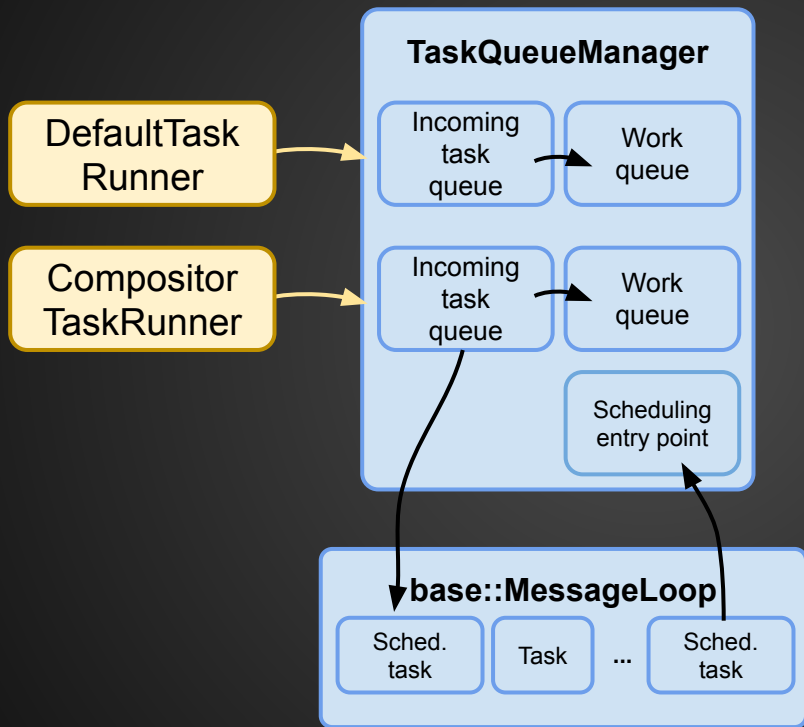
# Architecture



# Architecture: Posting API



# Architecture: Posting API



Chromium:

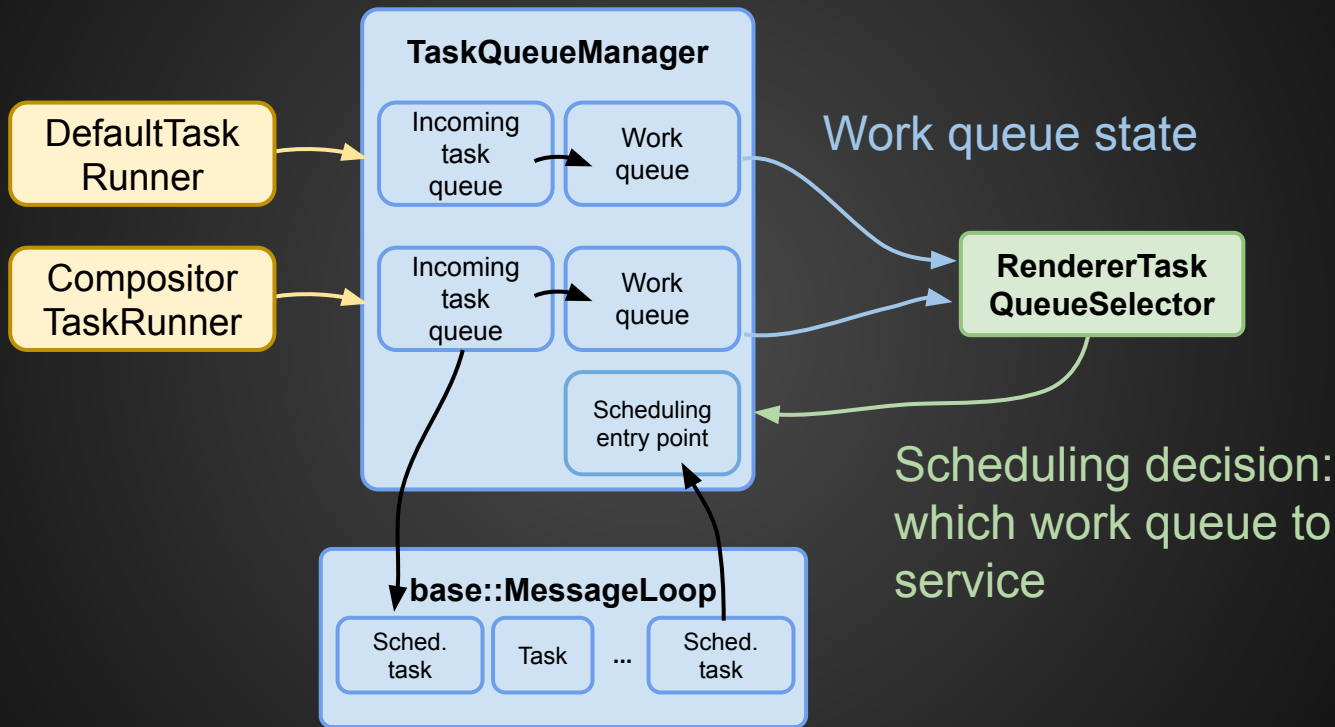
`DefaultTaskRunner()->PostTask()`  
`CompositorTaskRunner()->PostTask()`  
`IdleTaskRunner()->PostIdleTask()`

Blink:

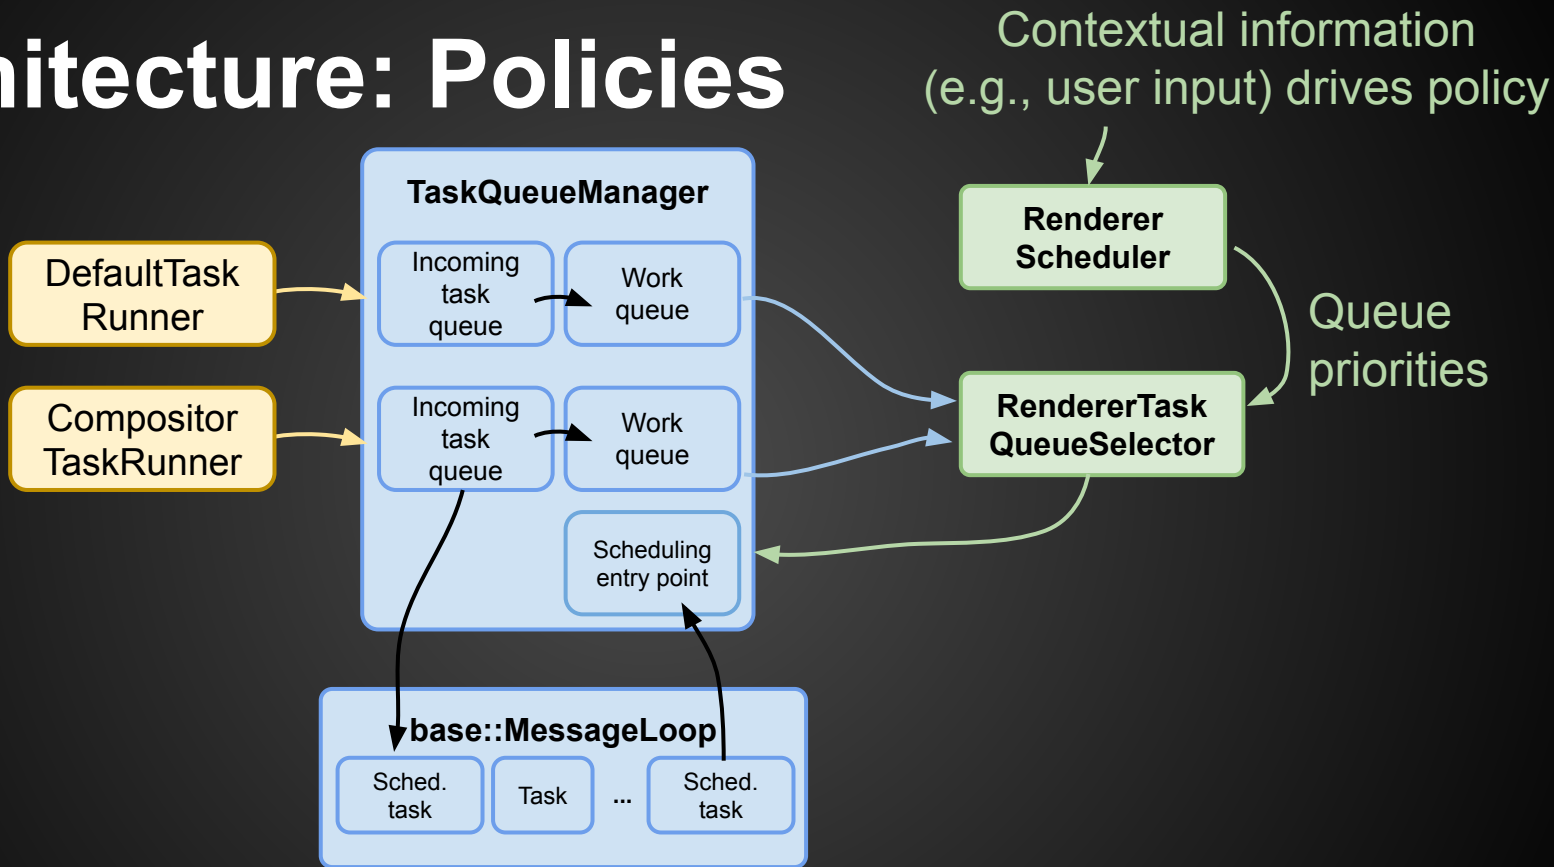
`Scheduler::shared()->scheduler()`  
`Platform::current()->callOnMain()`  
`WebThread::postTask()`

Other tasks that aren't  
scheduled (e.g. IPC)

# Architecture: Queue selector



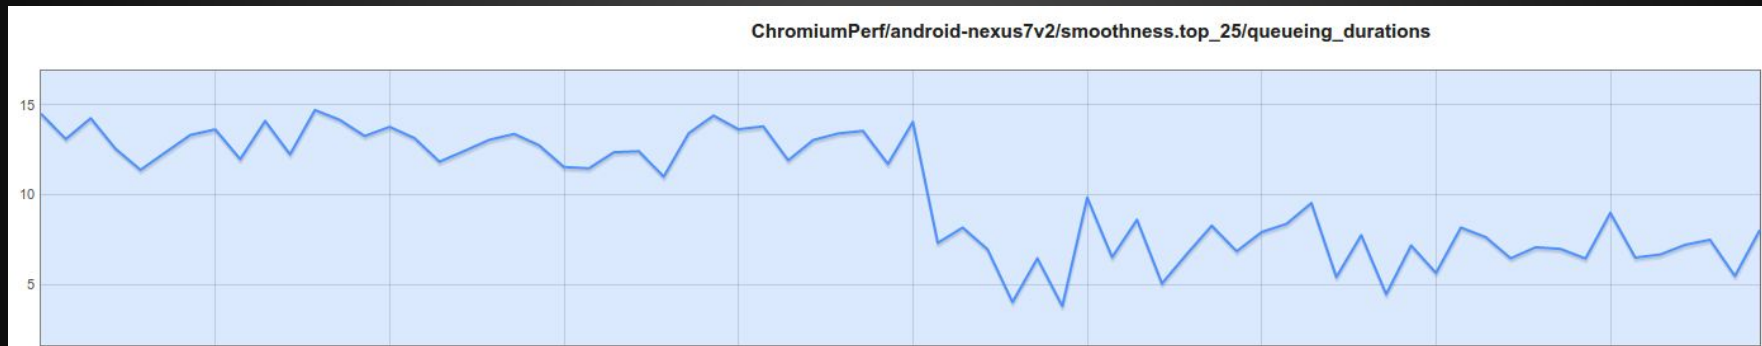
# Architecture: Policies



# Results

# Results from Scheduler V1

(Nexus 7v2)

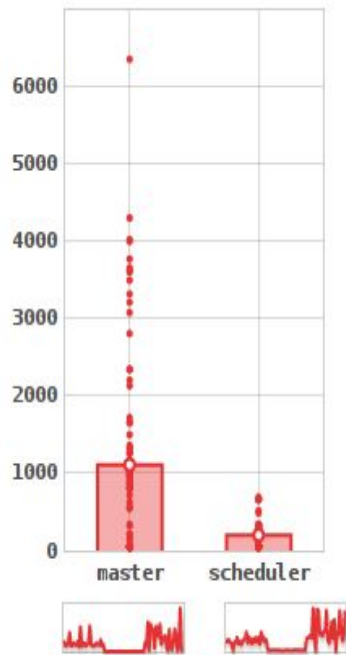


Queueing\_durations ~12ms => ~7ms

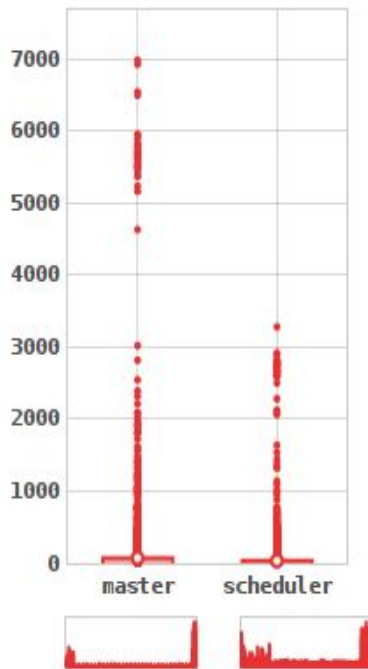


# Preliminary results for V3 (Nexus 7v2)

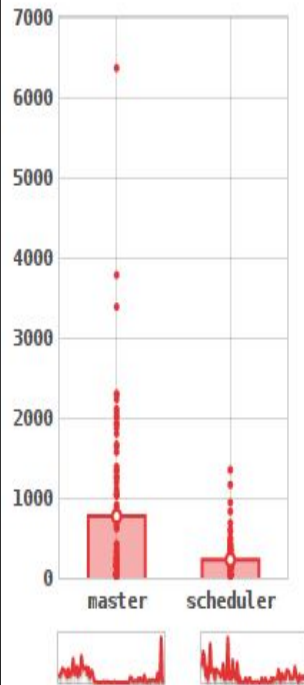
▼ smoothness.top\_25:mean\_input\_event\_latency



▼ smoothness.top\_25:queueing\_durations



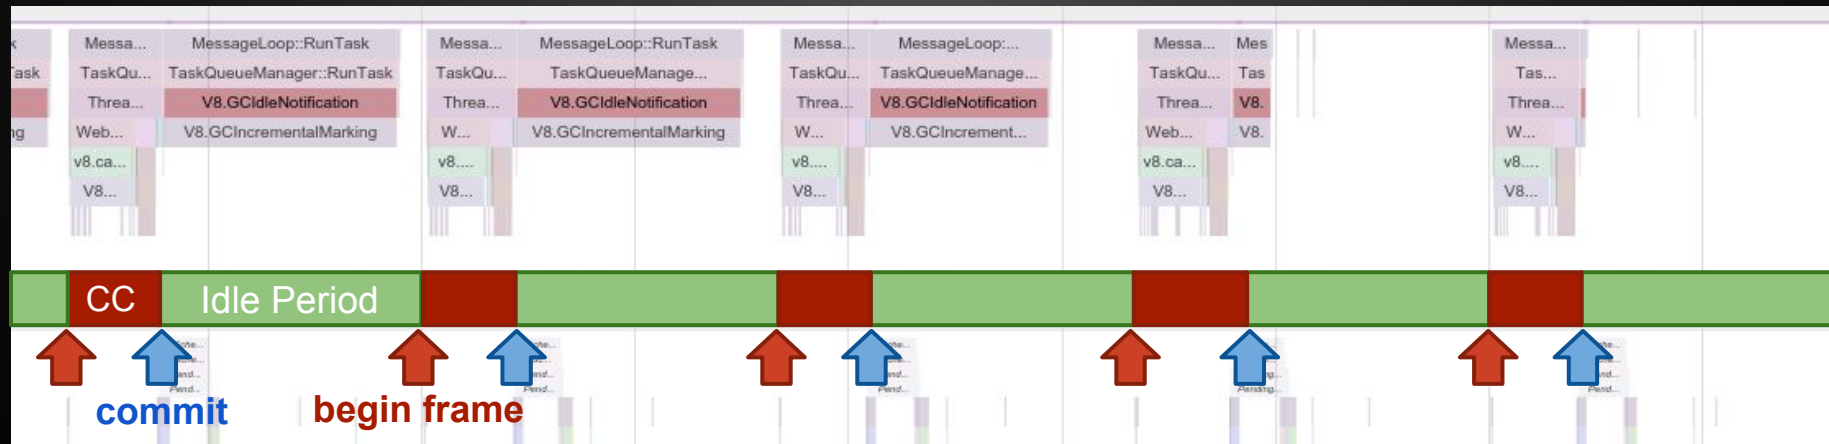
▼ smoothness.top\_25:first\_gesture\_scroll\_update\_latency



**Future steps**

# Future steps: Idle tasks

- Get non-critical work off the critical path
- Only executed when frame has been committed



# Idle task characteristics

- May be starved for arbitrarily long
- Take a deadline argument
  - Should finish by the deadline or earlier
- Posted tasks won't run until next idle period

# First idle task customer: V8 GC

- Do GC work off critical path
- Idle task performs GC incremental marking step
- Estimates how much heap it can mark before deadline
- Repost task if more to do
- Early result: 50% reduction in GC on the critical path

# Future Idle task customers...?

- V8 GC sweeping
- V8 optimising compiler
- HTML parsing
- Oilpan incremental marking

# Future steps: LoadingTaskQueue

- Adding a queue for page loading tasks
  - Resource loading
  - HTML parsing
  - Javascript compilation
- Add a loading priority policy
  - Prioritises loading tasks to optimise first paint 'above the fold'
- Deprioritise when in compositor priority policy

# Future steps: Half-baked ideas

- Provide scheduling primitives to JS
  - `setTimeout(0)` should mean `setTimeout(0)`
  - Posting idle tasks
- Power optimisation
  - Synchronize idle/background ticking in renderers with other wakeups
- Feedback for developers
  - Point out scheduling problems



# Future steps: Half-baked ideas

- Provide scheduling primitives to JS
  - `setTimeout(0)` should mean `setTimeout(0)`
  - Posting idle tasks
- Power optimisation
  - Synchronize idle/background ticking in renderers with other wakeups
- Feedback for developers
  - Point out scheduling problems

# Lessons learned

- How (not) to write a scheduler
  - Existing code unintentionally relies on implicit ordering
  - It's easy to break low level things in subtle ways
- Failing fast and iterating reaches good solutions
- Live and die by metrics
- Having MTV contacts (eseidel@) was invaluable

# Making your tasks scheduler aware

- Avoid blocking the main thread >8ms
- Check `shouldYieldForHighPriorityWork()` from long running tasks
- Post to the most appropriate queue (or add a new one?)
- Do background work using idle tasks
- Join [scheduler-dev@chromium.org](mailto:scheduler-dev@chromium.org) to get involved.