



# V8 Garbage Collection



Reducing Memory Consumption and Latency

Hannes Payer  
Chrome/V8 Team



# Overview

- Metrics we care about
- How does garbage collection work in V8?
- The three deadly sins of garbage collection
- Idle time garbage collection

## **V8 Garbage Collection**

Ulan Degenbaev  
Jochen Eisinger  
Michael Lippautz  
and me.

## **Blink Scheduler**

Ross McIlroy  
Sami Kyostila  
& others...

## **Compositor**

Manfred Ernst

# Metrics we care about



**Throughput**

60 frames per second

=

16.66ms per frame



**Latency**



**Battery**

⇒

ideally max GC pause < 6ms

while keeping >  **Memory** footprint small

# V8 Garbage Collection

- V8's garbage collector interrupts JavaScript execution regularly to collect unused memory
- Generational Garbage Collector
  - Generational hypothesis: most objects die shortly after their allocation
  - Small young generation (up to 16M)  
Semi-space Scavenger: touches live objects only
  - Large old generation (up to 1.4G)  
Mark-compact collector with incremental marking, concurrent sweeping, and compaction: touches the whole committed memory

# V8 Garbage Collection

- Generational Garbage Collector
  - Young Generation: Semi-space Scavenger
  - Old Generation: Mark&Sweep

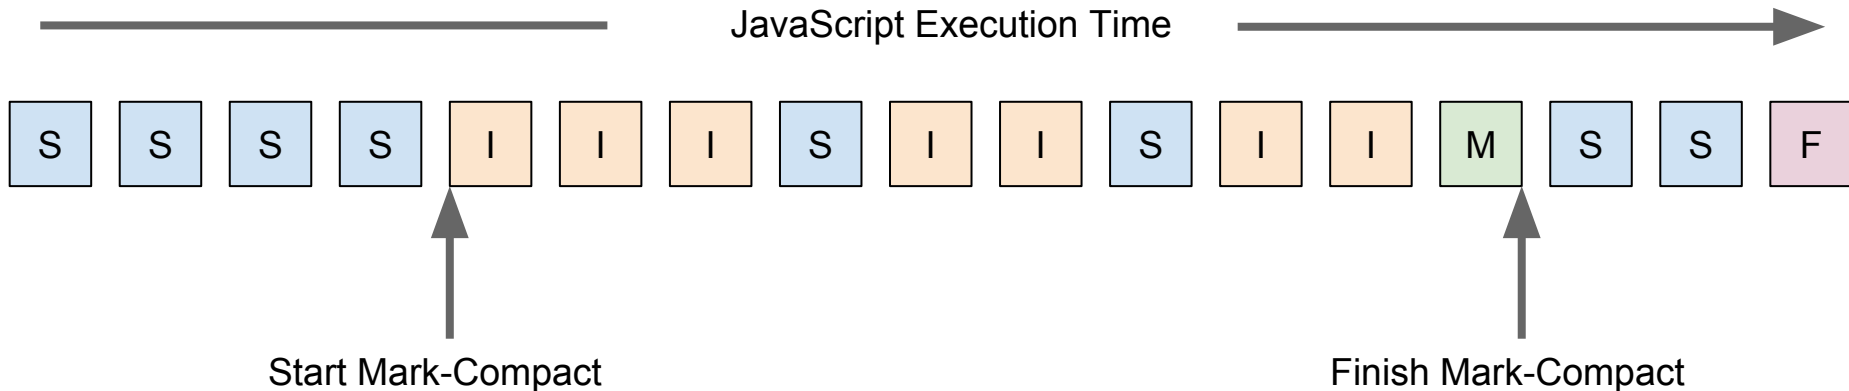


Young Generation

Old Generation

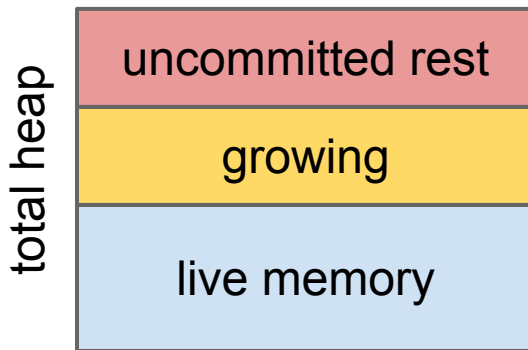
# V8 Garbage Collection

- S** Scavenger (~0-10 ms)
- I** Incremental Marking (~0.01-CONFIGURABLE ms)
- M** Final Mark-Compact Collection (~4-40 ms)
- F** Full Mark-Compact Collection (>40ms)



# V8 Garbage Collection

- Young generation garbage collection trigger
  - when semi-space is full
- Old generation garbage collection trigger
  - based on previous live memory size (a growing factor of 1.1 - 4)
  - start incremental marking close to full state



# The three deadly sins of garbage collection



... or frequent  
feature requests...



# Sin 1: Should I use object pooling to avoid garbage collection?



# Sin 1: Should I use object pooling to avoid garbage collection?

- Why do users want that?
  - “Because it is a known fact that garbage collection is slow, so I try to avoid allocations.”
- Why is it a bad idea?
  - It defeats compiler optimizations (escape analysis, allocation folding, write barrier elimination, generational hypothesis) => it may make your code slower
  - Increases application code complexity (malloc/free style)

## Sin 2: Can I turn off the garbage collector?



## Sin 2: Can I turn off the garbage collector?

- Why do users want that?
  - “Because now there is a time critical phase and garbage collection would result in jank”
- Why is it a bad idea?
  - Difficult to maintain in application code base
  - Complicates the garbage collector (always allocate mode)
  - Out of memory bugs when you forget to turn on or turn on too late

## Sin 3: Can I explicitly invoke the garbage collector?



# Sin 3: Can I explicitly invoke the garbage collector?

- Why do users want that?
  - “Because I know that now would be a good time and later it is super critical that there is no garbage collection going on to avoid jank.”
  - “Because I would like to reduce memory consumption.”
- Why is it a bad idea?
  - Garbage collection heuristics get confused
  - An explicit invocation may take really long and may result in jank

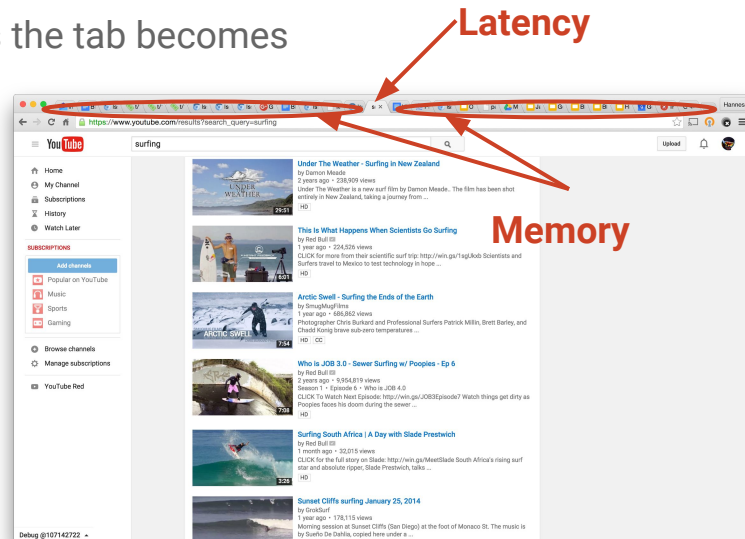
**You are just saying**

**NO NO NO**

**but how can you  
help us???**

# Chrome is funky place for garbage collection...

- Foreground tab
  - Latency matters
  - New frames are drawn every 16.66 ms when animation or scrolling happens
  - Reducing memory becomes important as soon as the tab becomes inactive
- Background tab
  - Latency is secondary
  - But we can reduce memory consumption






# Idea: Make garbage collection invisible



## When is the best time to do a GC?

When nobody is looking.

Using camera to track eye movement  
When subject looks away do a GC.

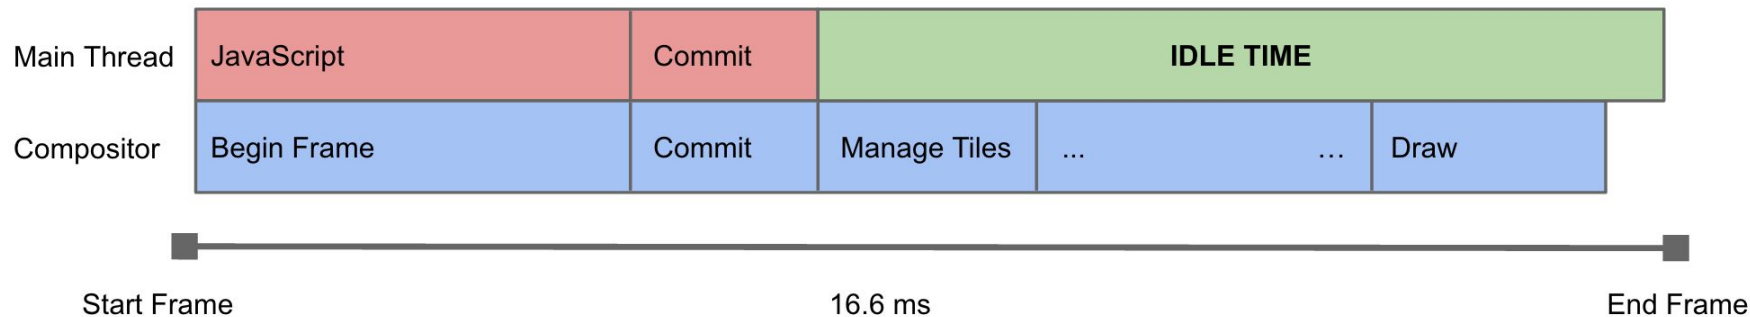


[https://upload.wikimedia.org/wikipedia/commons/3/35/Computer\\_Workstation\\_Variables.jpg](https://upload.wikimedia.org/wikipedia/commons/3/35/Computer_Workstation_Variables.jpg)

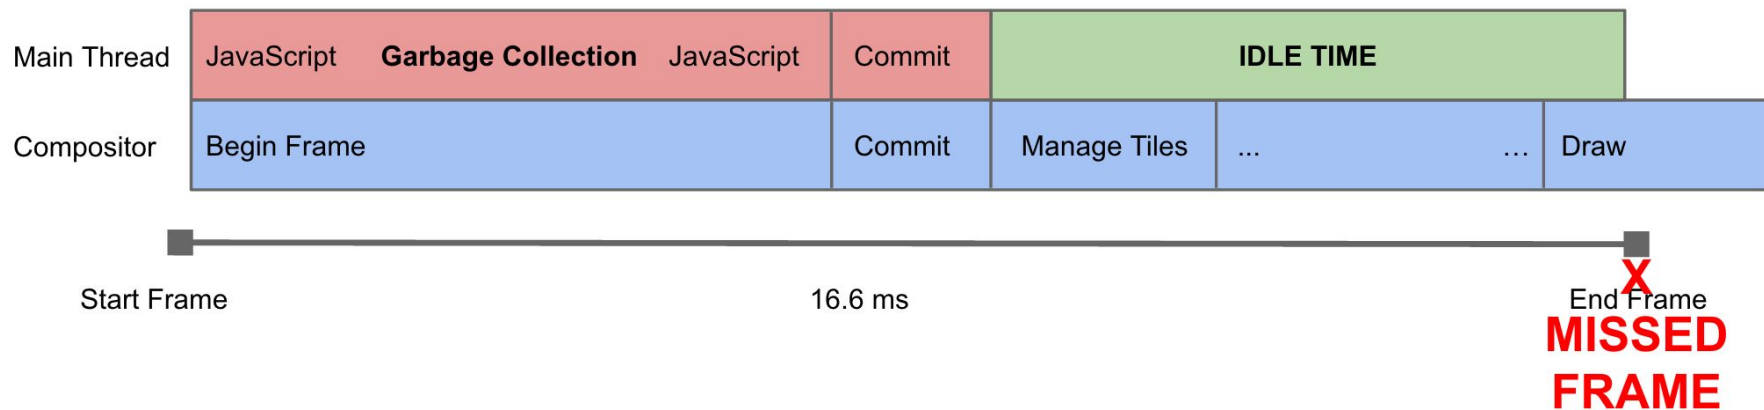
# Idle Time Garbage Collection

# Latency

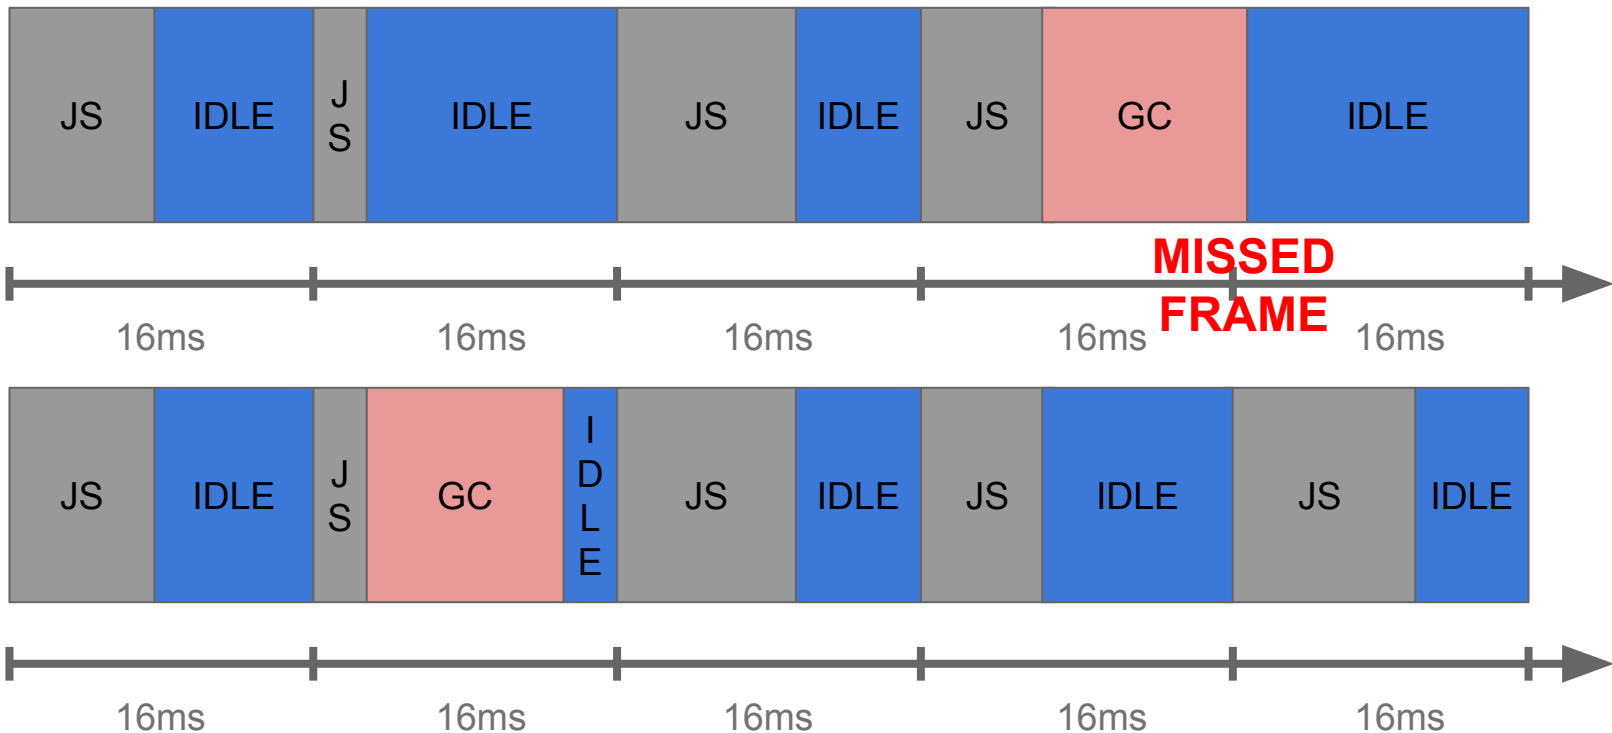
# Life of a Frame



# Life of a Frame




# Life of a frame



# Latency-driven Idle Time GC Scheduling

- We have a bunch of heuristics and monitoring code that will try to estimate:
  - average young generation collection speed/MB
  - average incremental marking speed/MB
  - average finalization of mark-compact speed/MB
- We register an idle garbage collection task in the Blink scheduler when a given garbage collection operation should happen soon.
- The task scheduler will schedule us as soon as there is idle time (and will give us up to 50ms to perform garbage collection)



# History of Idle Garbage Collection Systems

- M38
  - V8 was directly called from the Compositor
    - No global notion of idleness
    - Maximum idle time 16.66ms
    - Constant rendering overhead after each rendered frame
- M41
  - Integrated into Blink task scheduler
    - Whenever Blink task scheduler detected idleness a V8 idle tasks was scheduled -> many unnecessary V8 invocations, because most of the time there was no garbage collection work pending
- M45-M46
  - V8 registers idle tasks in the Blink task scheduler when there is work pending





## M41 vs M46

### M46 --disable-v8-idle-tasks

Scavenge 289.3 (388.6) -> 283.3 (388.6) MB, 18.8 / 0 ms [allocation failure].  
Scavenge 294.2 (388.6) -> 288.1 (388.6) MB, 20.7 / 0 ms [allocation failure].  
Scavenge 299.0 (388.6) -> 291.6 (388.6) MB, 15.4 / 0 ms [allocation failure].  
Scavenge 303.9 (388.6) -> 295.8 (388.6) MB, 10.2 / 0 ms [allocation failure].  
Scavenge 307.4 (388.6) -> 299.6 (388.6) MB, 9.8 / 0 ms [allocation failure].  
Scavenge 311.5 (388.6) -> 304.6 (388.6) MB, 16.5 / 0 ms [allocation failure].  
Scavenge 315.4 (388.6) -> 309.9 (388.6) MB, 19.3 / 0 ms [allocation failure].  
...


### M46

Scavenge 363.1 (401.9) -> 358.6 (402.9) MB, 4.5 / 0 ms [idle task: scavenge].  
Scavenge 362.1 (402.9) -> 358.6 (402.9) MB, 3.3 / 0 ms [idle task: scavenge].  
Scavenge 364.8 (402.9) -> 363.3 (402.9) MB, 5.7 / 0 ms [idle task: scavenge].  
Scavenge 363.5 (402.9) -> 363.3 (406.9) MB, 10.7 / 0 ms [idle task: scavenge].  
Scavenge 366.9 (406.9) -> 363.4 (406.9) MB, 5.3 / 0 ms [idle task: scavenge].  
Scavenge 367.0 (406.9) -> 363.5 (406.9) MB, 5.3 / 0 ms [idle task: scavenge].  
Scavenge 370.4 (406.9) -> 368.2 (406.9) MB, 6.4 / 0 ms [idle task: scavenge].  
...

### Improvements M46 over --disable-v8-idle-tasks

Frames per second:	17.73ms (2.70% better)
Frame time discrepancy:	138.65ms (44.13% better)

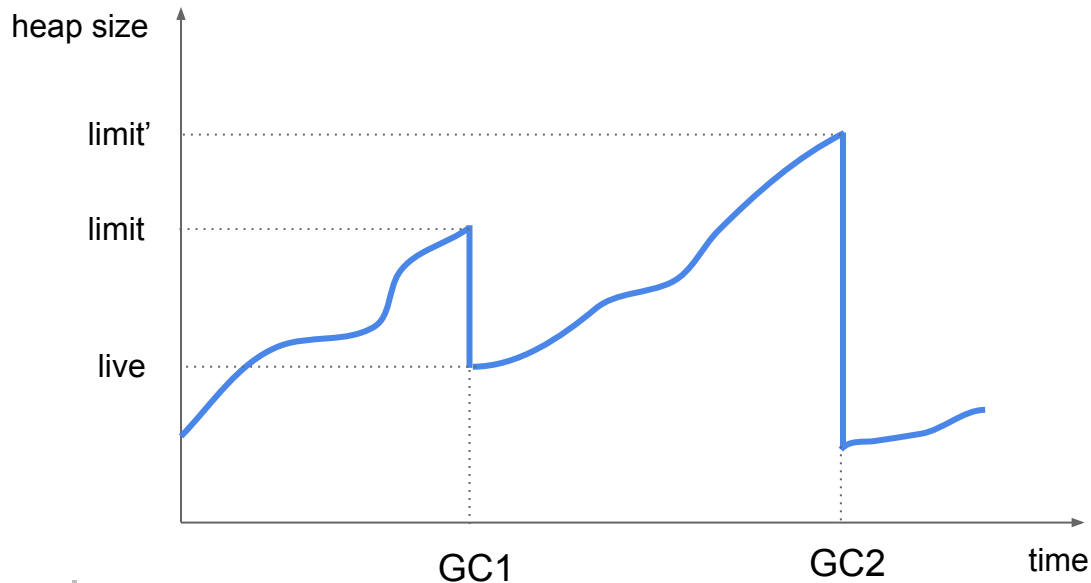
# Telemetry Infinite Scrolling Benchmarks



# Memory

# Heap Growing Strategy

A new garbage collection is triggered (incremental marking is started) as soon as we reach a certain heap size (limit).




$$limit' = live\ memory \times [1.1, 4]$$

Assumption: The allocation rate is constant. Hence  $limit'$  is reached soon which triggers garbage collection.

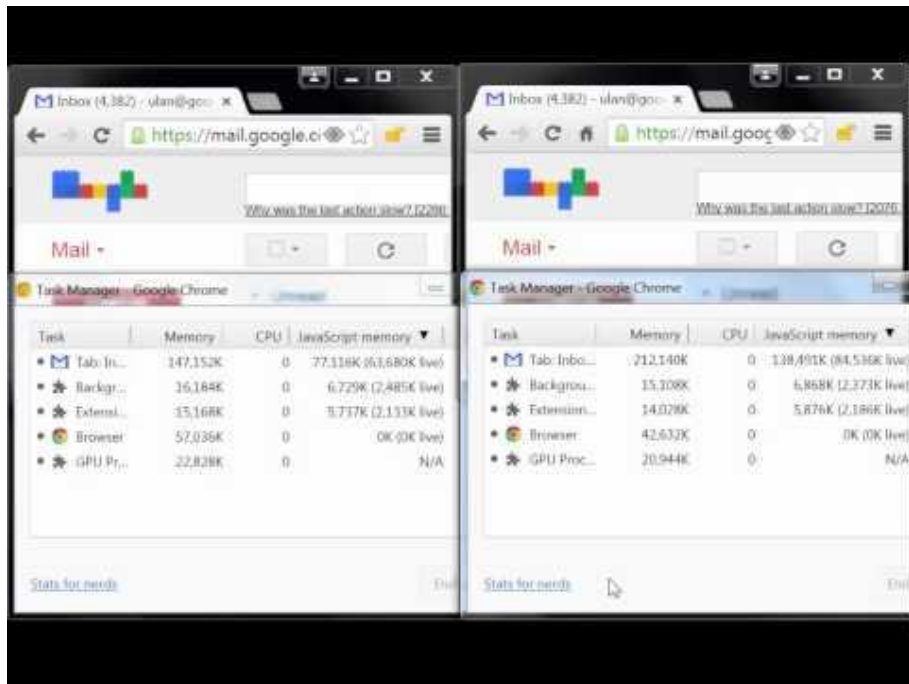
But: What if we had a high allocation rate, set a high  $limit'$ , and suddenly the website became idle, i.e. allocation rate became low (or zero)?

# Heap Growing Strategy



We will never reach limit' and may unnecessarily hold on to a lot of memory.

# Memory-reducing Idle GC Scheduling



Gmail M45 vs M43

- **Memory Reducer**  
Track recent allocation rate and check back periodically if a tab became idle
- If idle:
  - start incremental marking,
  - schedule incremental marking steps,
  - perform more aggressive compaction,
  - shrink and uncommit young generation
  - set a small heap growing limit.
- Don't drain the battery!
- ~40% less memory on average

# Challenge: Benchmarks

We like long-running (>30 seconds) benchmarks that perform *interesting* actions on real websites.

How do we get benchmark that represent the web of tomorrow?



# Thank You! Questions?

+Hannes Payer

hpayer@google.com

<http://research.google.com/pubs/HannesPayer.html>

