

Jumbo

A unity build system for Chromium

Daniel Bratell

TL;DR

Jumbo: 3 times faster* builds now
and possibly 9 times faster builds
in the future.


The problem

The compilation time for Chromium is long

Very long

Extremely long

Hours.



xkcd 303

<https://www.xkcd.com/303/>


(by Randall Munroe)

Compile time evolution

In five years

34 min → 4 hours


+1% per week




Will newer hardware solve it?

(Betteridge's law of headlines)

Hardware does
not keep up



Time vs repo size evolution



Time consumers

- Code generation
- Compiling
- Linking

Can use the ninja log (`out/something/.ninja_log`)

No file needs more than 0.04% of the total time but:

There are 44,000 .cc files and 43,000 .cpp files in Chromium (including generated code).

Files are small

Median length `.cc` files is 134 lines


Median length `.cpp` files is 78 lines [pre-Blink move].

Preprocess a Blink file of length **110 lines** and you get **244,000 lines of source code**.

Per file

In lines the
source file is
nothing!

In compile
time almost
nothing



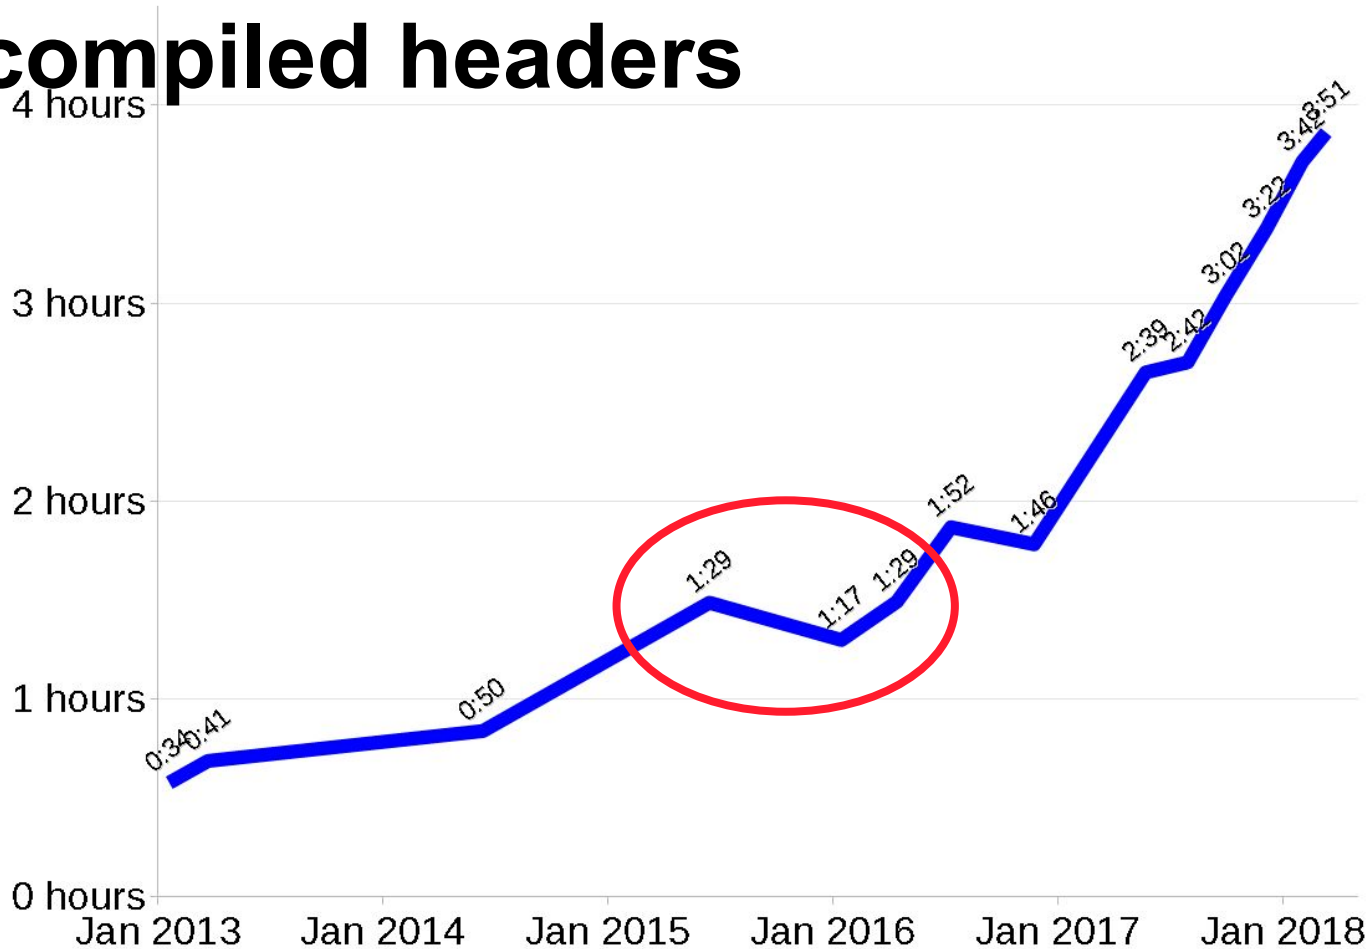
Precompiled headers

Been used for system headers for a long time

2015: Blink got more comprehensive precompiled headers

Saves 10-20% of the compile time


Effect precompiled headers



Unity builds

In Chromium: Jumbo builds

- Compile a lot of code in a single translation unit
- Common in large projects and the games industry
- Long used for blink's v8 bindings, and by third_party/sqlite
- Used in WebKit since February



Jumbo

Is it an elephant?

Is it a dutch
supermarket?

Is it a unity build
system?



Google

jumbo

All Images Maps Videos News More Settings Tools

Any time ▾ All results ▾ Clear

Jumbo - Wikipedia
<https://en.wikipedia.org/wiki/Jumbo> ▾
Jumbo (about Christmas 1860 – September 15, 1885), also known as **Jumbo the Elephant** and **Jumbo the Circus Elephant**, was a 19th-century male African bush elephant born in Sudan. **Jumbo** was exported to Jardin des Plantes, a zoo in Paris and then transferred in 1865 to London Zoo in England. Despite public protest ...
Years active: 1862–1885 in captivity **Cause of death:** Railway accident
Species: [African bush elephant](#) **Weight:** 6.15 tonnes (13,558 lb)
[History](#) · [Death](#) · [Legacy](#) · [References](#)

Jumbo: Altijd lage prijzen bij dé supermarkt van Nederland
<https://www.jumbo.com/> ▾ [Translate this page](#)
Ontdek alle aanbiedingen en het ruime assortiment van **Jumbo!** ✓ Laagste prijsgarantie ✓
Boodschappen thuisbezorgen ✓ Heerlijke recepten ✓ Handige **Jumbo** app.
[Aanbiedingen](#) · [Producten](#) · [Geen bestelkosten](#) · [Seizoensaanbiedingen](#)

Homepage - Games & Puzzles - Jumbo
www.jumbo.eu/ ▾
The most fantastic games and puzzles for all ages including Peppa Pig, Little Kingdom, Fireman Sam, Falcon, Jan van Haasteren and Wasgij.


Jumb
Elephant

Jumbo, also
and Jumbo t
19th-century
born in Suda
Jardin des P
transferred i
England. Wi

Born: 1861,
Died: Septe
[Canada](#)
Species: Lo
Cause of de

Jumbo in BUILD.gn

```
source_set("my_code") {  
    sources = ["file1.cc", "file2.cc", "file3.cc"]  
}
```



```
import("//build/config/jumbo.gni")  
jumbo_source_set("my_code") {  
    sources = [ "file1.cc", "file2.cc", "file3.cc"];  
}
```

Jumbo in the file system

Template action generates `gen/.../my_code_jumbo_1.cc`:

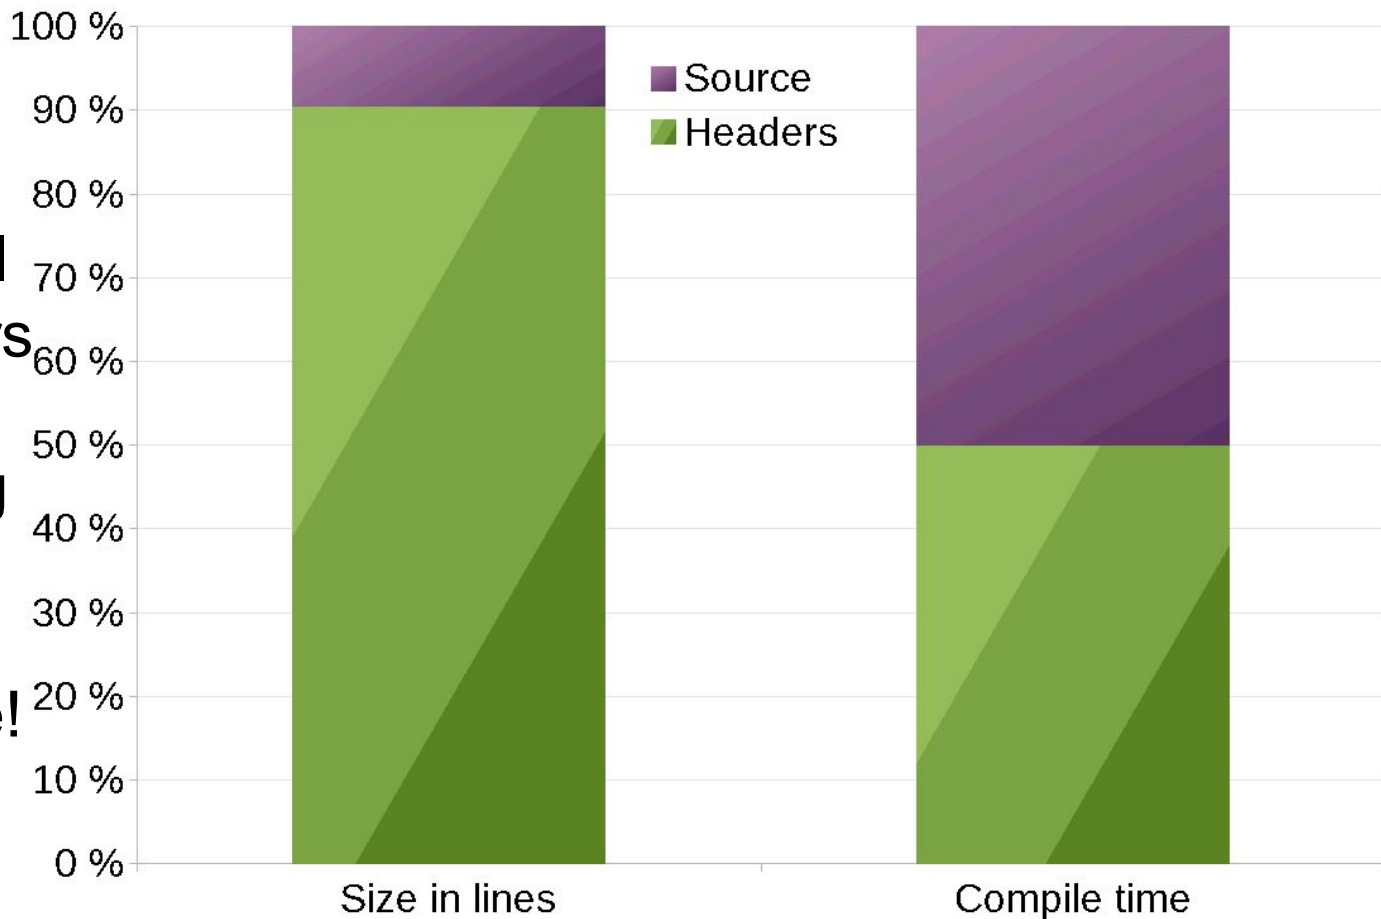
```
#include "../..../chrome/my_code/file1.cc"  
#include "../..../chrome/my_code/file2.cc"  
#include "../..../chrome/my_code/file3.cc"  
#include "../..../chrome/my_code/otherfile.cc"
```

`my_code_jumbo_1.cc` compiled as usual

Jumbo

90% of the
lines are still
from headers

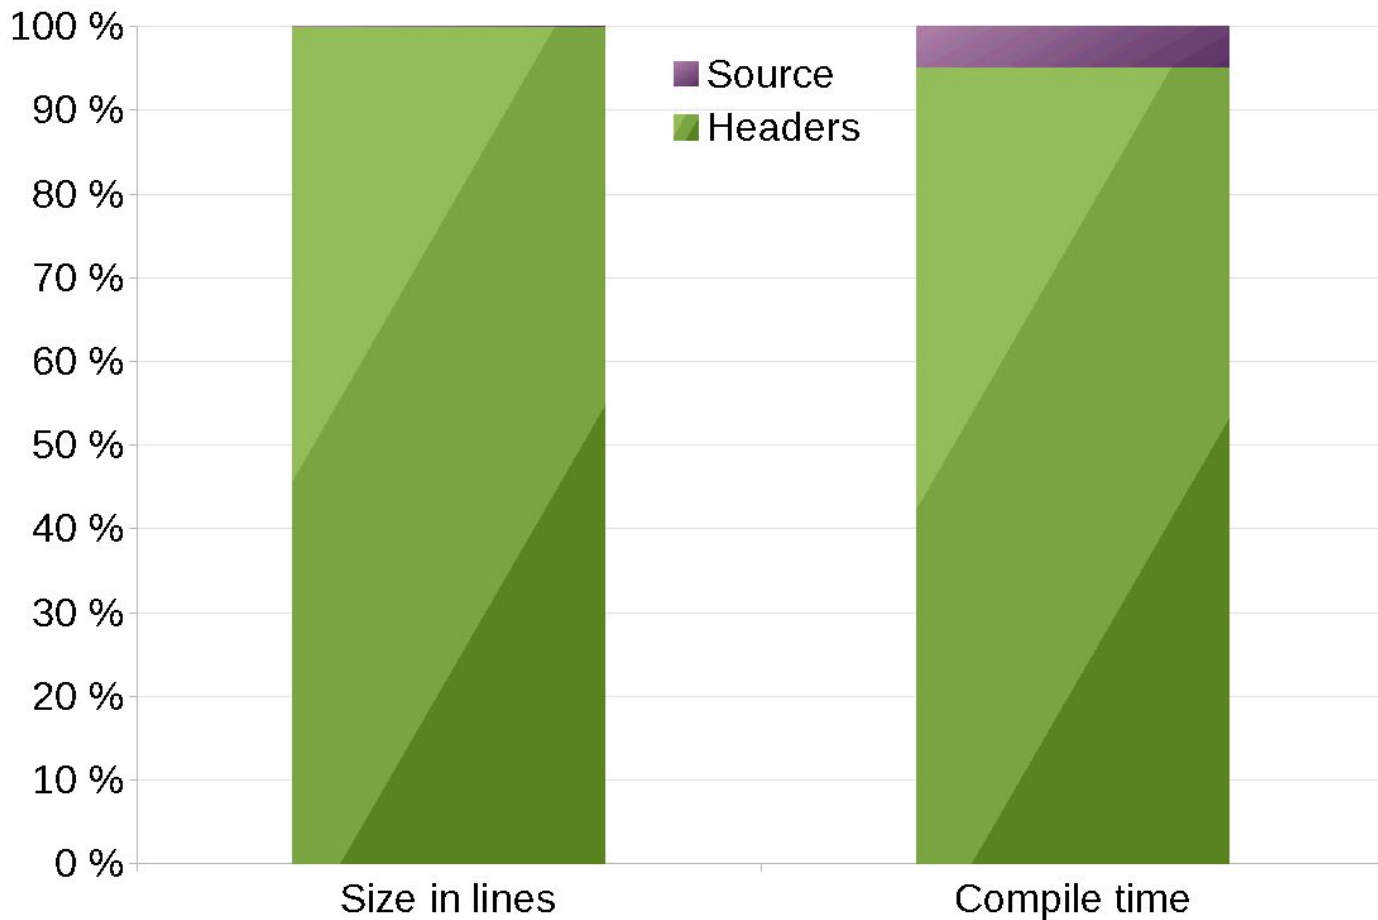
... compiling
the source
code is now
half the time!



Per file

In lines the
source file is
nothing!


In compile
time almost
nothing



Jumbo

90% of the
lines is still
from headers

... compiling
the source
code is now
half the time!



Metaphor

If you have a lot of cargo, some tools are more efficient



Metaphor


If you have a lot of cargo, some tools are more efficient



Jumbo

Compiling:
chrome+
content_shell+
blink_tests


4 core/8 thread



Jumbo

Compiling:
chrome+
content_shell+
blink_tests

On a 4 core/8
thread computer



Not always faster

- Always more efficient
 - Saves roughly 150 Wh per full build
- Less parallel
 - Concern if >100-1000 cores
 - Longer executing bottleneck tasks
 - Idling waiting for dependencies
 - Not an issue for "normal" hardware
- Tunable
 - Currently tuned for ~10-20 cores or ~100 cores if goma

Metaphor

If you have a lot of cargo, some tools are more efficient



Not always faster

- Always more efficient
 - Saves roughly 150 Wh per full build
- Less parallel
 - Concern if >100-1000 cores
 - Longer executing bottleneck tasks
 - Idling waiting for dependencies
 - Not an issue for "normal" hardware
- Tunable
 - Currently tuned for ~10-20 cores or ~100 cores if goma

Good things

Main target

- Average compilation time per file reduced by 80-95%

Accidental positive effects

- Less compiler output (debug component build tree 26 GB → 14 GB)
- Faster linking (due to less data to process)
- Cheap "global" optimization (per jumbo unit)

More positive side effects

- Duplicate code removal
- Dead code removal + dead member removal
- Addition of include guards
- Solved the X11 header problem (no more `#undef None`)
- Namespace renamings, ~~no more~~ fewer `::`-prefixes
- Easier to find symbols when their names are more distinct
than `content:: {anonymous namespace}::kValue`

Complications

- With jumbo: "1 cc file" \neq "1 translation unit"
 - local names become less local.

To a lesser degree:

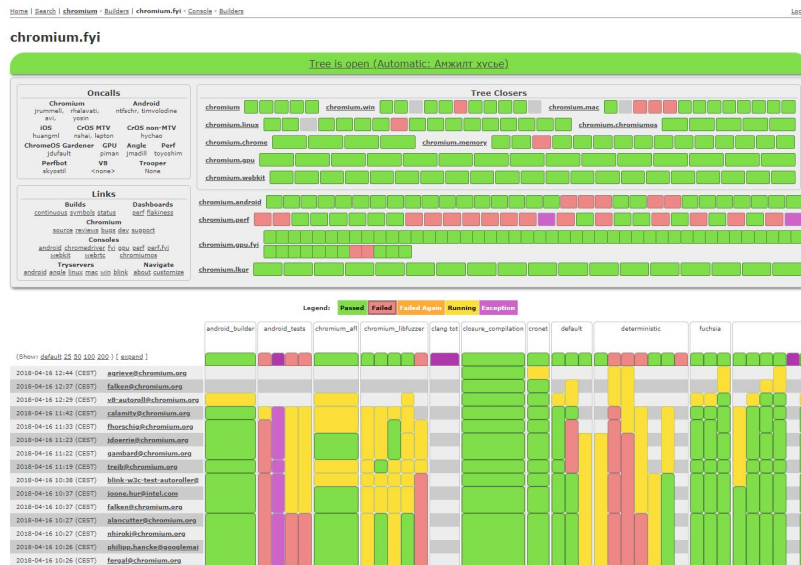
- More code is exposed to system headers
- Stress testing tools (~~"object copy $O(n^2)$ in number of sections"~~)
- Worse IDE support in gn
- Increased single file rebuild times (unless linking makes up for it)

Infrastructure

3 fyi bots

Mail a selected few when one of them break (sometimes twice a day, sometimes once in a week)

CQ support coming any day/week



To do

CQ support

Star crbug.com/782863 for updates (unless it's already fixed)

Add support to the rest of Chromium (after CQ)

Tests, components, services, third_party

Native gn support

For better IDE support

PoC/patch by Tomasz Moniuszko exists

Star crbug.com/772918 for updates



Assistance by clang (PoC covered in lightning talk)

Create sandboxes for each file and prevent some of the problems

PoC: Add pragma that hides/disables current anonymous namespaces. Can not be developed purely as a plugin.

Doesn't solve all problems but prevents the more annoying ones

Time frame: Uncertain, months to infinity



Deprecating other systems

- `split_static_library`
 - Needed because libs > 2 GB. Not the case in jumbo builds
- component builds
 - Solving long link times. Linking is (a bit) faster in jumbo
- Putting a lot of source in one file
 - v8 has not split up its code because of compile times. Faster in jumbo?

Jumbo alternatives

Distributed compilation (icecc, distcc, an open and independent goma)

Can sometimes be combined with jumbo for win-win

Faster compilers

clang can be 20% faster if compiled with other flags
compilers that cache state

C++ modules?

Can be combined with jumbo for win-win

Another implementation language than C++?



Open source goma instead of clang

goma client open source while server still proprietary

If depends on Google's infrastructure

- Intellectual property

- Price

- Dependence

- "Privacy"

Else

Access to hardware (works for large companies but not individuals and small companies)



- Maintenance

How to make it faster?

Less code

No

Precompiled headers

Saves only 10-20%

Faster compilers

No silver bullet, maybe 20%

Faster hardware

Not fast enough

Distributed compilation

Really helpful **if** ...

Massive hardware available

Distribution systems work

Can save 90+% of the time

Unity builds

Requires code changes

Can save 90% of the time

Questions

"If you are wondering something, there are probably many others that would like to know the answer as well."

/ Me - right now

Thanks!

Daniel Bratell @ Opera

Mostyn Bramley-Moore @ Vewd

Bruce Dawson @ Google

(Windows/goma)

Dirk Pranke @ Google

(infrastructure)

Tomasz Moniuszko @ Opera

(native gn, IDE support)

Jens Widell @ Opera

(clang support)

And many, many others (haraken, thakis, the_stig, pdr, fs, avi, kinuko, sky, sadrul, ...) who have tested, experimented, reviewed and added jumbo support to code

