



Perf Insights, RAIL & Happiness

<https://goo.gl/n2l4Fd>

10 Nov 2015

nduca@chromium.org





rail score plz



RESPONSE

100ms



ANIMATION

10ms



IDLE

50ms



LOAD

1000ms

how do we measure RAIL?




"Just"....


On CrRenderer main, track which stage we're in...

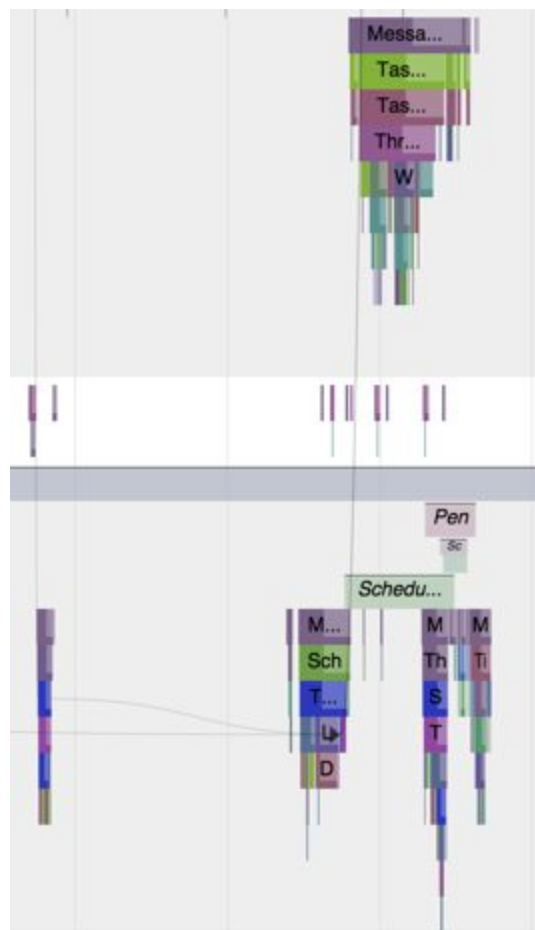
Time when the stage begins, when it ends


If it exceeds the goal, then you've had a violation

Track the # of violations, or the amount










Mes...	Me
Tas...	Tas
Th T...	Thr
Sc Sch	
v8 v...	
V V8.	


Me
Ta
Thr

actually measuring RAIL

RAILScoring

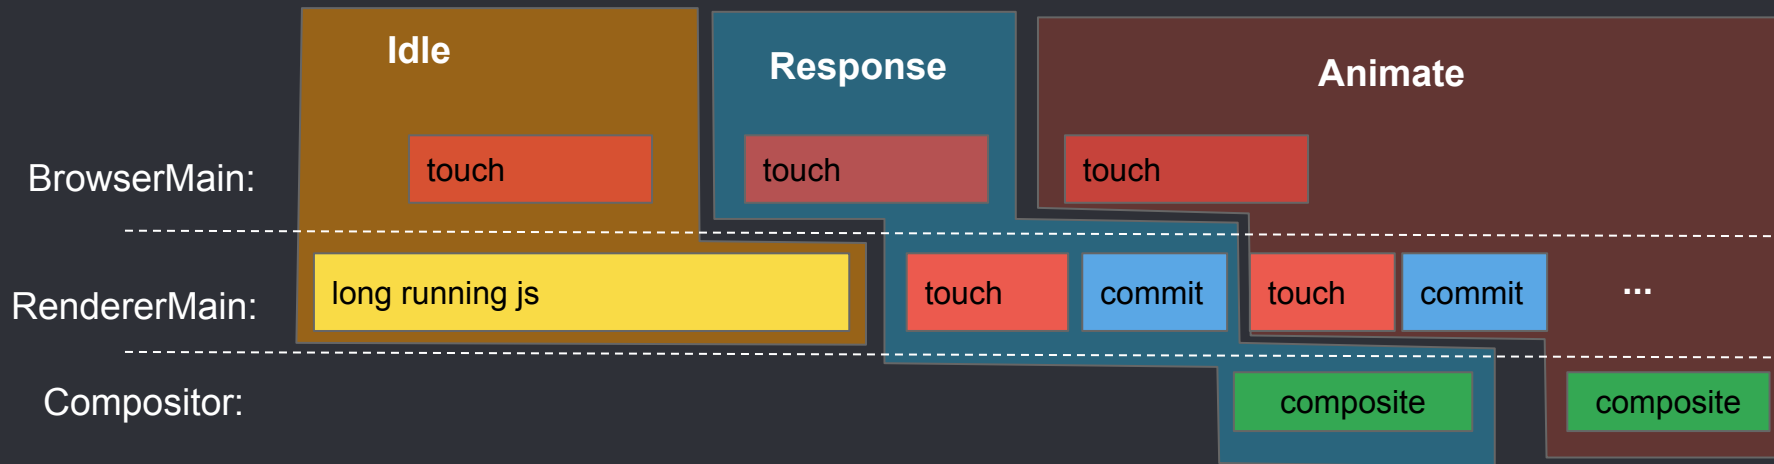


~ = RAIL Stage ... naming is hard.

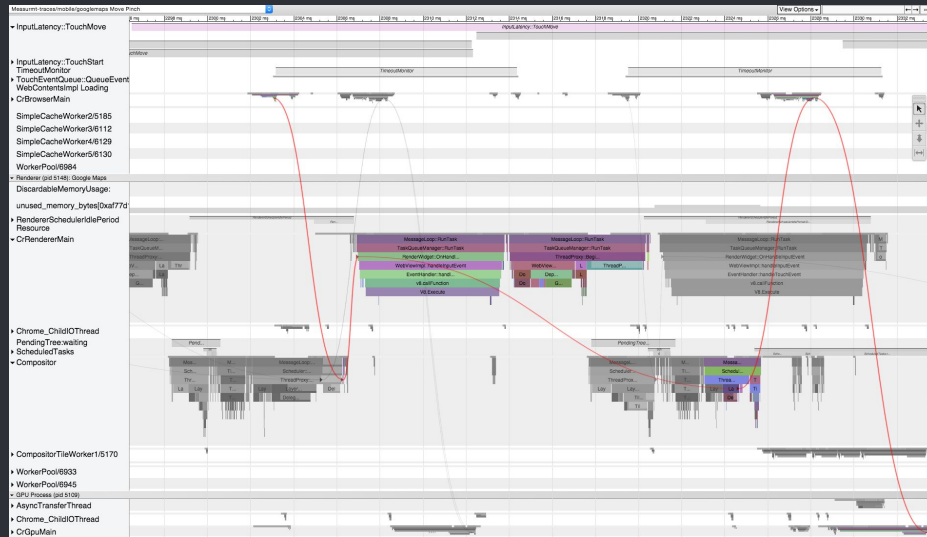
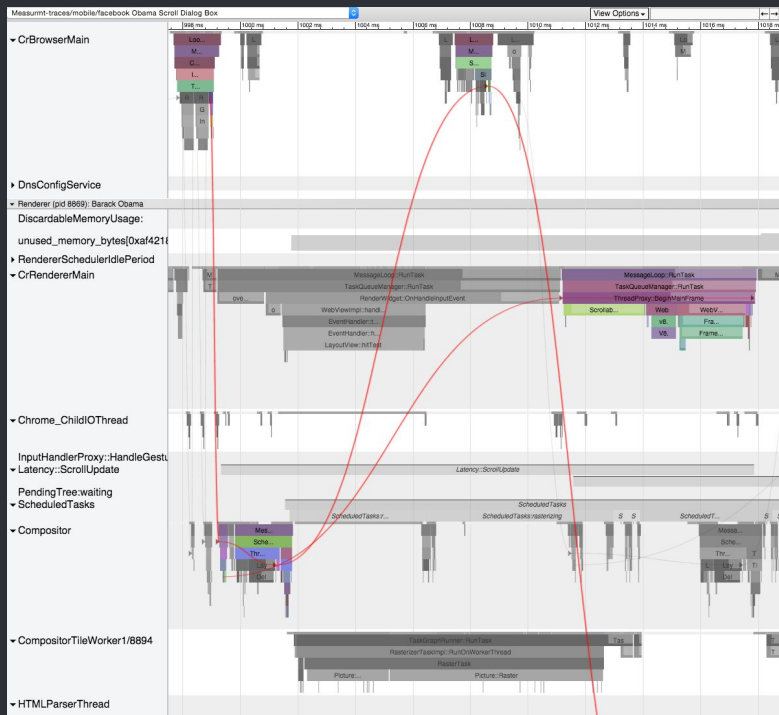


- InputLatency::MouseDown
 - InputLatency::MouseMove
 - InputLatency::MouseWheel

Threading means that IRs aren't nice pretty squares:



So we reconstruct what work is associated with an IR,
using flow events and other tricks:



Trace

IR Finding

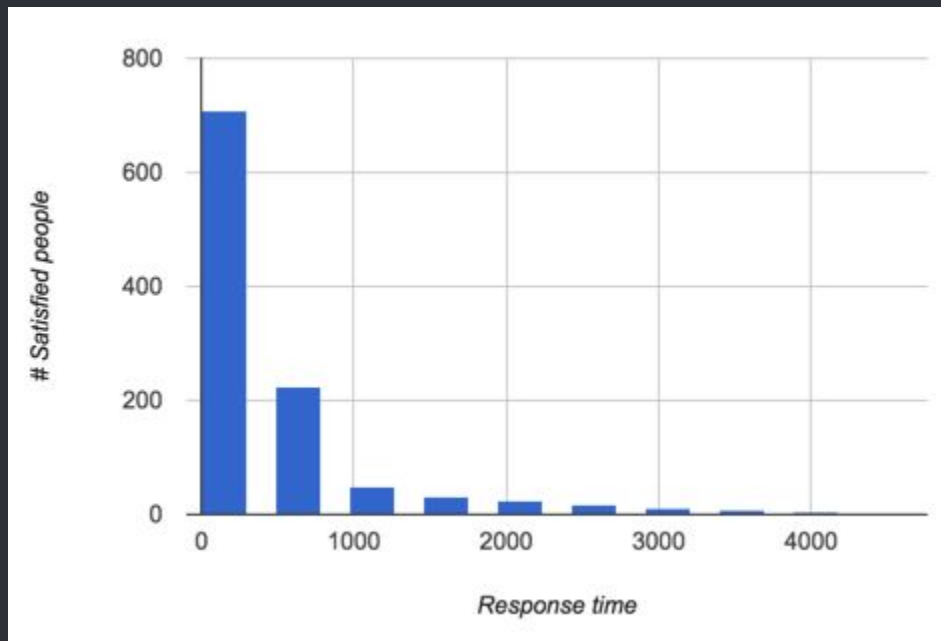
Associated Events

IR Scoring

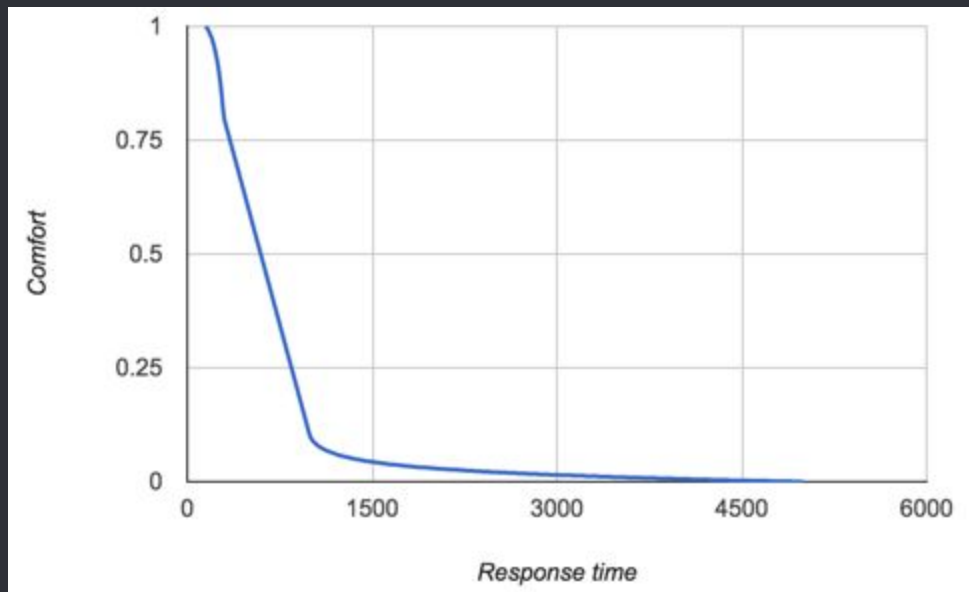
Peak-end
Scoring

Great metrics are continuous, not discrete
discrete metrics are noise amplifying.

Key insight: though 150ms response is the ideal, the underlying data is more like this:



Bit of interpolation gives you comfort:



If you're bored, you can help out!

<http://goo.gl/Z3XM79>

But what about efficiency?

Eg:



is better than




Efficiency \sim cputime used / #cpus

Nicely catches this:



Final IR scoring: blend comfort & efficiency



*we're still fiddling with weights of comfort vs efficiency

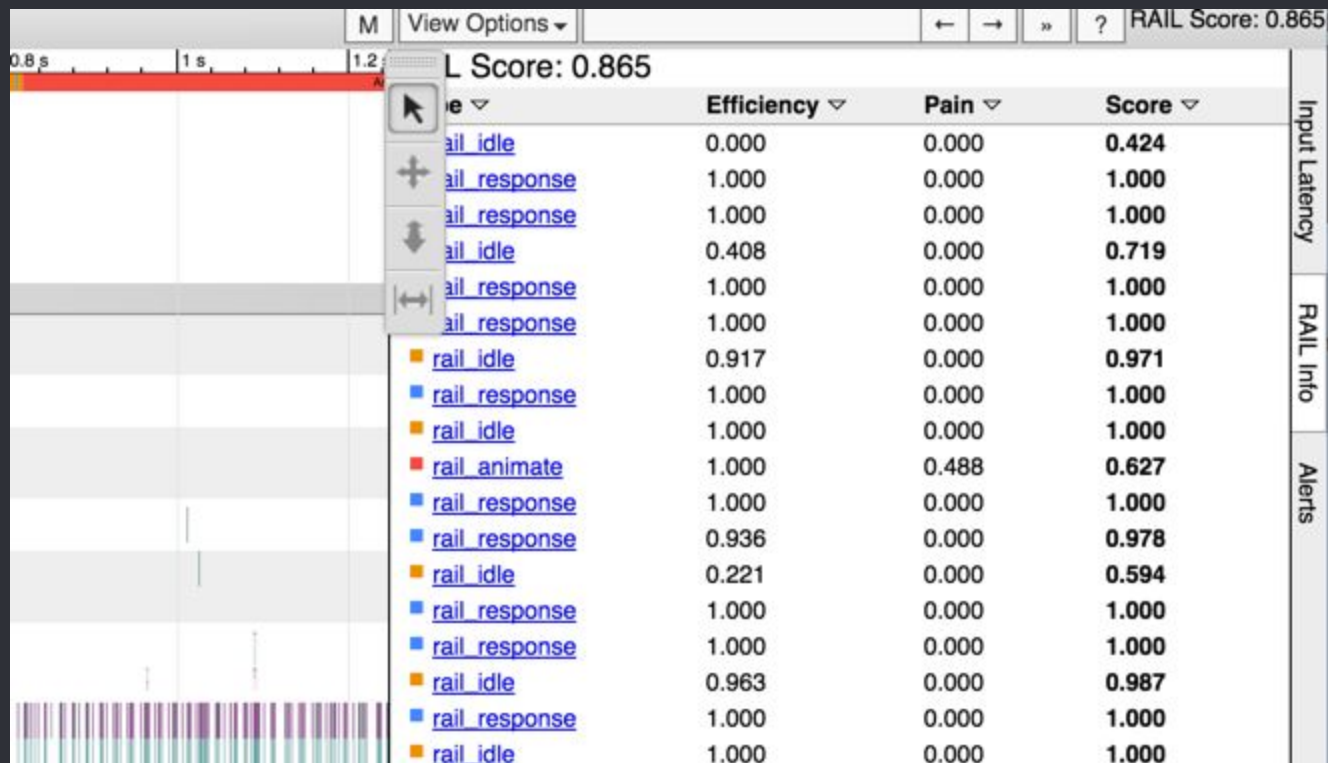
Final RAILscore based on peak end rule

~= worst experience dominates your opinion of the overall

E.g. a weighted average

- Each IR gets a score from [0-1]
- Better scores are damped, e.g. $\text{pow}(\text{ir.score}, 2)$
- Sum up and divide

Demo



Computing RAILscores at scale

Perf Insights & Pi Reports

Core idea

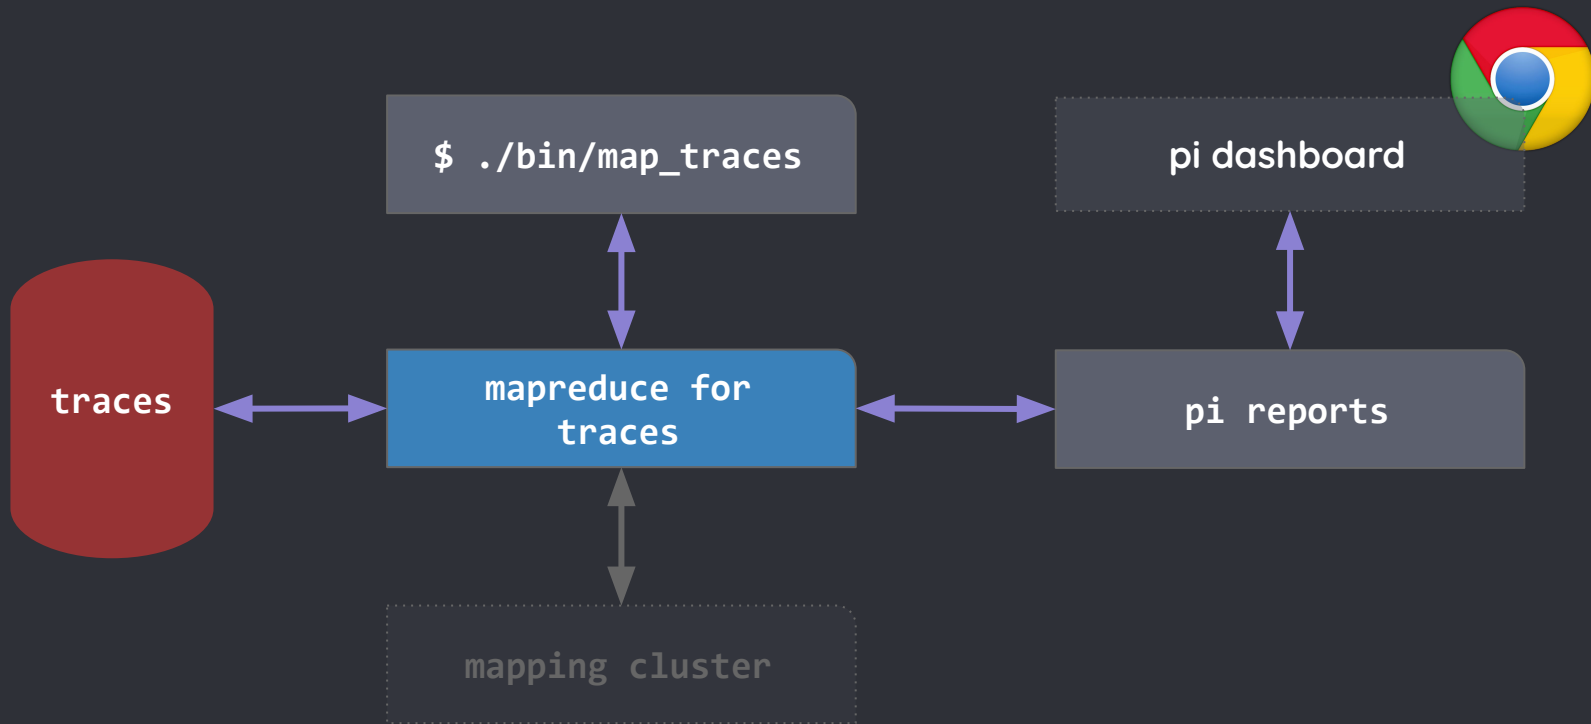
Lets bulk process traces.

LOTS of traces.

```
$ ./map_traces -j500 task_durations_histogram.html
```

```
0  -----0 (8 = 22.2%)
1  -----0 (9 = 25.0%) {22.2%}
2  -----0 (5 = 13.9%) {47.2%}
3  -----0 (2 = 5.6%) {61.1%}
4  ...
6  -----0 (2 = 5.6%) {66.7%}
7  ...
9  -----0 (3 = 8.3%) {72.2%}
10 -----0 (1 = 2.8%) {80.6%}
11 ...
```


Lẽ block diagram!



Writing a mapper

```
function myMapFunction(mapInput, mapResults, args) {  
  ...  
}
```

mapInput An object { **traceHandle**, **model** }.

model Where the real stuff lives: events, flows, metadata...

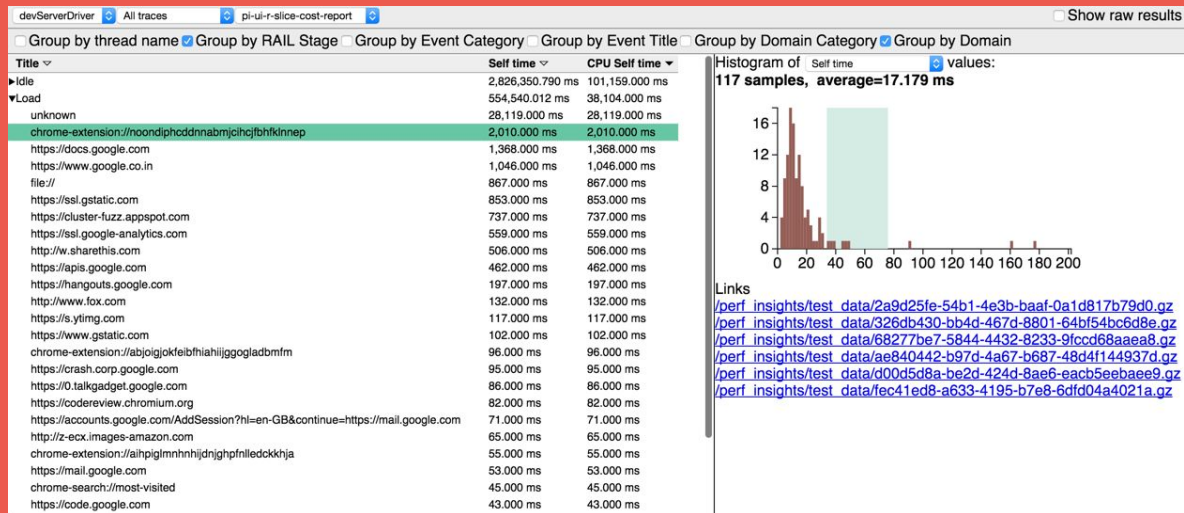
Learn about it at go/perfinsights-reports

mapResults Call **addResults(strKey, jsonableObj)** as needed.

Writing a mapper

```
function myMapFunction(mapInput, mapResults, args) {  
    mapInput.model.iterateAllSlices(function(slice) {  
        var cat = UserFriendlyCategory.getForSlice(slice);  
        if (!slice.isTopLevel) return;  
        var qd = getQueuingDuration(slice);  
        mapResults.addValue(cat, qd);  
    });  
}
```

Demo!



how to get lots of traces?



Contact fmeawad@chromium.org and rschoen@chromium.org for more information

Try map_traces yourself...

<https://goo.gl/6gZ8vC>

What next?

Memory

Battery

Emerging Markets

Great User Experience is Not Just RAIL

Security

Connectivity Resilience

Annoyance

RAIL++

User Happiness Index?

(name TBD)

And, if we can...

Analysis at a frame-level: which iframes are heaviest?

Make bulk trace analyses available to everyone, probably via a dashboard :)

Figure out how to expose some of this to web devs...