Speed Focus for 2018

Load faster, be more responsive, crash less.

tdresser@, sullivan@, bengr@

Mobile networks are slow

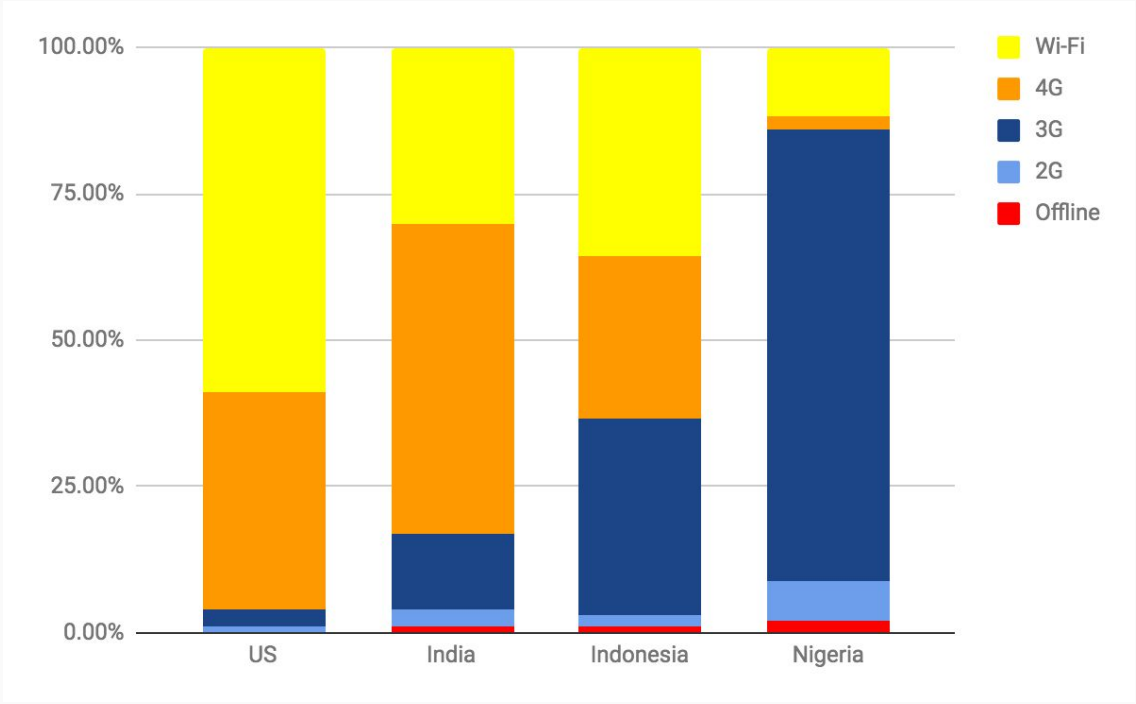

Fraction of mobile browsing sessions on each network technology

Chrome UMA (Android, Q4 2017)

Mobile devices are constrained

24% of Android devices shipped globally in 2017 H1 are **<$100**, often having **limited memory**.

IDC Quarterly Mobile Phone Tracker

Pages are bigger and more complicated

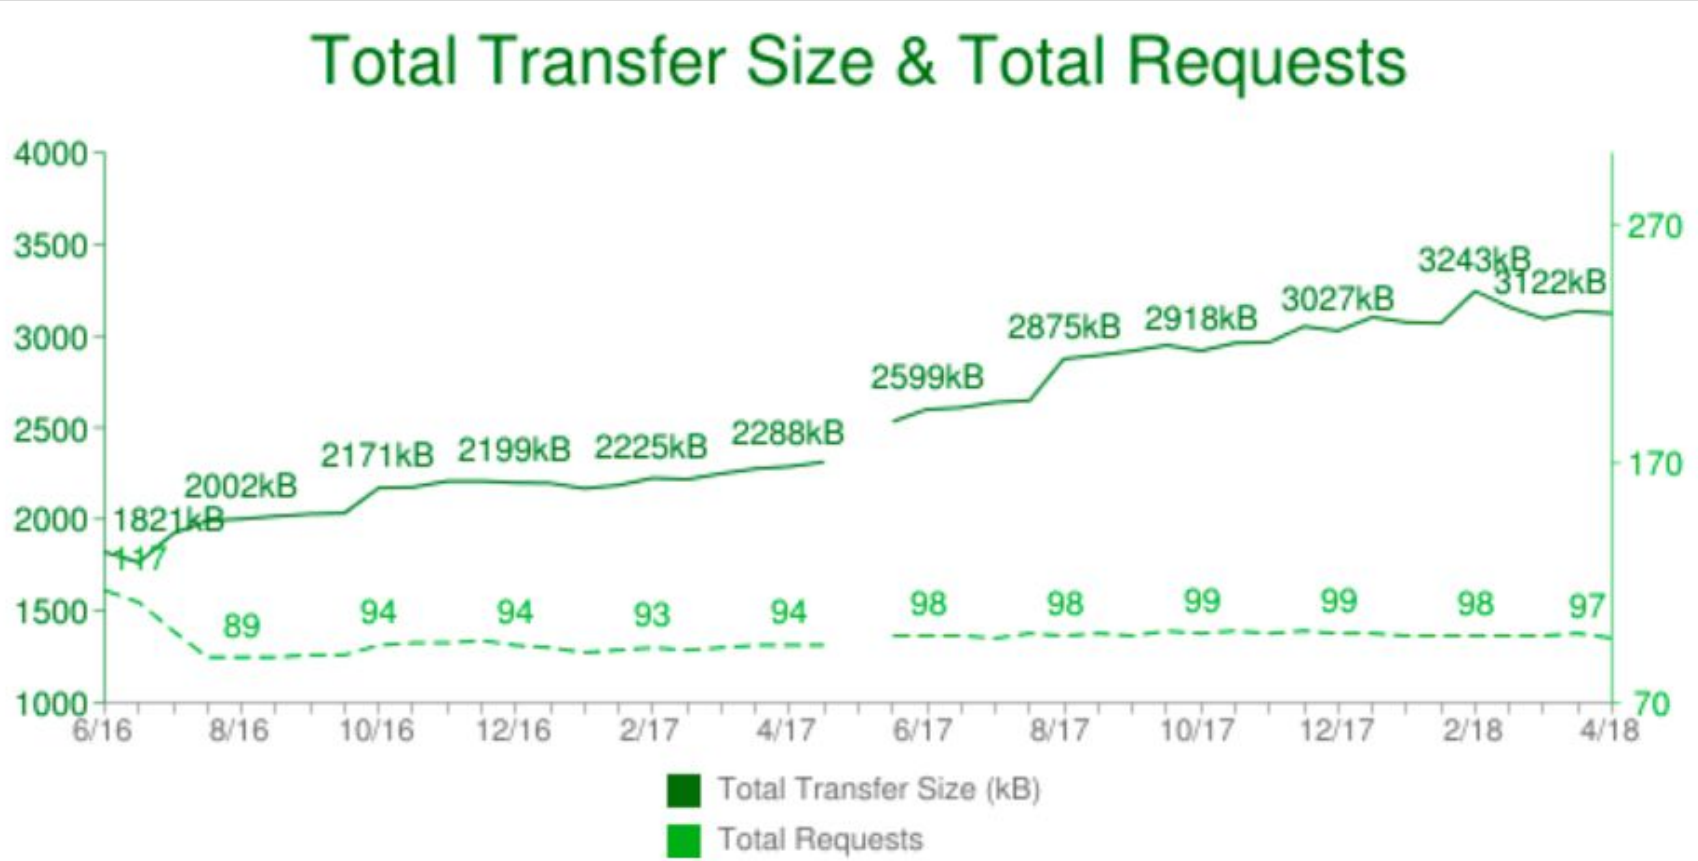

httparchive.org

Pages load slowly, are sluggish, and they crash.

How bad is it?

Example: Load latency in Nigeria (3G)

How much would you use the web if every page took this long to load?

But I thought the web was fast?

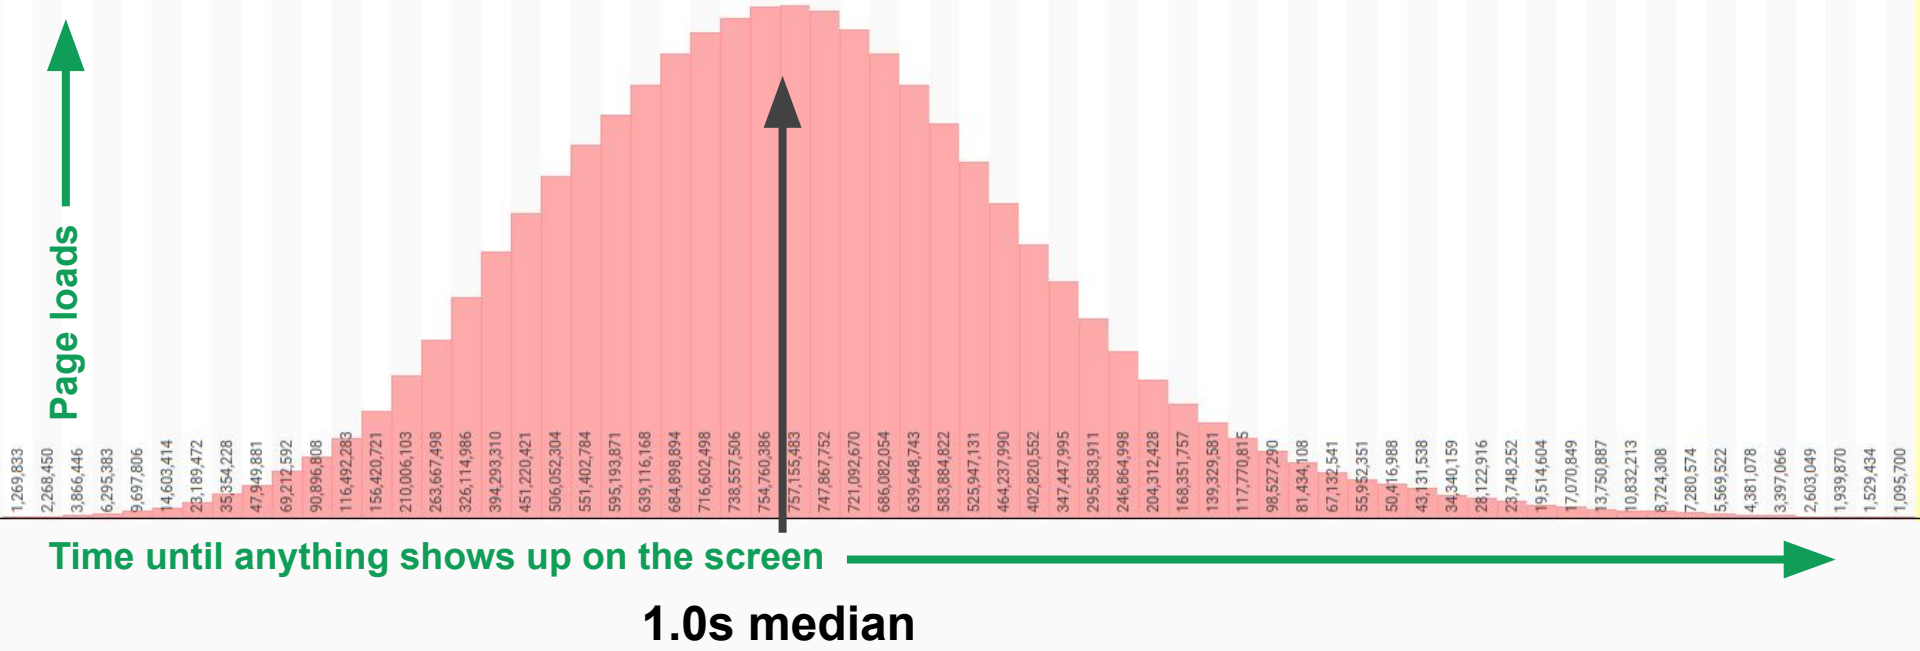

Chrome UMA (Android, Stable Channel, 4/18)

Tail latency is slow

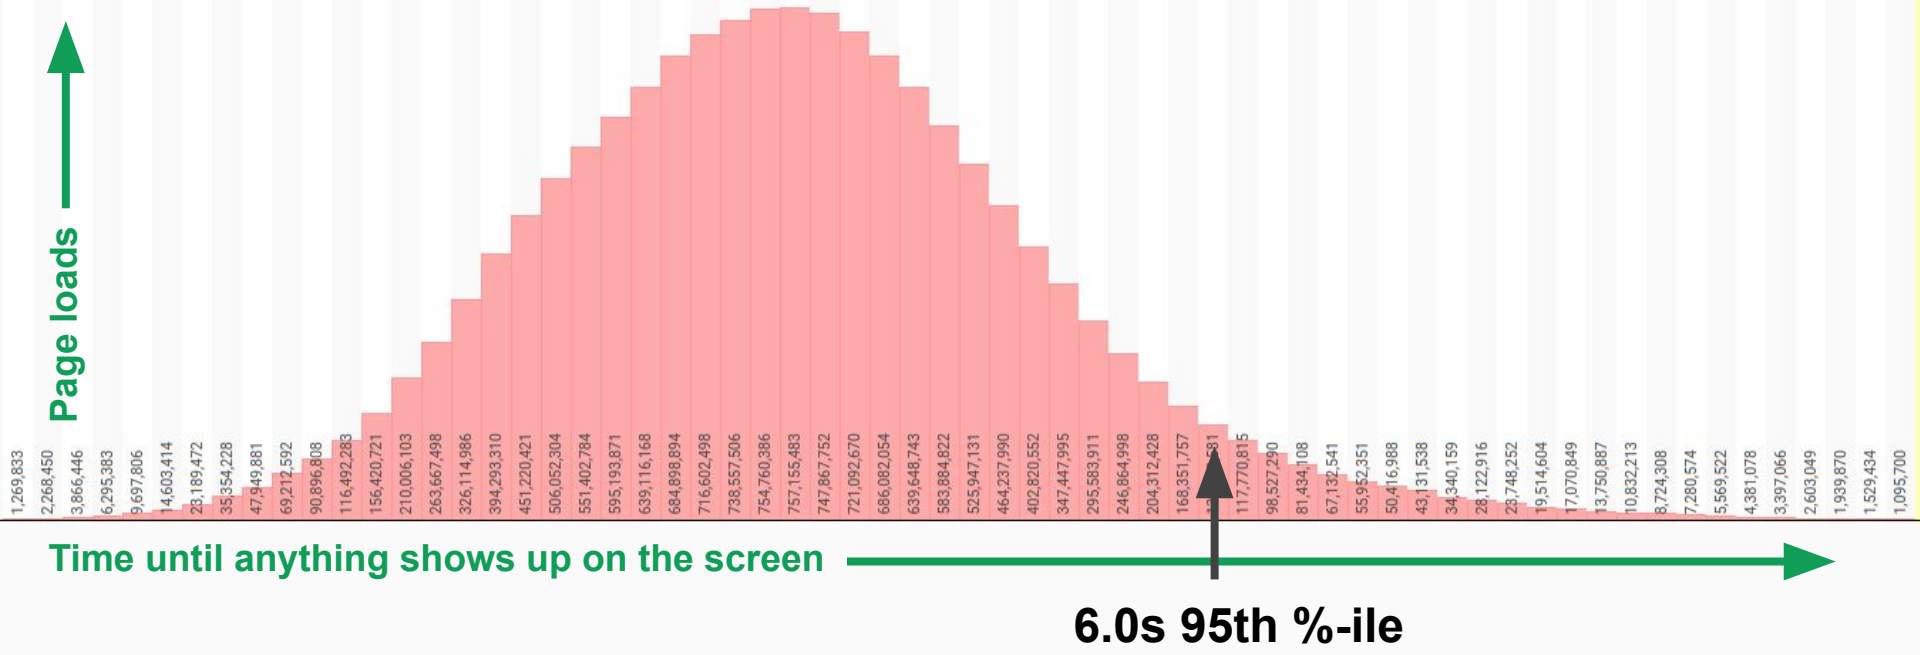

It’s even slower in emerging markets

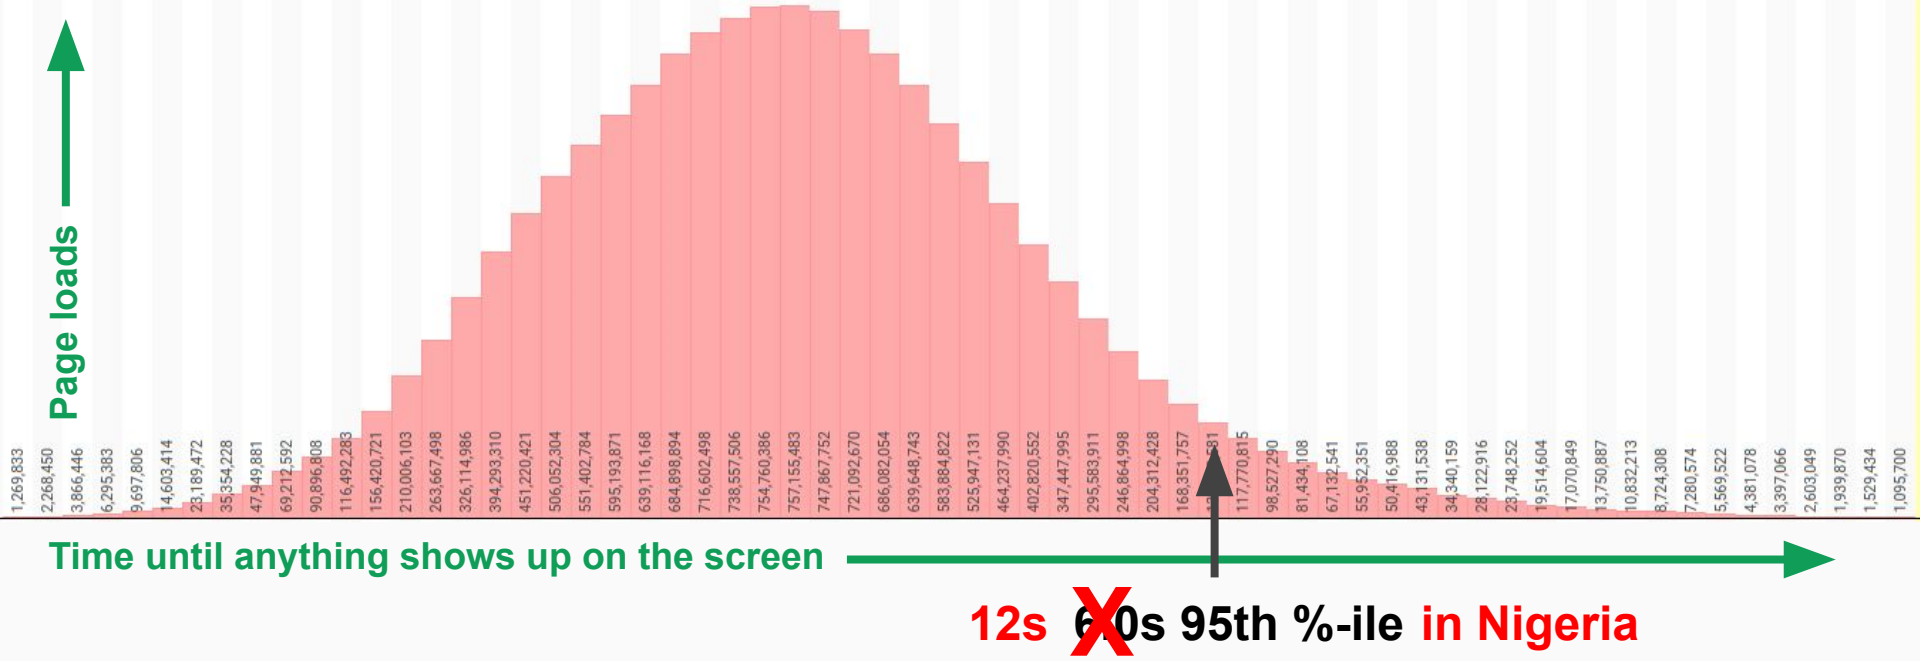

What is “Speed”?

“Speed” is all user facing aspects of performance.

How do we ensure we’re tackling the speed issues which impact users most?

2018 Speed Goals

- Load Fast
- Be Responsive
- Don’t run Out Of Memory

See details on metrics in the Speed Launch Metrics Survey.

Load Fast - **F**irst **C**ontentful **P**aint

When the browser first rendered any text, image (including background images), non-white canvas or SVG. This includes text with pending webfonts.

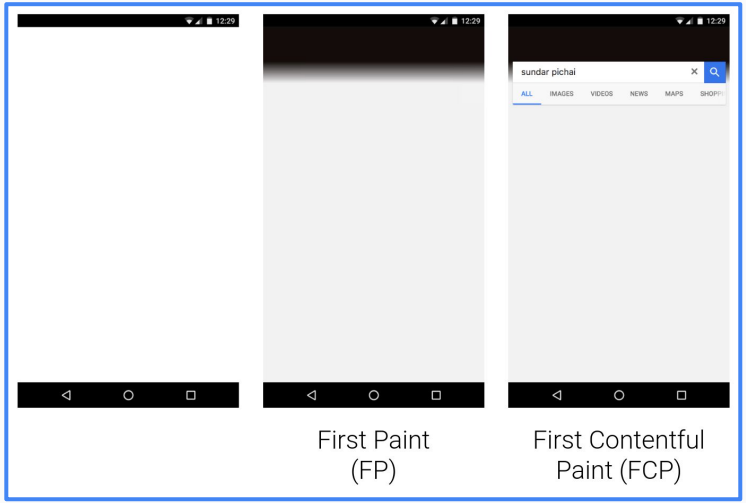

Load Fast - **F**irst **C**ontentful **P**aint

Overall Android - 99’th percentile First Contentful Paint

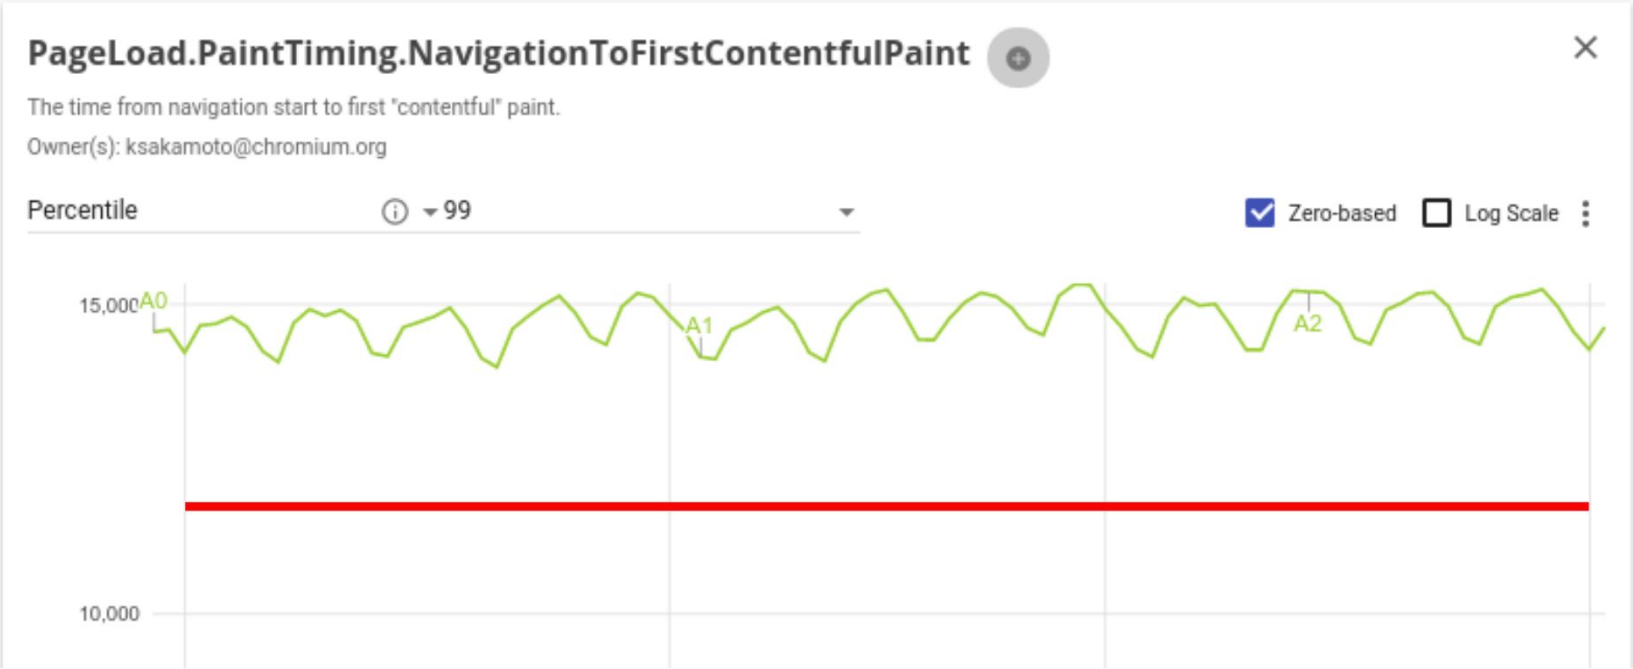

Be Responsive - **E**xpected **Q**ueueing **T**ime

When a user interacts, we want to respond instantly.

EQT: The expected duration an input event would be queued on the main thread, for an input arriving at a random time within a window.

Be Responsive - **E**xpected **Q**ueueing **T**ime

<=512MB Android - 99’th percentile Expected Queueing Time

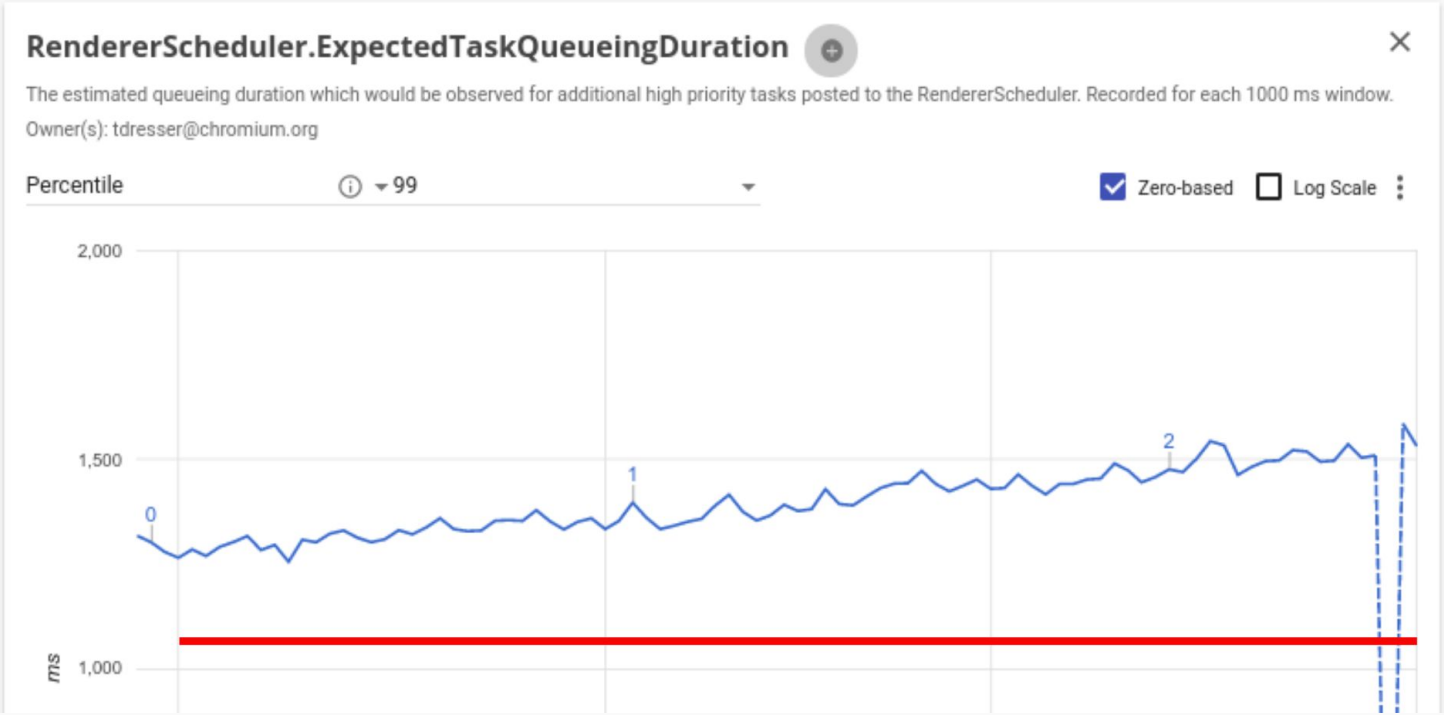

Don’t Run Out of Memory - **OOMs**

Renderer crashes from running out of memory on Android.

We track OOMs / million page loads.

Don’t Run Out of Memory - OOMs

<=512MB Android - OOMs / millino page loads

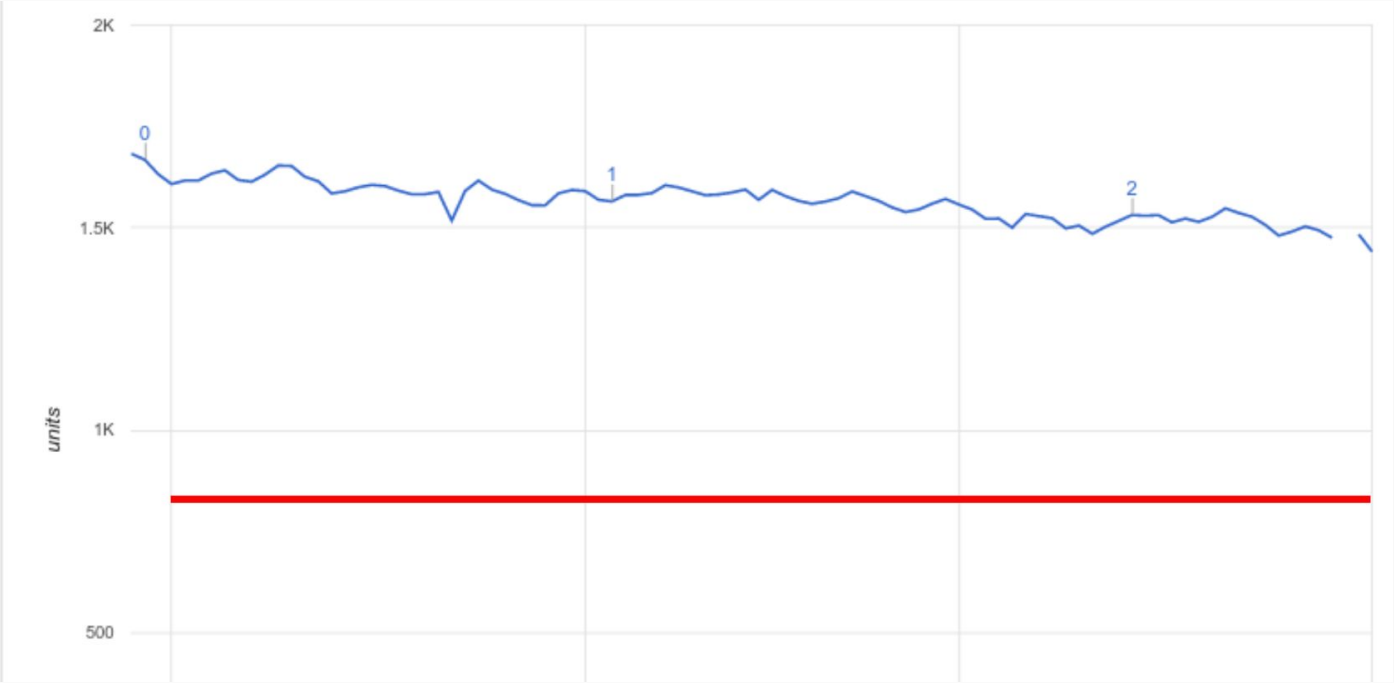

Metrics - Scroll Latency

- We track Scroll Begin Latency and Scroll Update Latency separately.
- These are measured for touch & wheel input.
- In all cases, we measure the duration from the event's hardware timestamp[1] until we update the display with the scrolled offset.

1) Except on Windows, where drivers can lie to us, and we just use the time when the browser first saw the event.

Metrics - Interactivity at Page Load

- Time to Interactive: When the page is ready to process input within 50ms.
  - Can’t be measured accurately in the wild.
- First Input Delay: The queueing delay of the first user input.
  - Correlates fairly well with TTI.

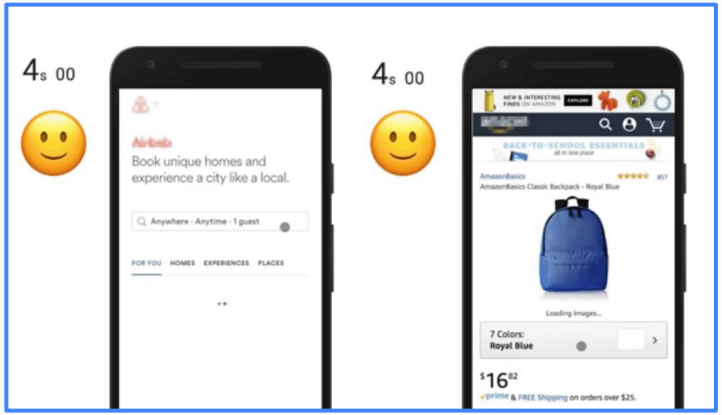

Types of Speed Projects

There are three main types of speed projects within Blink.

- Optimizations: Making existing pages faster.
- New Primitives: Shipping new features that enable faster pages.
- Activation: Making features fast enough that pages can actually use them.

Criteria for investing in speed projects

- Optimizations
  - Will this contribute to our overall speed goals?
  - Does an A/B test (via Finch) show that we're making a real improvement?
- New Primitives
  - Do we have a launch partner for an origin trial?
  - Did they see concrete improvements in metrics we care about during the origin trial?
- Activation
  - Do we have a launch partner?
  - Will they be able to adopt the feature, given the proposed performance improvements?

Evaluating impact of a performance optimization project

We evaluate the

end user impact

on the

speed launch metrics.

Start with a Finch Trial

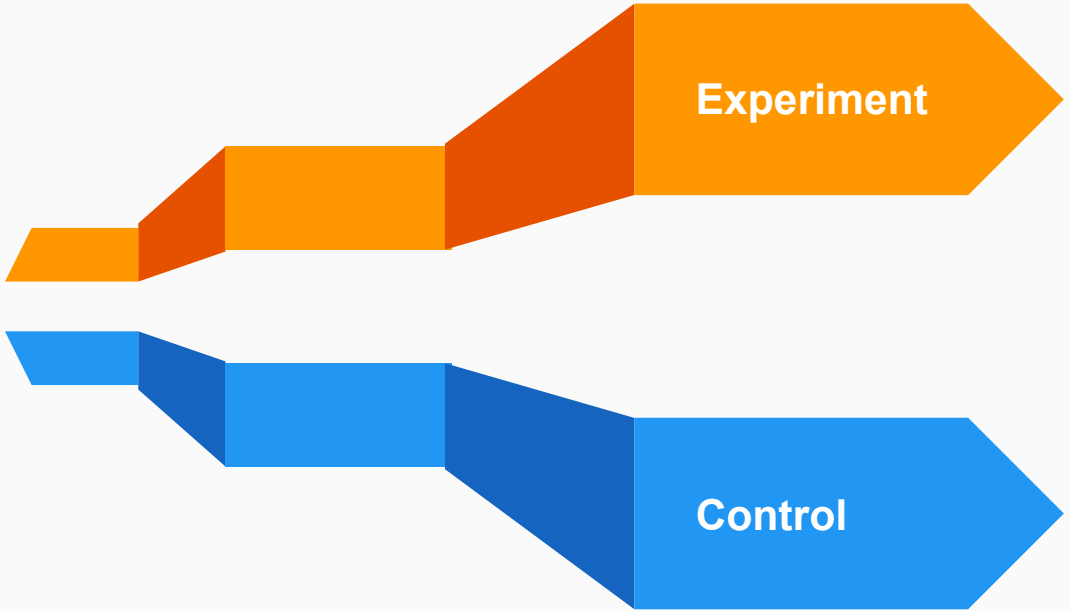

Finch Trials: Best practices

- Check for statistical significance
- Balance the right amount of volume with enabling for too many users
- Turn on experiment for **correctness and performance testing** on bots.
- More Info:
  - Googlers: **go/finch101**
  - Non-Googlers: **speed@chromium.org**

Story of a Speed Launch

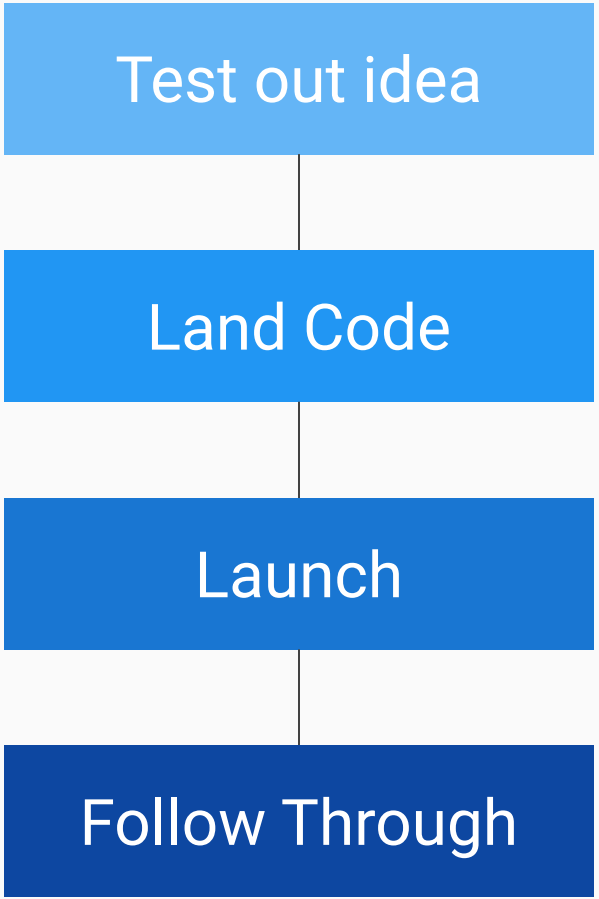

Before you submit:
Trying an idea

Understanding performance: trace viewer

Use trace-viewer to get an **understanding of Chrome's performance.**

bit.ly/trace-viewer

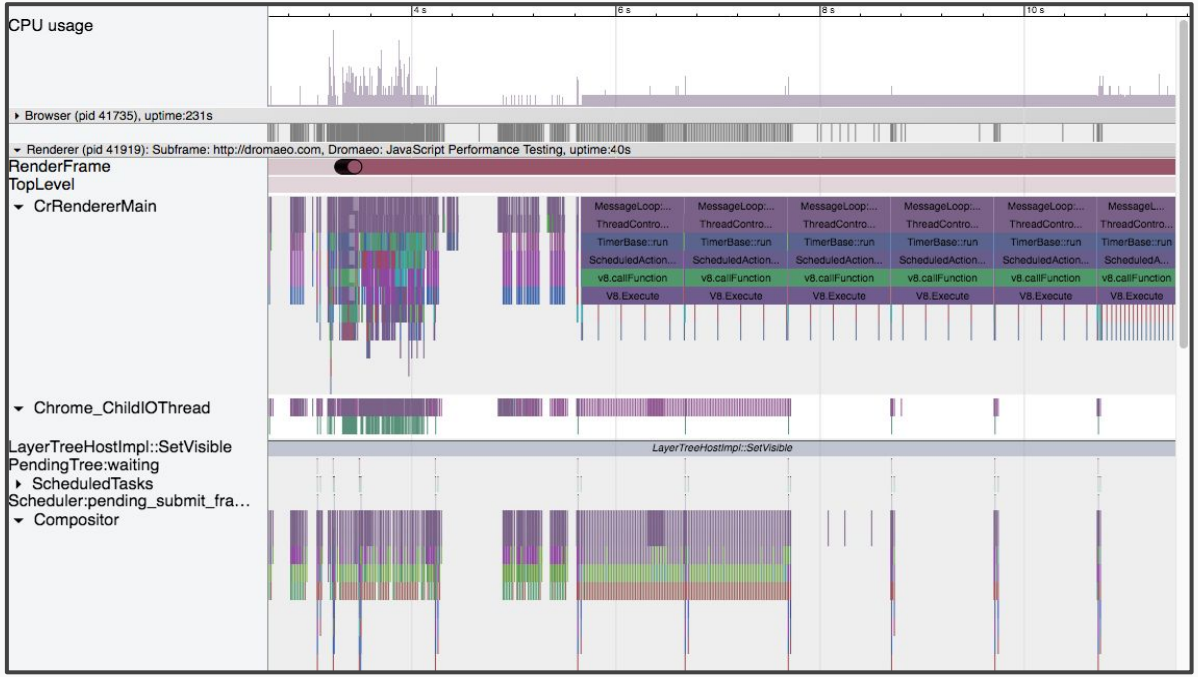

Benchmarks: system health

Run **system health benchmarks** to measure **speed launch metrics** on real web content (90 minutes of popular desktop and mobile use cases).

Use the **perf try bots** to get performance numbers from a variety of real devices across Windows, Mac, Linux, and Android.

You can view tracers either locally or if running on perf try jobs.

bit.ly/chrome-system-health-benchmarks
bit.ly/chrome-perf-trybots

Expanding to more pages: cluster telemetry

You can get **speed launch metrics** as well as many more fine-grained performance metrics on a **much larger corpus of sites** through cluster telemetry.

Traces are available as well.

Googlers: go/cluster-telemetry
Non-Googlers: ping speed@chromium.org for help

Landing the code

Landing the code: Perf Waterfall

After landing, code goes through a deeper battery of tests:

- ~30 hours of tests (bit.ly/chrome-benchmark-harnesses)
- 12 device configurations

Finch trials are turned on/off on perf waterfall via field trial testing config file: bit.ly/field-trial-testing-config

Landing the code: Perf Dashboard

Perf dashboard automatically **detects regressions**.

Perf sheriffs file bugs.

Landing the code: Perf Dashboard

See the full effects of your change at
https://chromeperf.appspot.com/group_report?rev=COMMIT_POS

		Bug ID	Revisions ▼	Master	Bot	Test Suite	Test	Delta %	Abs Delta
7		830762 ✗	548863 - 548995	ChromiumPerf	chromium-rel-win7-dual	jetstream	Score	8.3%	10.500 score
13		830771 ✗	548863 - 548995	ChromiumPerf	chromium-rel-win7-dual	blink_perf.events	hit-test-lots-of-layers	46.7%	0.256 runs/s
		830746 ✗	548863 - 548995	ChromiumPerf	chromium-rel-win7-dual	blink_perf.layout	multicol_tall-content-short-columns-realistic	11.5%	47.363 runs/s
20		830772 ✗	548844 - 548999	ChromiumPerf	android-one	thread_times.key_silk_cases	thread_other_cpu_time_per_frame/http___groupcloned.com_test_plain_sticky-using-webkit-backface-visibility.html	52.6%	0.016 ms
4			548844 - 548999	ChromiumPerf	android-one	rasterize_and_record_micro.top_25	record_time/file___static_top_25_google.html	6.6%	0.075 ms

Launching

In the wild: Finch trials

The **heartbeat tab** lists speed launch metrics impact of your finch trial.

Can slice by different platforms, device characteristics, etc.

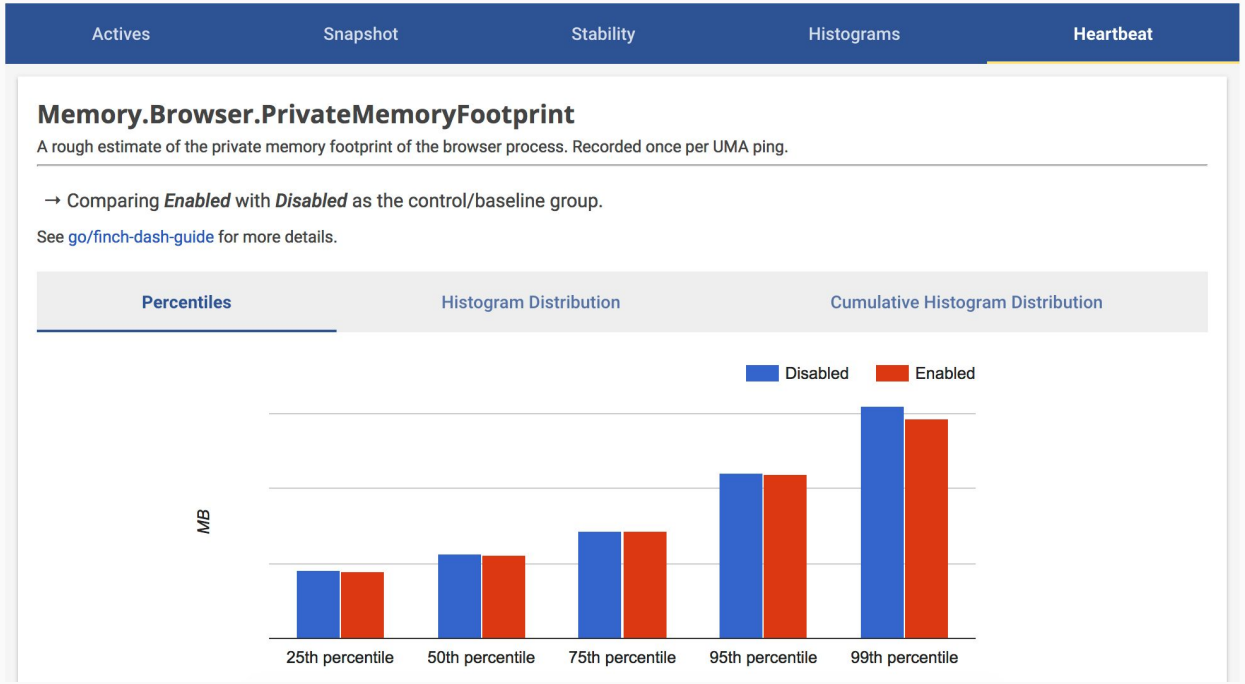

Follow Through

In the wild: We need your help!

If performance regressions slip through the cracks, we try to catch them!

- UMA speed launch metric regressions
  - Requires a lot of manual digging.
  - Very difficult to analyze. **Please use finch trials!**
- Out of process heap profiling - anonymized heap dumps from the wild
- Sampling profiler detects increases in execution time during startup

Directions to improve on Speed

- **Optimize** the loading stack and renderer
- **Intervene** on poor page construction
- Provide **new APIs** (or restrictions) to encourage using **best practices**

Talks and brainstorms:

- **Optimizing image decoding on ARM: 3 gens of silicon w/ software**
- **Accelerate graphics performance with ozone-gbm on Linux desktop systems**
- **Web Lifecycle & Interventions for CPU suspension, Throttling & Tab discarding**
- **near-OOM intervention**
- **Improving the Loading Experience (in EM)**
- **How to make developers care about memory?**
- **...**