

Optimizing Chromium image decoding for



An **inflated** story

Adenilson Cavalcanti
BS. MSc.
Staff Engineer - **Arm** San Jose (CA)

Optimizing Chromium image decoding for



Or 3 gens of silicon thanks to NEON

Adenilson Cavalcanti
BS. MSc.
Staff Engineer - **Arm** San Jose (CA)

What to optimize in Chromium



What to optimize in Chromium

- Too big.
- Too many areas.
- What would be helpful?



What to optimize in Chromium

Bulk of content still is:

- Text.
- Images.



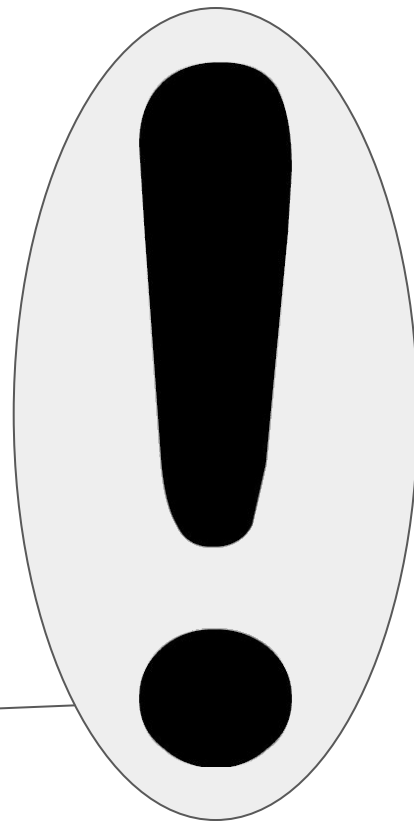
What to optimize in Chromium

Bulk of content still is:

- Text.
- Images.

Text

Image



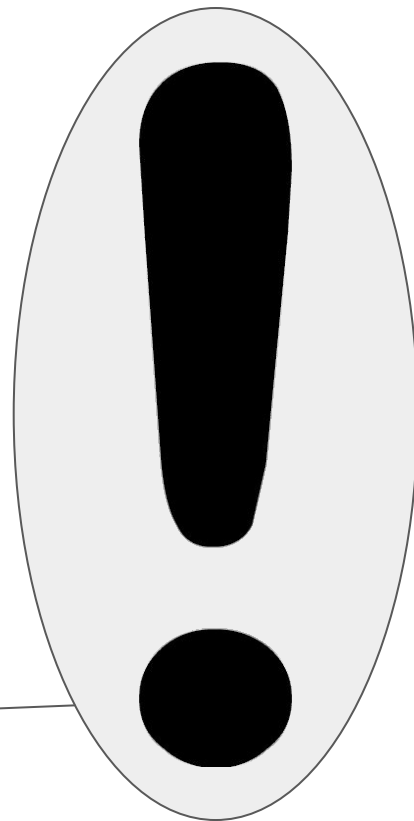
What to optimize in Chromium

Bulk of content still is:

- Text.
- Images.

Text

Image



PNG

- Powerful format: Palette, pre-filters, compressed.
- Encoder affects behavior.
- Libpng and zlib are 'Bros!'.

Meet Mr. Parrot



Source: <https://upload.wikimedia.org/wikipedia/commons/3/3f/ZebraHighRes.png>

Parrots are not created equal



Parrots are not created equal



Original: 2.7MB



Palette: 0.8MB



Zopfli: 2.6MB

Features affect hotspots

== Image has pre-compression filters (2.7MB) ==

Lib	Command	SharedObj	method	CPU (%)
zlib	TileWorker	liblink	inflate_fast	1.96
zlib	TileWorker	libblnk	adler32	0.88
blink	TileWorker	liblink	ImageFrame::setRGBAPremultiply ..	0.45
blink	TileWorker	liblink	png_read_filter_row_up.....	0.03*

== Image was optimized using zopfli (2.6MB) ==

Lib	Command	SharedObj	method	CPU (%)
zlib	TileWorker	liblink	inflate_fast	3.06
zlib	TileWorker	libblnk	adler32	1.36
blink	TileWorker	liblink	ImageFrame::setRGBAPremultiply ..	0.70
blink	TileWorker	liblink	png_read_filter_row_up.....	0.48*

== Image has no pre-compression filters (0.9MB) ==

Lib	Command	SharedObj	method	CPU (%)
libpng	TileWorker	liblink	cr_png_do_expand_palette	0.88
zlib	TileWorker	liblink	inflate_fast	0.62
blink	TileWorker	liblink	ImageFrame::setRGBAPremultiply ..	0.49
zlib	TileWorker	libblnk	adler32	0.31

Candidates

- Inflate_fast: **zlib**.
- Adler32: **zlib**.
- ImageFrame: Blink.
- png_do_expand_palette:
libpng.
- Crc32: **zlib**.

Why zlib?

Zlib

Used everywhere (libpng, Skia, freetype, **cronet**, blink, chrome, linux kernel, etc).

Old code base released in 1995.

Written in K&R C style.

Context

Lacks any optimizations for ARM CPUs.

Problem statement

Identify potential optimization candidates and verify positive effects in Chromium.

Implementation

Adler-32

$$A = 1 + D_1 + D_2 + \dots + D_n \pmod{65521}$$

$$B = (1 + D_1) + (1 + D_1 + D_2) + \dots + (1 + D_1 + D_2 + \dots + D_n) \pmod{65521}$$

$$= n \times D_1 + (n-1) \times D_2 + (n-2) \times D_3 + \dots + D_n + n \pmod{65521}$$

$$\text{Adler-32}(D) = B \times 65536 + A$$

Adler-32: simplistic implementation

```
// From: https://en.wikipedia.org/wiki/Adler-32
const int MOD_ADLER = 65521;
unsigned long naive_adler32(unsigned char *data,
                           unsigned long len)
{
    uint32_t a = 1, b = 0;
    unsigned long index;

    for (index = 0; index < len; ++index) {
        a = (a + data[index]) % MOD_ADLER;
        b = (b + a) % MOD_ADLER;
    }

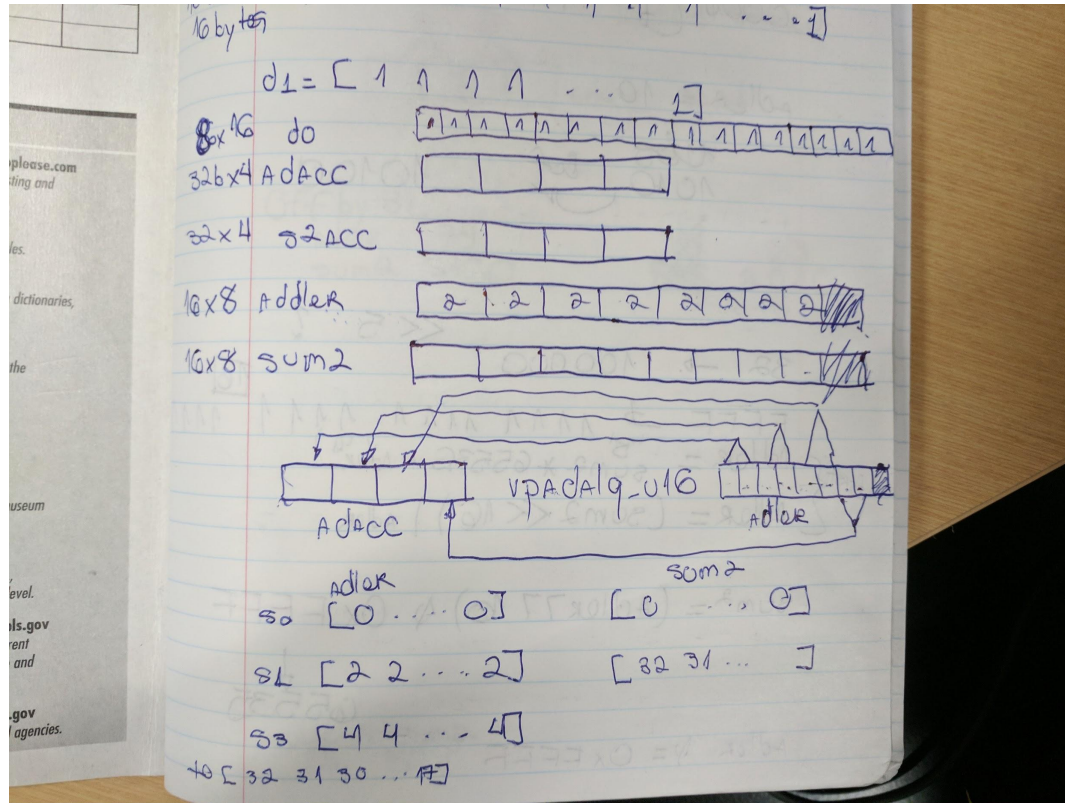
    return (b << 16) | a;
}
```

Problems

- Zlib's Adler-32 was more than **7x faster** than naive implementation.
- It is hard to vectorize the following computation:

```
void accum(uint32_t *pair, const unsigned char *buf,
           unsigned int len)
{
    unsigned int i;
    for (i = 0; i < len; ++i) {
        pair[0] += buf[i];
        pair[1] += pair[0];
    }
}
```

Highly technical drawing (Jan 2017)



'Taps' to the rescue

Assembly:

<https://godbolt.org/g/KMeBAJ>

```
static void NEON_accum32(uint32_t *s, const unsigned char *buf,
                        unsigned int len)
{
    static const uint8_t taps[32] = {
        32, 31, 30, 29, 28, 27, 26, 25,
        24, 23, 22, 21, 20, 19, 18, 17,
        16, 15, 14, 13, 12, 11, 10, 9,
        8, 7, 6, 5, 4, 3, 2, 1 };

    uint32x2_t adacc2, s2acc2, as;
    uint8x16_t t0 = vld1q_u8(taps), t1 = vld1q_u8(taps + 16);

    uint32x4_t adacc = vdupq_n_u32(0), s2acc = vdupq_n_u32(0);
    adacc = vsetq_lane_u32(s[0], adacc, 0);
    s2acc = vsetq_lane_u32(s[1], s2acc, 0);

    while (len >= 2) {
        uint8x16_t d0 = vld1q_u8(buf), d1 = vld1q_u8(buf + 16);
        uint16x8_t adler, sum2;
        s2acc = vaddq_u32(s2acc, vshlq_n_u32(adacc, 5));
        adler = vpaddlq_u8(d0);
        adler = vpadalq_u8(adler, d1);
        sum2 = vmull_u8(vget_low_u8(t0), vget_low_u8(d0));
        sum2 = vmlal_u8(sum2, vget_high_u8(t0), vget_high_u8(d0));
        sum2 = vmlal_u8(sum2, vget_low_u8(t1), vget_low_u8(d1));
        sum2 = vmlal_u8(sum2, vget_high_u8(t1), vget_high_u8(d1));
        adacc = vpadalq_u16(adacc, adler);
        s2acc = vpadalq_u16(s2acc, sum2);
        len -= 2;
        buf += 32;
    }
}
```

Happy end! Up to 12% performance gain in PNG

```
commit d400b38e450a71e9ed75bd4c1dda6329267e9ce0
```

```
Author: Adenilson Cavalcanti <adenilson.cavalcanti@arm.com>
```

```
Date: Mon Jan 30 15:30:38 2017 -0800
```

NEON implementation for Adler32

The checksum is calculated in the uncompressed PNG data and can be made much faster by using SIMD.

Tests in ARMv8 yielded an improvement of about 3x (e.g. walltime was 350ms x 125ms for a 4096x4096 bytes executed 30 times).

This alone yields a performance boost for PNG decoding ranging from 5% to 18% depending on a few factors (SoC, battery status, big/little, etc).

<https://bugs.chromium.org/p/chromium/issues/detail?id=688601>

Shipping on M63: Intel got some love too!

author Noel Gordon <noel@chromium.org> Fri Sep 29
committer Commit Bot <commit-bot@chromium.org> Fri Sep 29
tree [a25de9dd3212b49c1d903e72289e424b72127c3e](#)
parent [6baf6221674f5a075f12f83e4262a4751b5d445b](#) [diff]

zlib adler_simd.c

Add SSE3 implementation of the Adler32 checksum, suitable for both large workloads, and small workloads commonly seen during PNG image decoding. Add a NEON implementation.

Speed is comparable to the serial Adler32 computation but near 64 bytes of input data, the SIMD code paths begin to be faster than the serial path: 3x faster at 256 bytes of input data, to ~8x faster for 1M of input data (~4x on ARMv8 NEON).

For the PNG 140 image corpus, PNG decoding speed is ~8% faster on average on the desktop machines tested, and ~2% on an ARMv8 Pixel C Android (N) tablet, <https://crbug.com/762564#c41>

Update x86.{c,h} to runtime detect SSE3 support and use it to enable the Adler32_simd code path and update inflate.c to call x86_check_features(). Update the name mangler file names.h for the new symbols added, add FIXME about simd.patch.

Ignore data alignment in the SSE3 case since unaligned access is no longer penalized on current generation Intel CPU. Use it in the NEON case however to avoid the extra costs of unaligned memory access on ARMv8/v7.

NEON credits: the v_s1/s2 vector component accumulate code was provided by Adenilson Cavalcanti. The uint16 column vector sum code is from libdeflate with corrections to process NMAX input bytes which improves performance by 3% for large buffers.

<https://bugs.chromium.org/p/chromium/issues/detail?id=688601>

Inffast (Simon Hosie)

- Second candidate in the perf profiling was **inflate_fast**.
- Very **high level** idea: perform long loads/stores in the byte array.
- Average **20% faster!**
- Shipping on M62.

```
+                                     */
+                                     out = chunkcopy_safe(out, from, len, limit);
+                                     }
+                                 }
+                                 else {
-                                     from = out - dist;          /* copy direct from output */
-                                     do {                          /* minimum length is three */
-                                         *out++ = *from++;
-                                         *out++ = *from++;
-                                         *out++ = *from++;
-                                         len -= 3;
-                                     } while (len > 2);
-                                     if (len) {
-                                         *out++ = *from++;
-                                         if (len > 1)
-                                             *out++ = *from++;
-                                     }
+                                     /* Whole reference is in range of current output. No
+                                     range checks are necessary because we start with room
+                                     for at least 258 bytes of output, so unroll and roundoff
+                                     operations can write beyond `out+len` so long as they
+                                     stay within 258 bytes of `out`.
+                                     */
+                                     out = chunkcopy_lapped_relaxed(out, dist, len);
+                                 }
+                             }
```

<https://bugs.chromium.org/p/chromium/issues/detail?id=697280>

Libpng (Richard Townsend)

- NEON optimization in libpng.
- From **10 to 30% improvement**.
- Depends on png using a palette.
- Shipping on M66.

Blink::setRGBAPreMultiply (Jonathan Wright)

- NEON optimization in Blink.
- Around **9% improvement**.
- Shipping on M66.

CRC32: zlib

- YMMV on PNGs (from 1 to 5%).
- Remember it is used while **decompressing** web content (29% boost for gzipped content).
- ARMv8-a has a crc32 instruction (from 3 to 10x faster than zlib's crc32 C code).
- Shipping on M66.

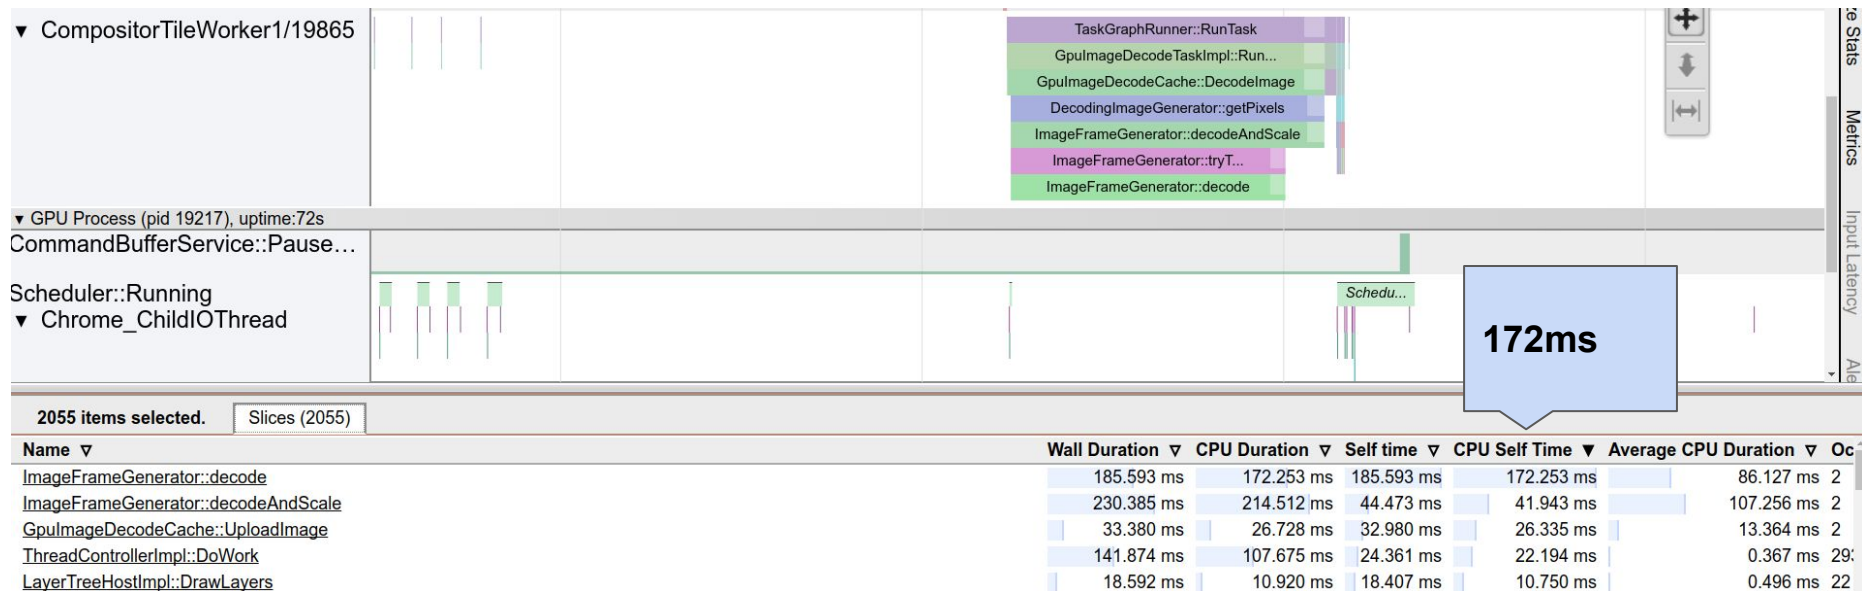
<https://bugs.chromium.org/p/chromium/issues/detail?id=709716>

Impact

Combined effect of 5 patches

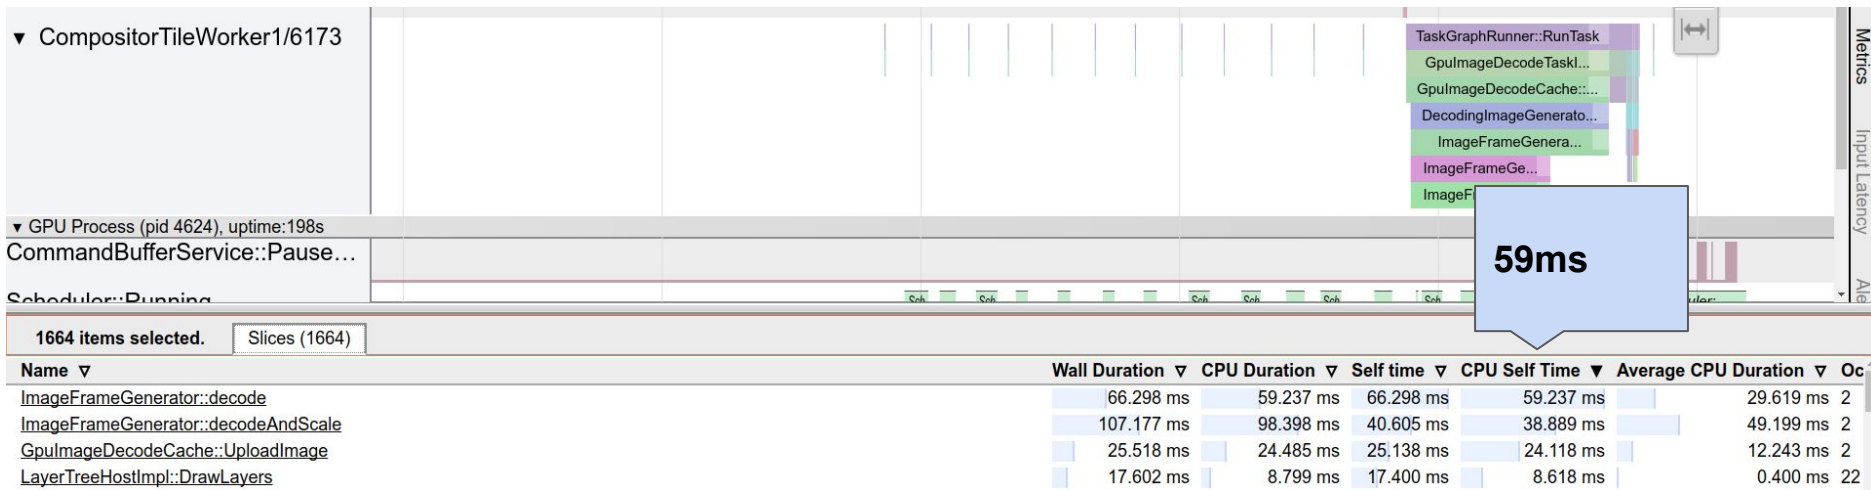


No patches: Chrome M66 trace in a Moto X4



172ms spent decoding parrot PNG with palette.

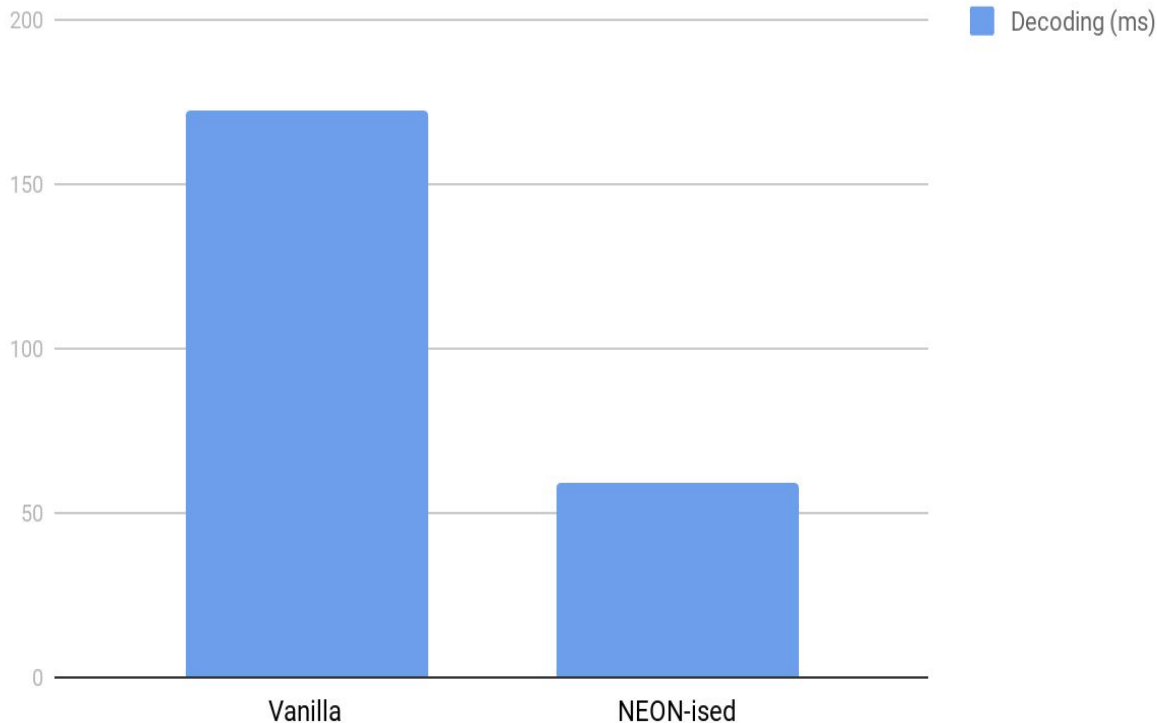
NEON: Chrome M66 trace in a Moto X4



59ms spent decoding parrot PNG with palette.

PNG decoding: 172ms X 59ms (or 2.9x faster)

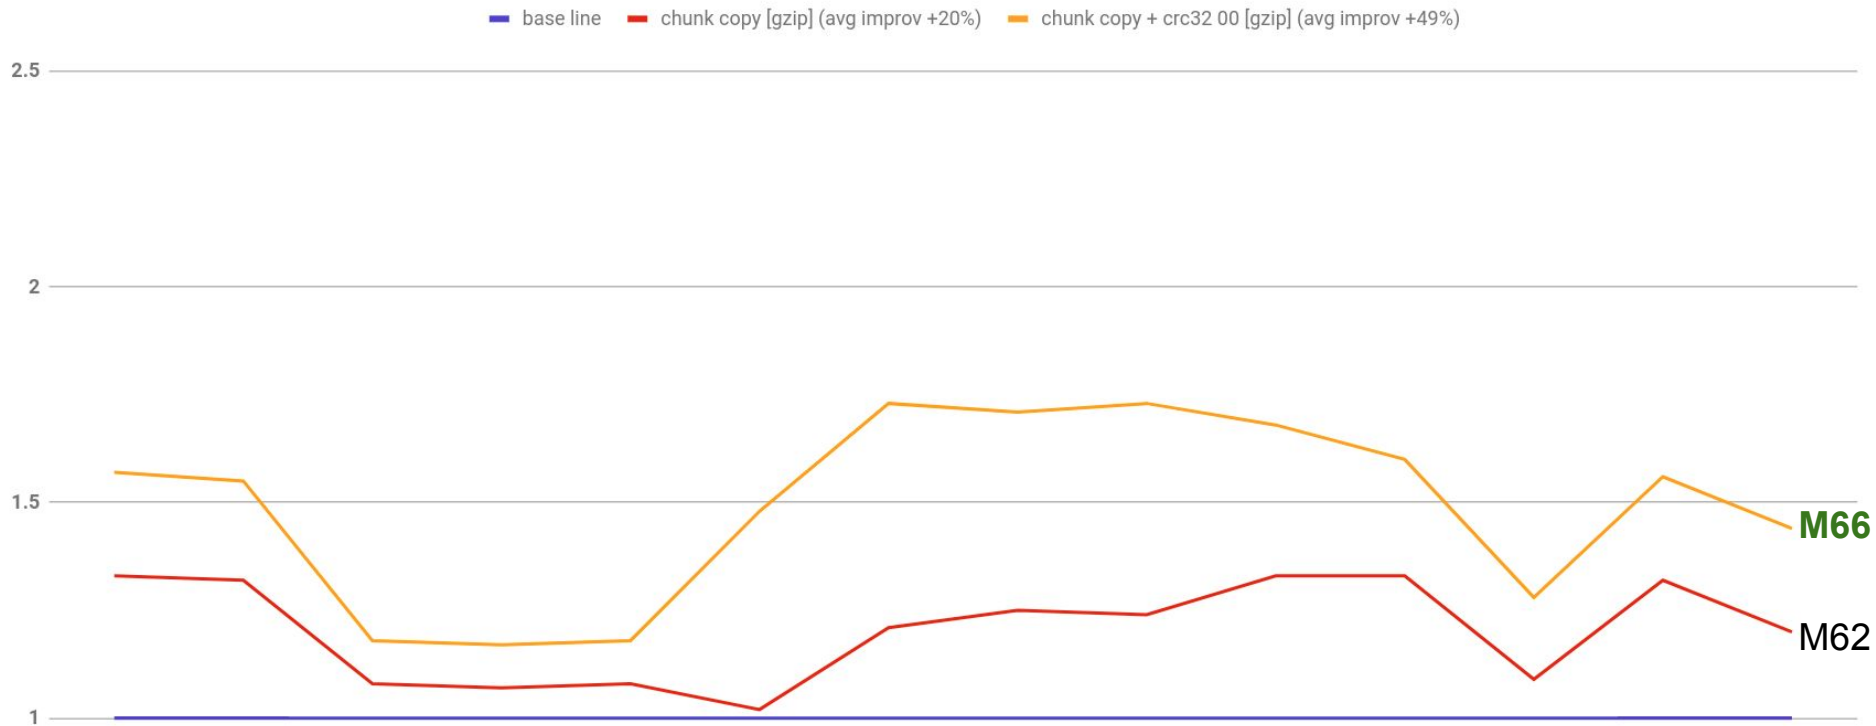
Decoding time (less is better)



Decompression: avg 1.4x for zlib format



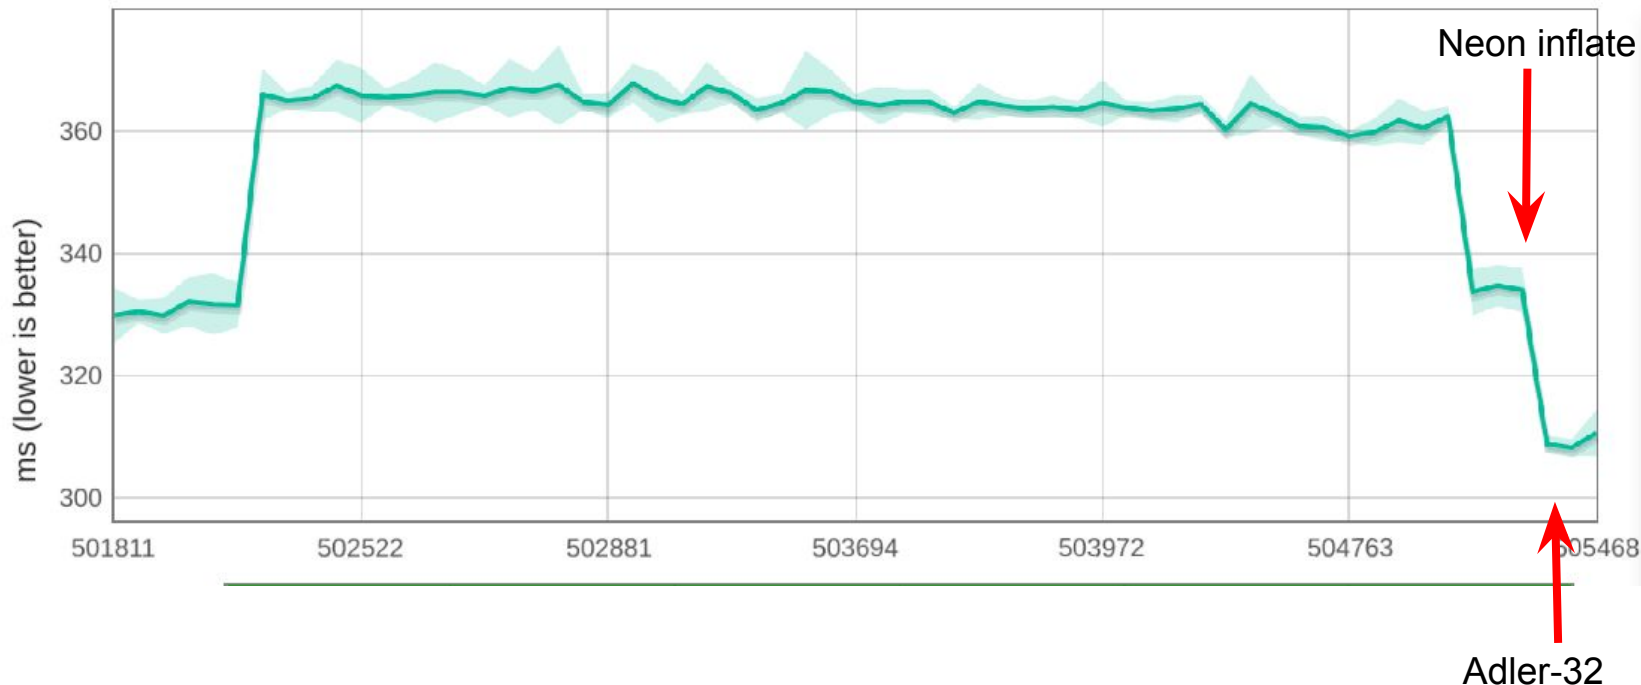
Decompression: avg 1.5x for gzip format



Possible change on compression research? Link:

<https://cran.r-project.org/web/packages/brotli/vignettes/brotli-2015-09-22.pdf>

Google bots confirm the effect of optimizations



Lessons learned

- **arm** cores can benefit **a lot** from NEON optimizations.
- Performance gains of up to 3 generations* of silicon.
- It pays off to work in a lower software layer (e.g. zlib/libpng).
- Upstreaming can be **really** slow.

What comes next

- Zlib: Compression.
- Port optimizations to Android?

Zlib users should consider migrating to Chromium's zlib.

Special Thanks

- **Arm** for sponsoring this research.
- Google: Chris Blume (reviews, fuzzers, perf bots, GMeet debugging*), Mike Klein (SIMD), Noel Gordon (tools, Intel port), Leon Scroggins (libpng).
- Team **Arm@UK**: Dave Rodgman, Richard Townsend, Stephen Kyle, Jonathan Wright.
- Team **Arm@US**: Adenilson Cavalcanti (Simon Hosie@google, Amaury Leleyzour@unity).
- Compiler explorer: <https://godbolt.org>

Questions?



Different kinds of feedback

